

CS3223 Lecture 1

Data Storage


What does a DBMS store?

- ▶ Relations
- ▶ System catalog (a.k.a. data dictionary) stores metadata about relations
 - ▶ Relation schemas - structure of relations, constraints, triggers
 - ▶ View definitions
 - ▶ Indexes - derived information to speed up access to relations
 - ▶ Statistical information about relations for use by query optimizer
- ▶ Log files - information maintained for data recovery

Memory Hierarchy


- ▶ **Primary memory**: registers, static RAM (caches), dynamic RAM (physical memory)
- ▶ **Secondary memory**: magnetic disks (HDD), solid-state disks (SSD)
- ▶ **Tertiary memory**: optical disks, tapes
- ▶ Tradeoffs:
 - ▶ capacity
 - ▶ cost
 - ▶ access speed
 - ▶ volatile vs non-volatile

Memory Hierarchy: Tradeoffs




Source: Hennessy & Patterson's Computer Architecture: A Quantitative Approach

DBMS Storage



- ▶ DBMS stores data on non-volatile disk for persistence
- ▶ DBMS processes data in main memory (RAM)
- ▶ Disk access operations:
 - ▶ read: transfer data from disk to RAM
 - ▶ write: transfer data from RAM to disk

Magnetic Hard-Disk Drive (HDD)

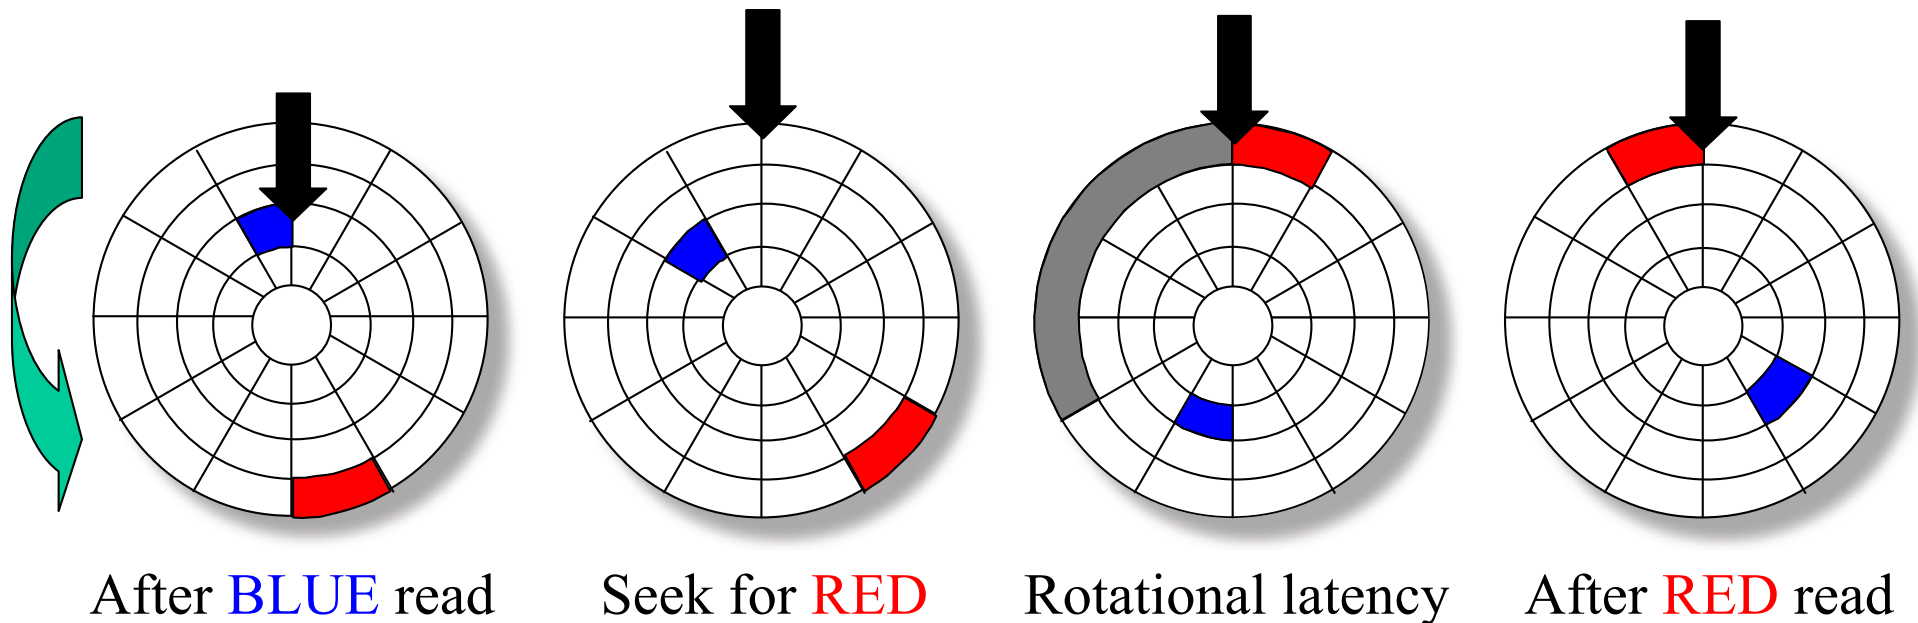


Source: R. Burns' slides on storage systems

Disk Access Time

- ▶ Disk access time:
 - ▶ **command processing time**: interpreting access command by disk controller
 - ▶ **seek time**: moving arms to position disk head on track
 - ▶ **rotational delay**: waiting for block to rotate under head
 - ▶ **transfer time**: actually moving data to/from disk surface
 - ▶ **access time = seek time + rotational delay + transfer time**
(command processing time is considered negligible)
- ▶ Response time for disk access = queuing delay + access time

Components of Disk Access Time



Source: R. Burns' slides on storage systems

Disk Access Time (cont.)

- ▶ Seek time
 - ▶ avg. seek time: 5-6 ms
- ▶ Rotational delay (or rotational latency)
 - ▶ Depends on rotation speed - measured in rotations per minute (RPM)
 - ▶ Average rotational delay = time for $\frac{1}{2}$ revolution
 - ▶ Example: For 10000 RPM, avg. rotational delay = $0.5 \left(\frac{60}{10000} \right) = 3 \text{ ms}$
- ▶ Transfer time
 - ▶ n = number of requested sectors on same track
 - ▶ transfer time = $n \times \frac{\text{time for one revolution}}{\text{number of sectors per track}}$
 - ▶ avg. sector transfer time: 100-200 μs
- ▶ Sequential vs random I/O

Where does the disk access time go?



Source: R. Burns' slides on storage systems

Solid-State Drive (SSD)


- ▶ Built with NAND flash memory without any mechanical/moving parts
- ▶ Faster random read: latency $\approx 20 - 100 \mu s$
- ▶ Higher data transfer rate
 - ▶ SATA interface: ≈ 500 MB/s
 - ▶ NVMEe PCIe interface: up to 3 GB/s
- ▶ Lower power consumption
- ▶ Write latency $\approx 100 \mu s$
- ▶ Disadvantages:
 - ▶ Update to a page requires erasure of multiple pages (≈ 5 ms) before overwriting page
 - ▶ Limited number of times a page can be erased ($\approx 10^5 - 10^6$)

Storage Manager Components

file & access methods manager
buffer pool manager
disk space manager

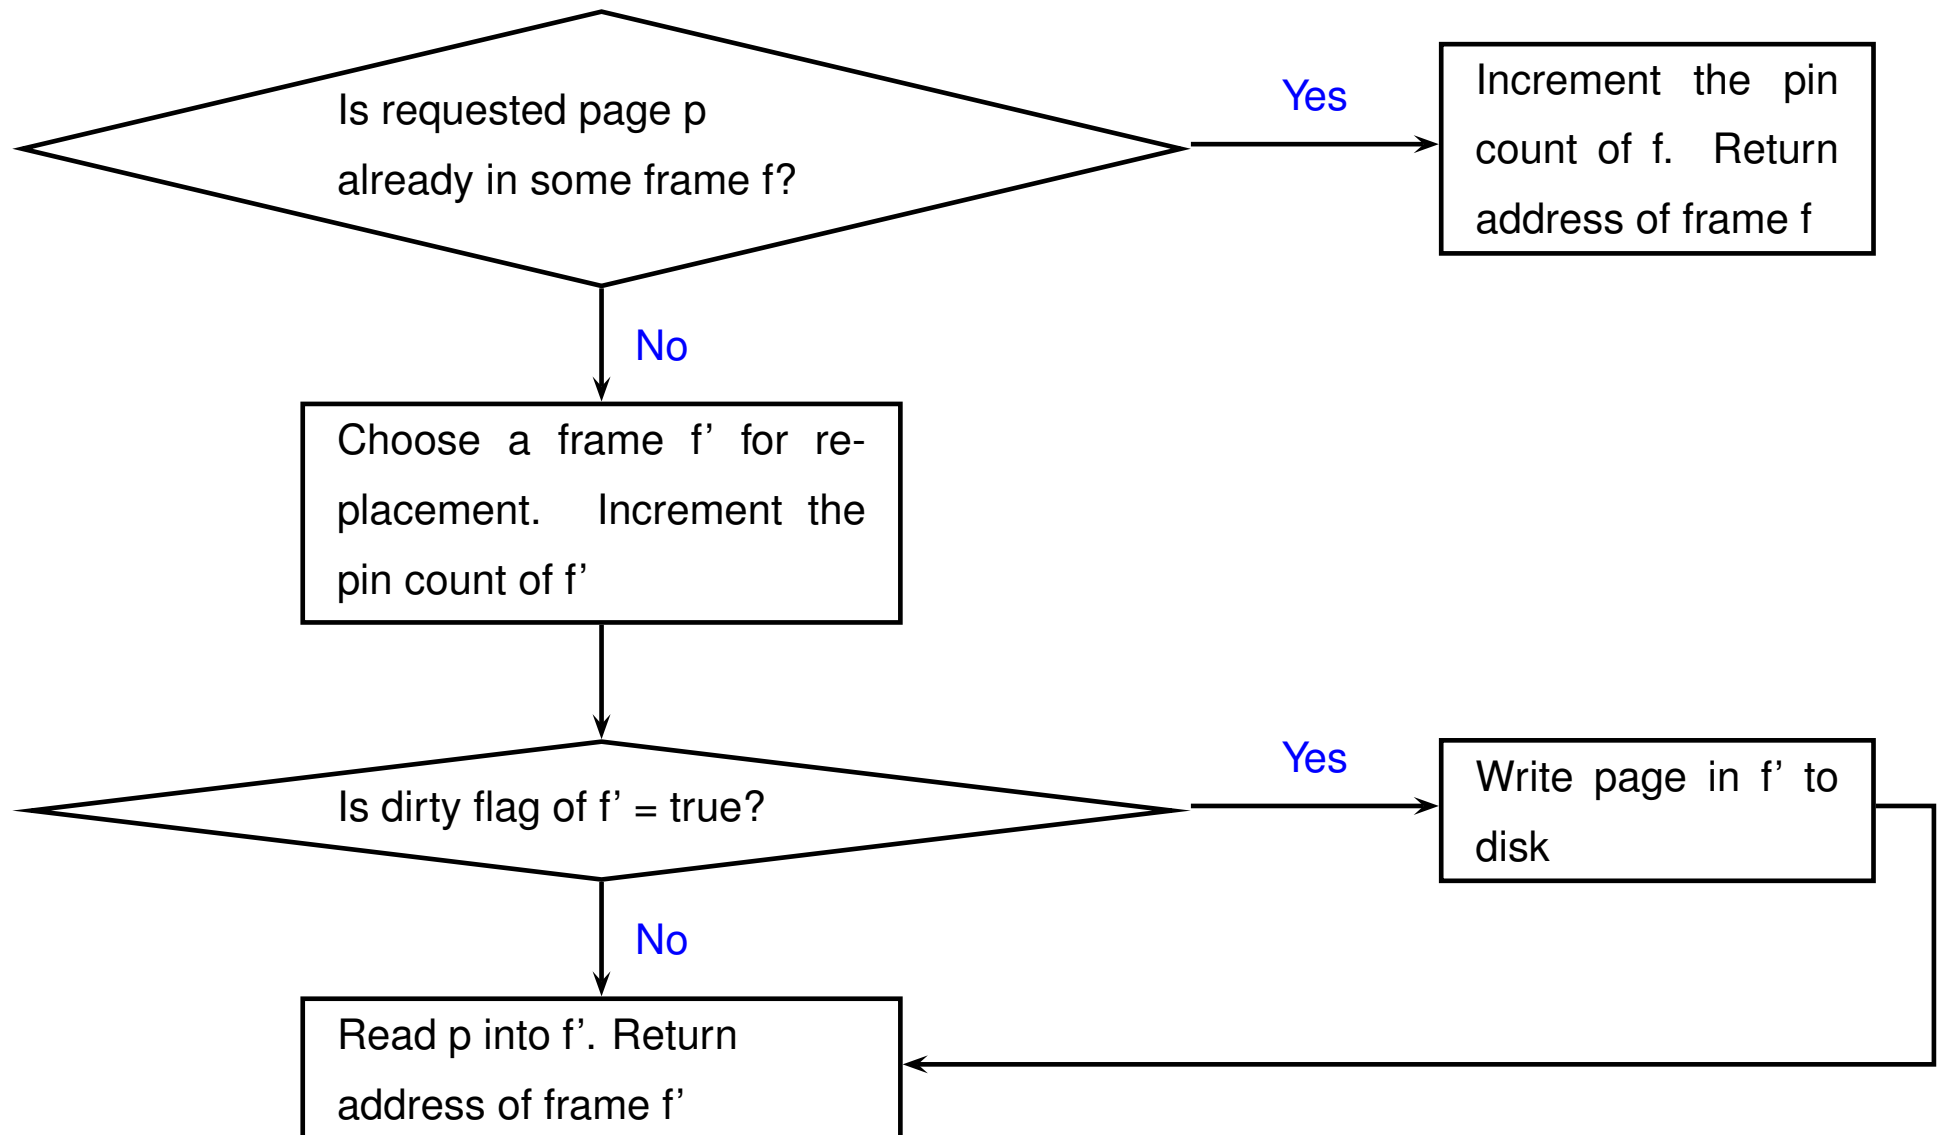
- ▶ Data is stored & retrieved in units called **disk blocks** (or **pages**)
 - Each block = sequence of one or more contiguous sectors
- ▶ **Files & access methods layer (aka file layer)** - deals with organization and retrieval of data
- ▶ **Buffer Manager** - controls reading/writing of disk pages
- ▶ **Disk Space Manager** - keeps track of pages used by file layer

Buffer Manager



- ▶ **Buffer pool** = main memory allocated for DBMS
- ▶ Buffer pool is partitioned into block-sized pages called **frames**
- ▶ Clients of buffer pool can
 - ▶ request for a disk page to be fetched into buffer pool
 - ▶ release a disk page in buffer pool
- ▶ A page in the buffer is **dirty** if it has been modified & not updated on disk
- ▶ Two variables are maintained for each frame in buffer pool:
 - ▶ **pin count** - number of clients using page (initialized to 0)
 - ▶ **dirty flag** - whether page is dirty (initialized to false)

Buffer Manager: Handling a request for page p



Buffer Manager (cont.)

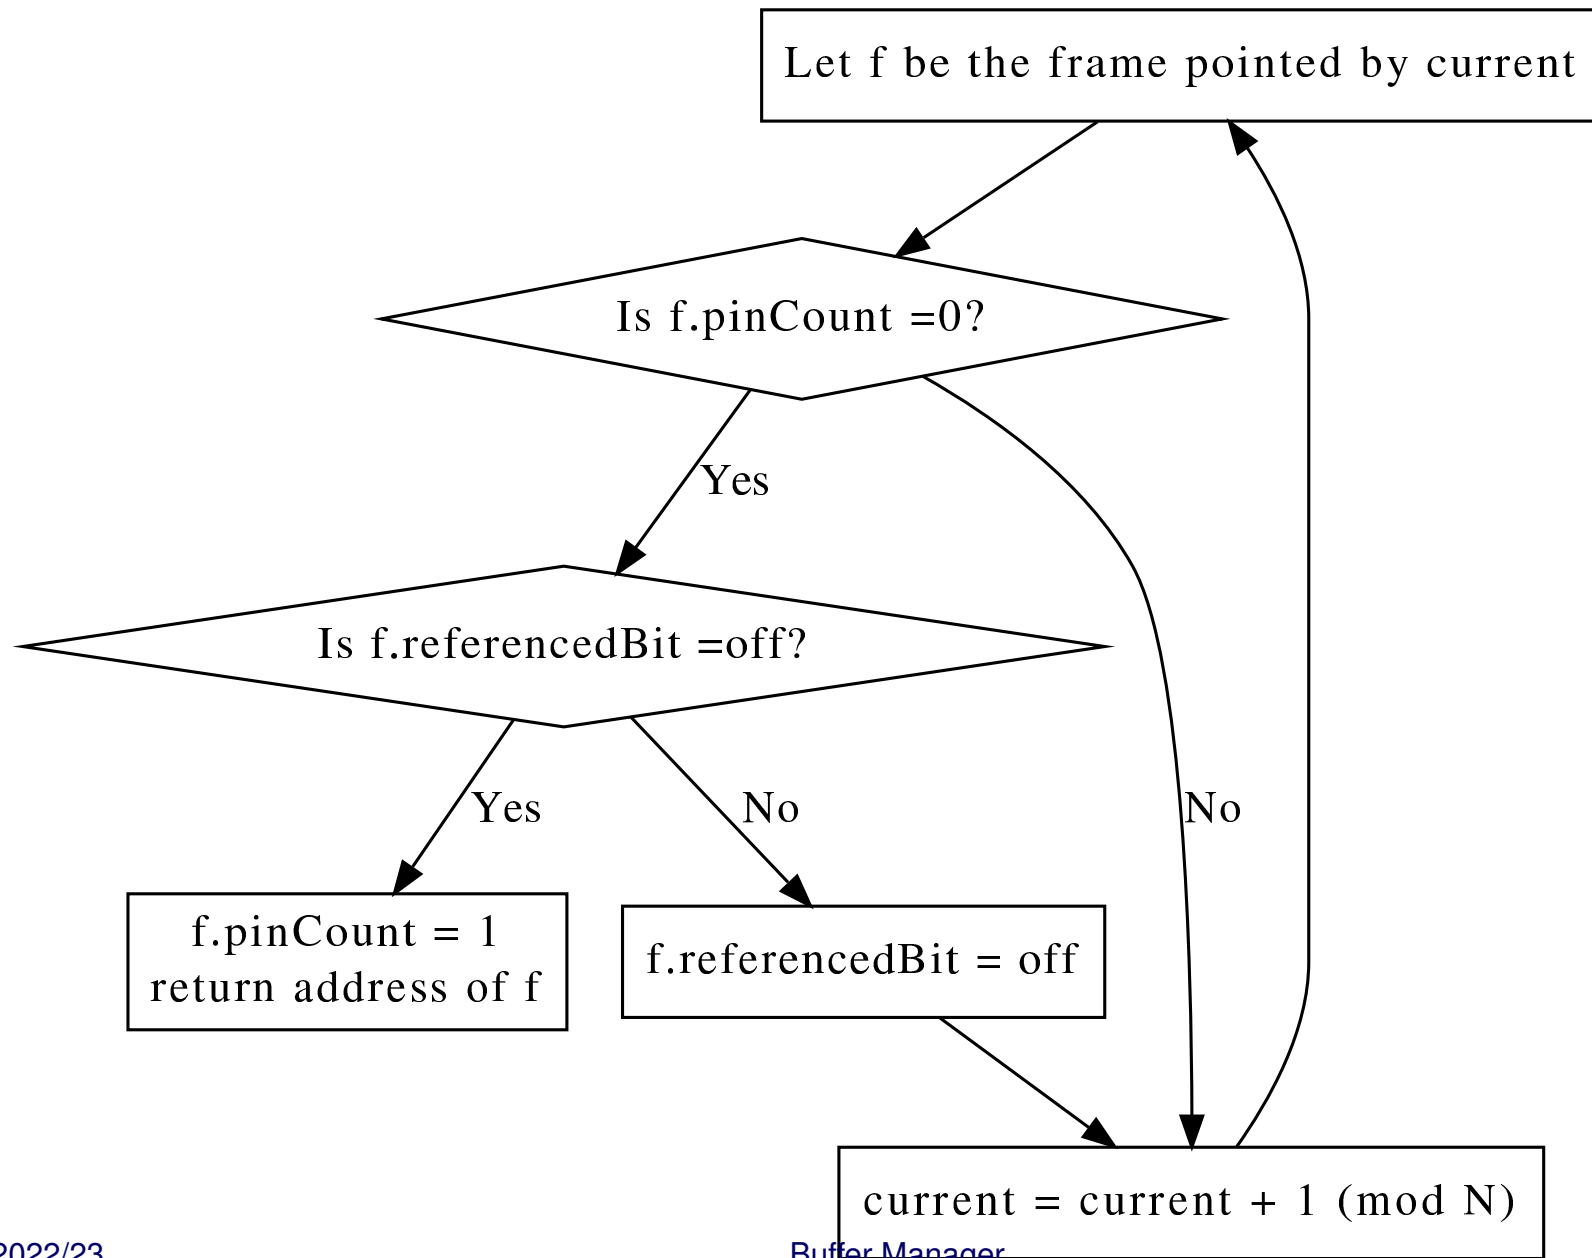
- ▶ Incrementing pin count is called **pinning** the requested page in its frame
- ▶ Decrementing the pin count is called **unpinning** the page
- ▶ When unpinning a page, its dirty flag should be updated to true if the page is dirty
- ▶ A page in buffer can be replaced only when its pin count is 0
- ▶ Before replacing a buffer page, it needs to be written back to disk if its dirty flag is true
- ▶ Buffer manager coordinates with transaction manager to ensure data correctness and recoverability

Buffer Manager: Replacement Policies

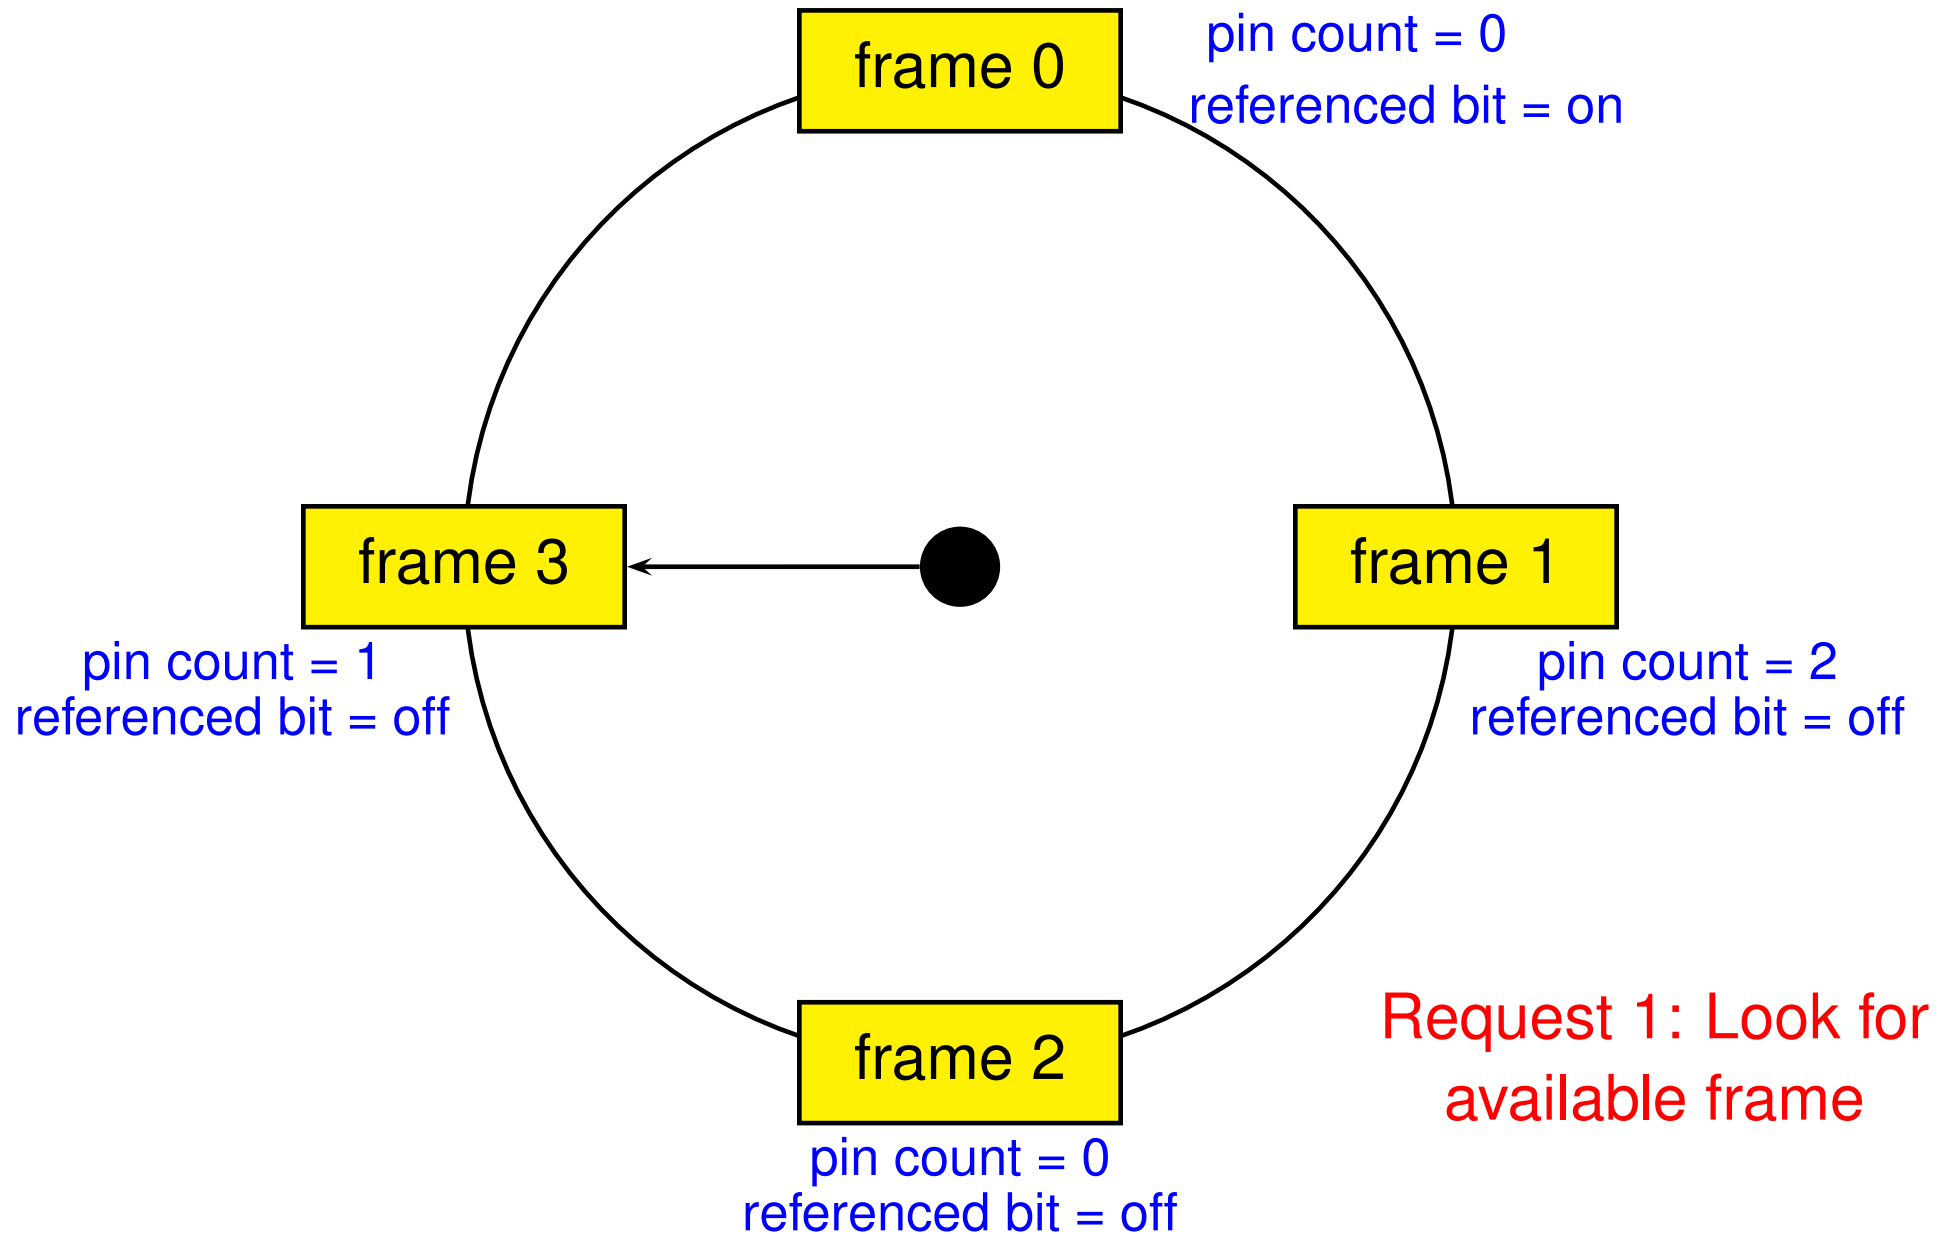
- ▶ **Replacement Policy**: decide which unpinned page to replace
 - ▶ Random
 - ▶ First In First Out (FIFO)
 - ▶ Most Recently Used (MRU)
 - ▶ Least Recently Used (LRU)
 - ★ Uses a queue of pointers to frames with pin count = 0
 - ▶ **Clock** - a variant of LRU
 - ★ **current** variable - points to some buffer frame
 - ★ Each frame has a **referenced bit** - turns on when its pin count becomes 0
 - ★ Replace a page that has referenced bit off & pin count = 0
 - ▶ etc.

Clock Replacement Policy

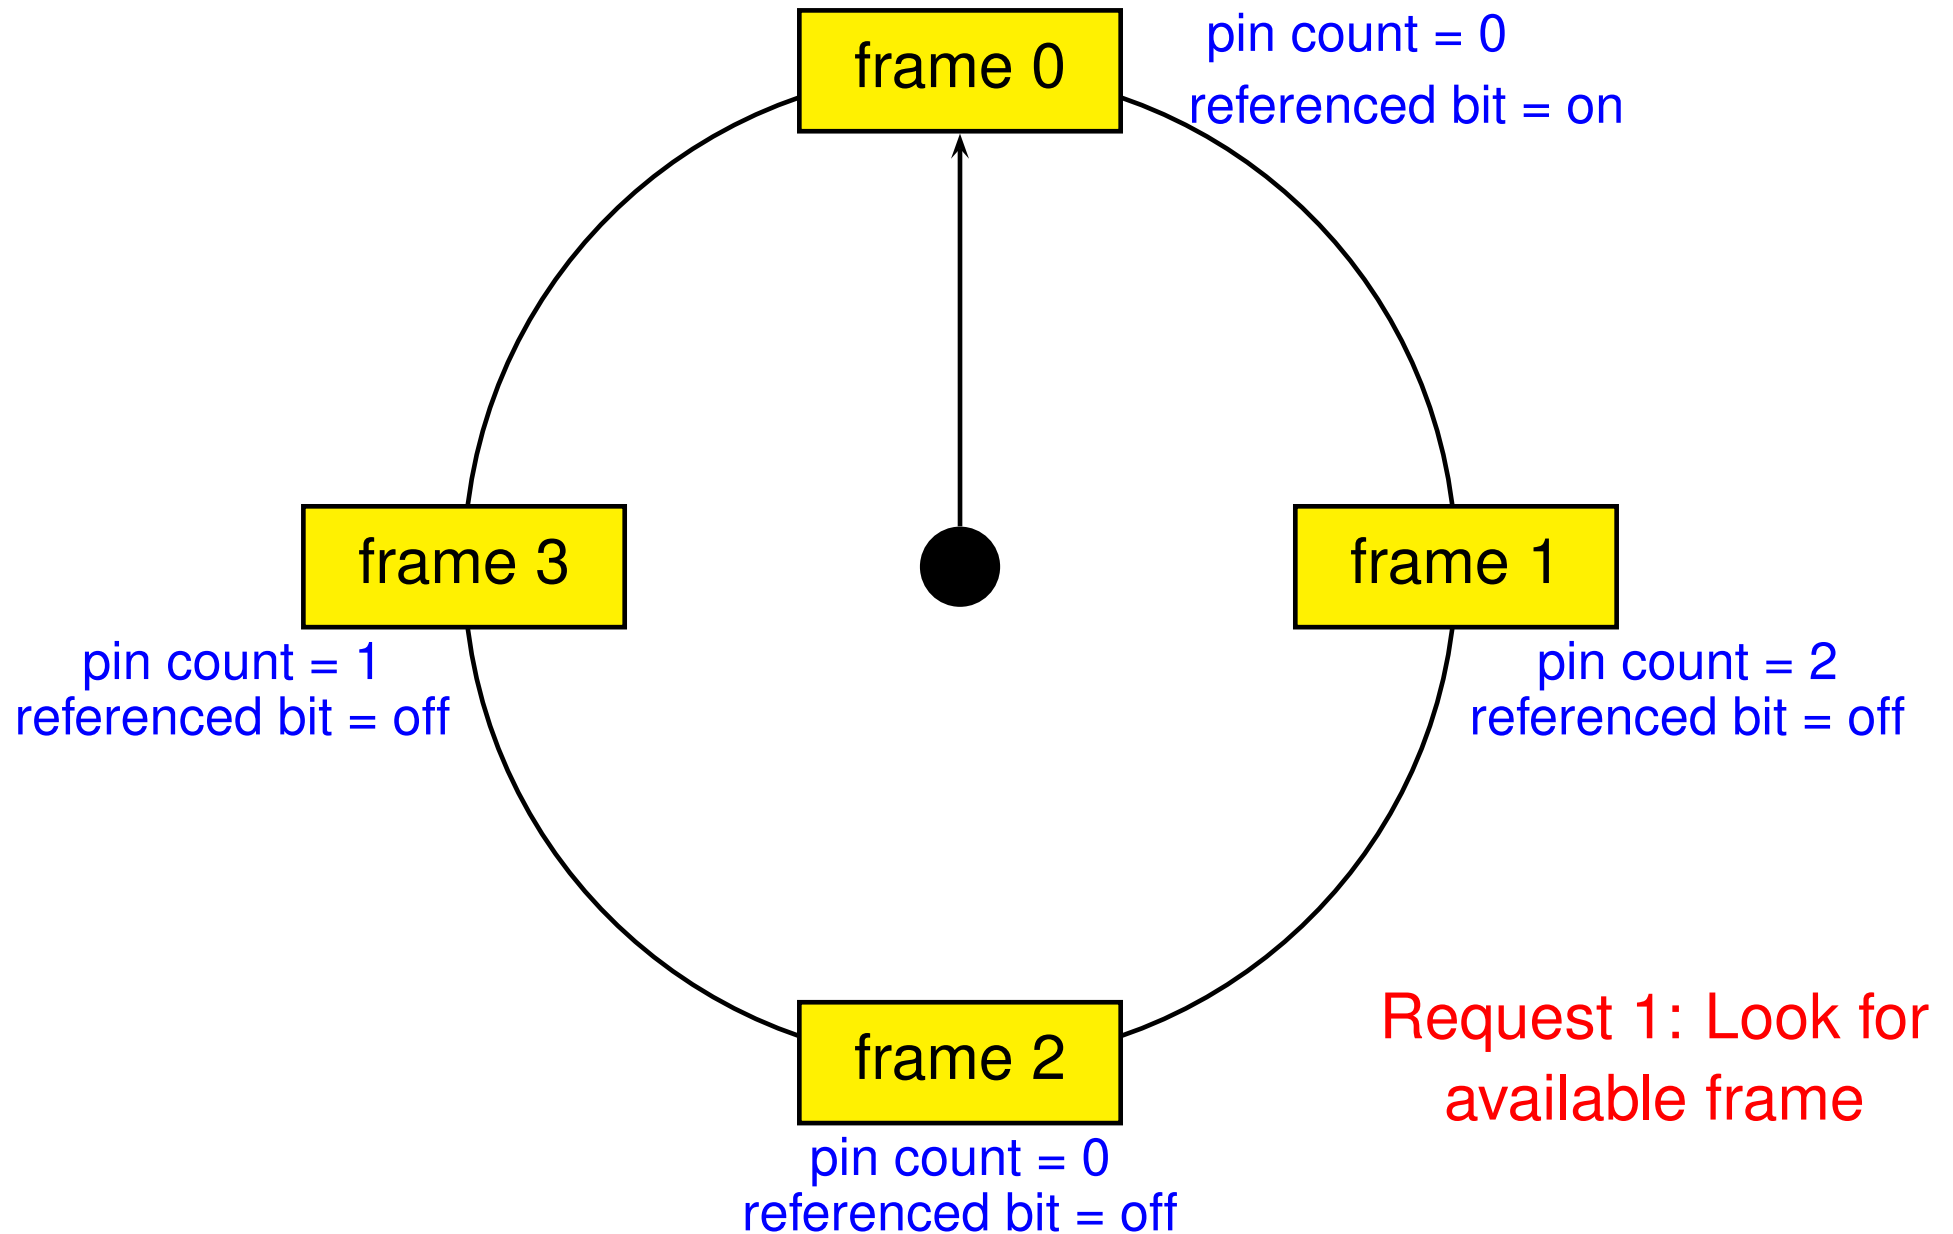
N = number of frames in buffer pool



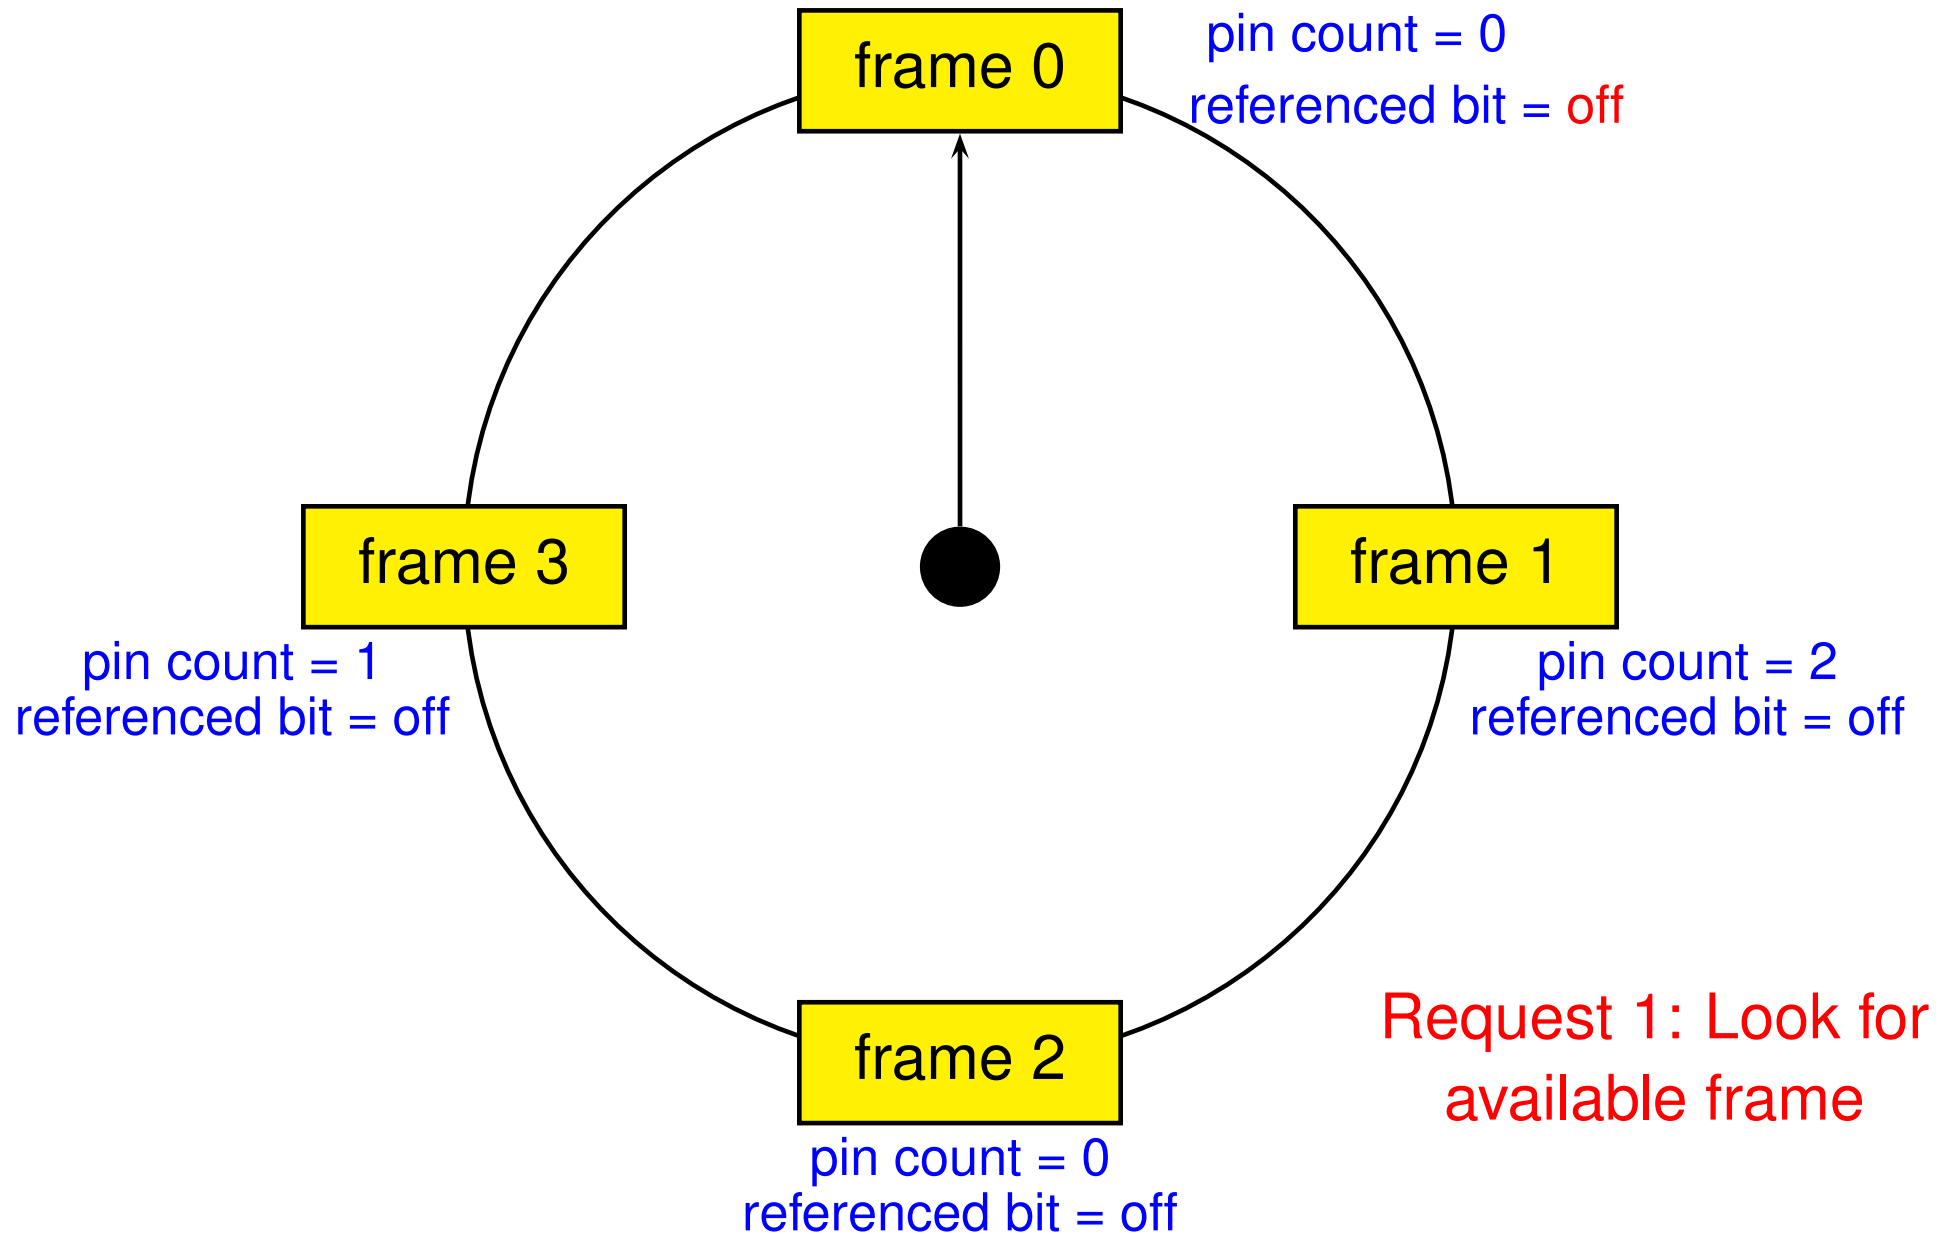
Clock Replacement Policy: Example



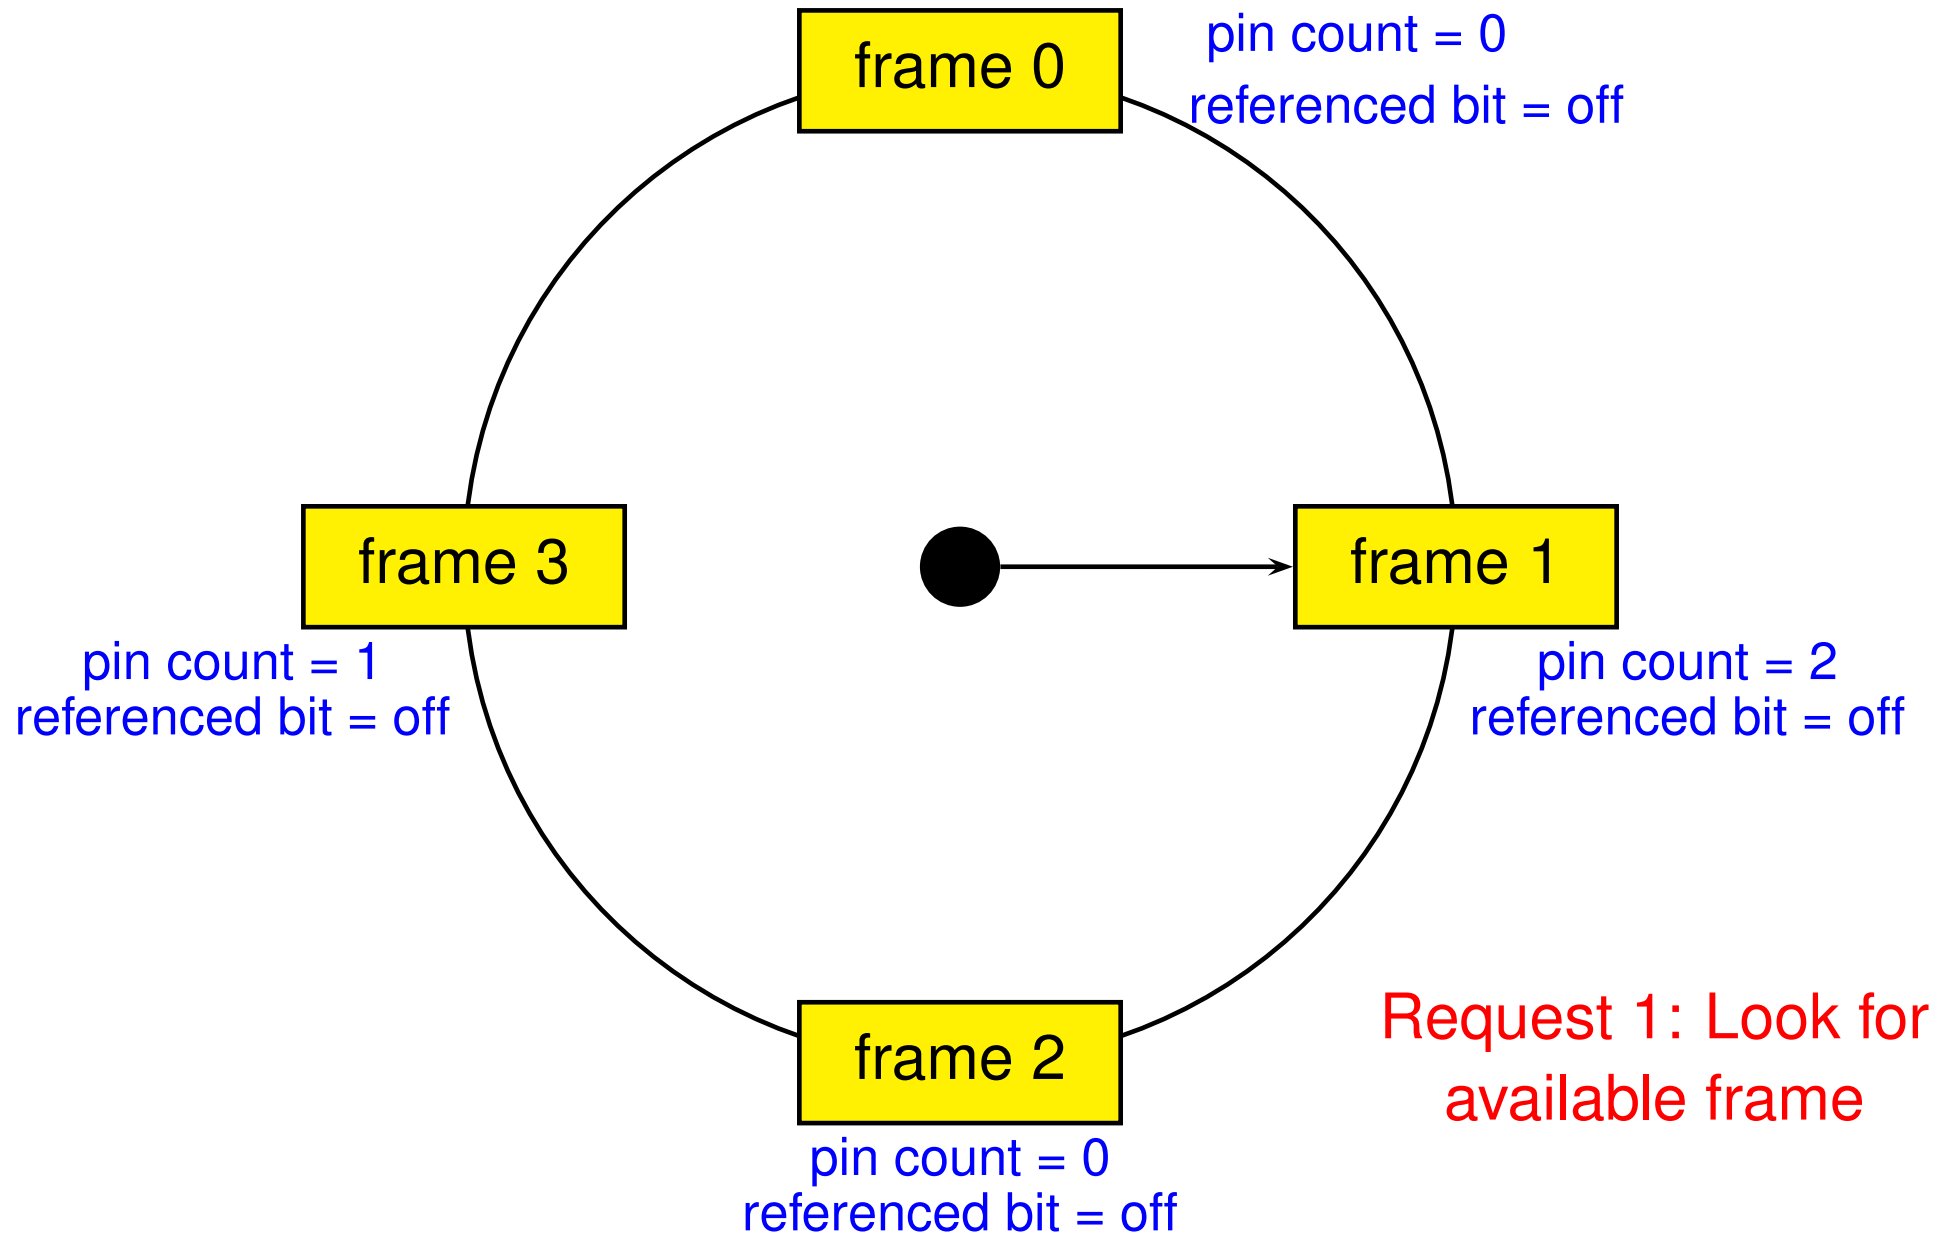
Clock Replacement Policy: Example



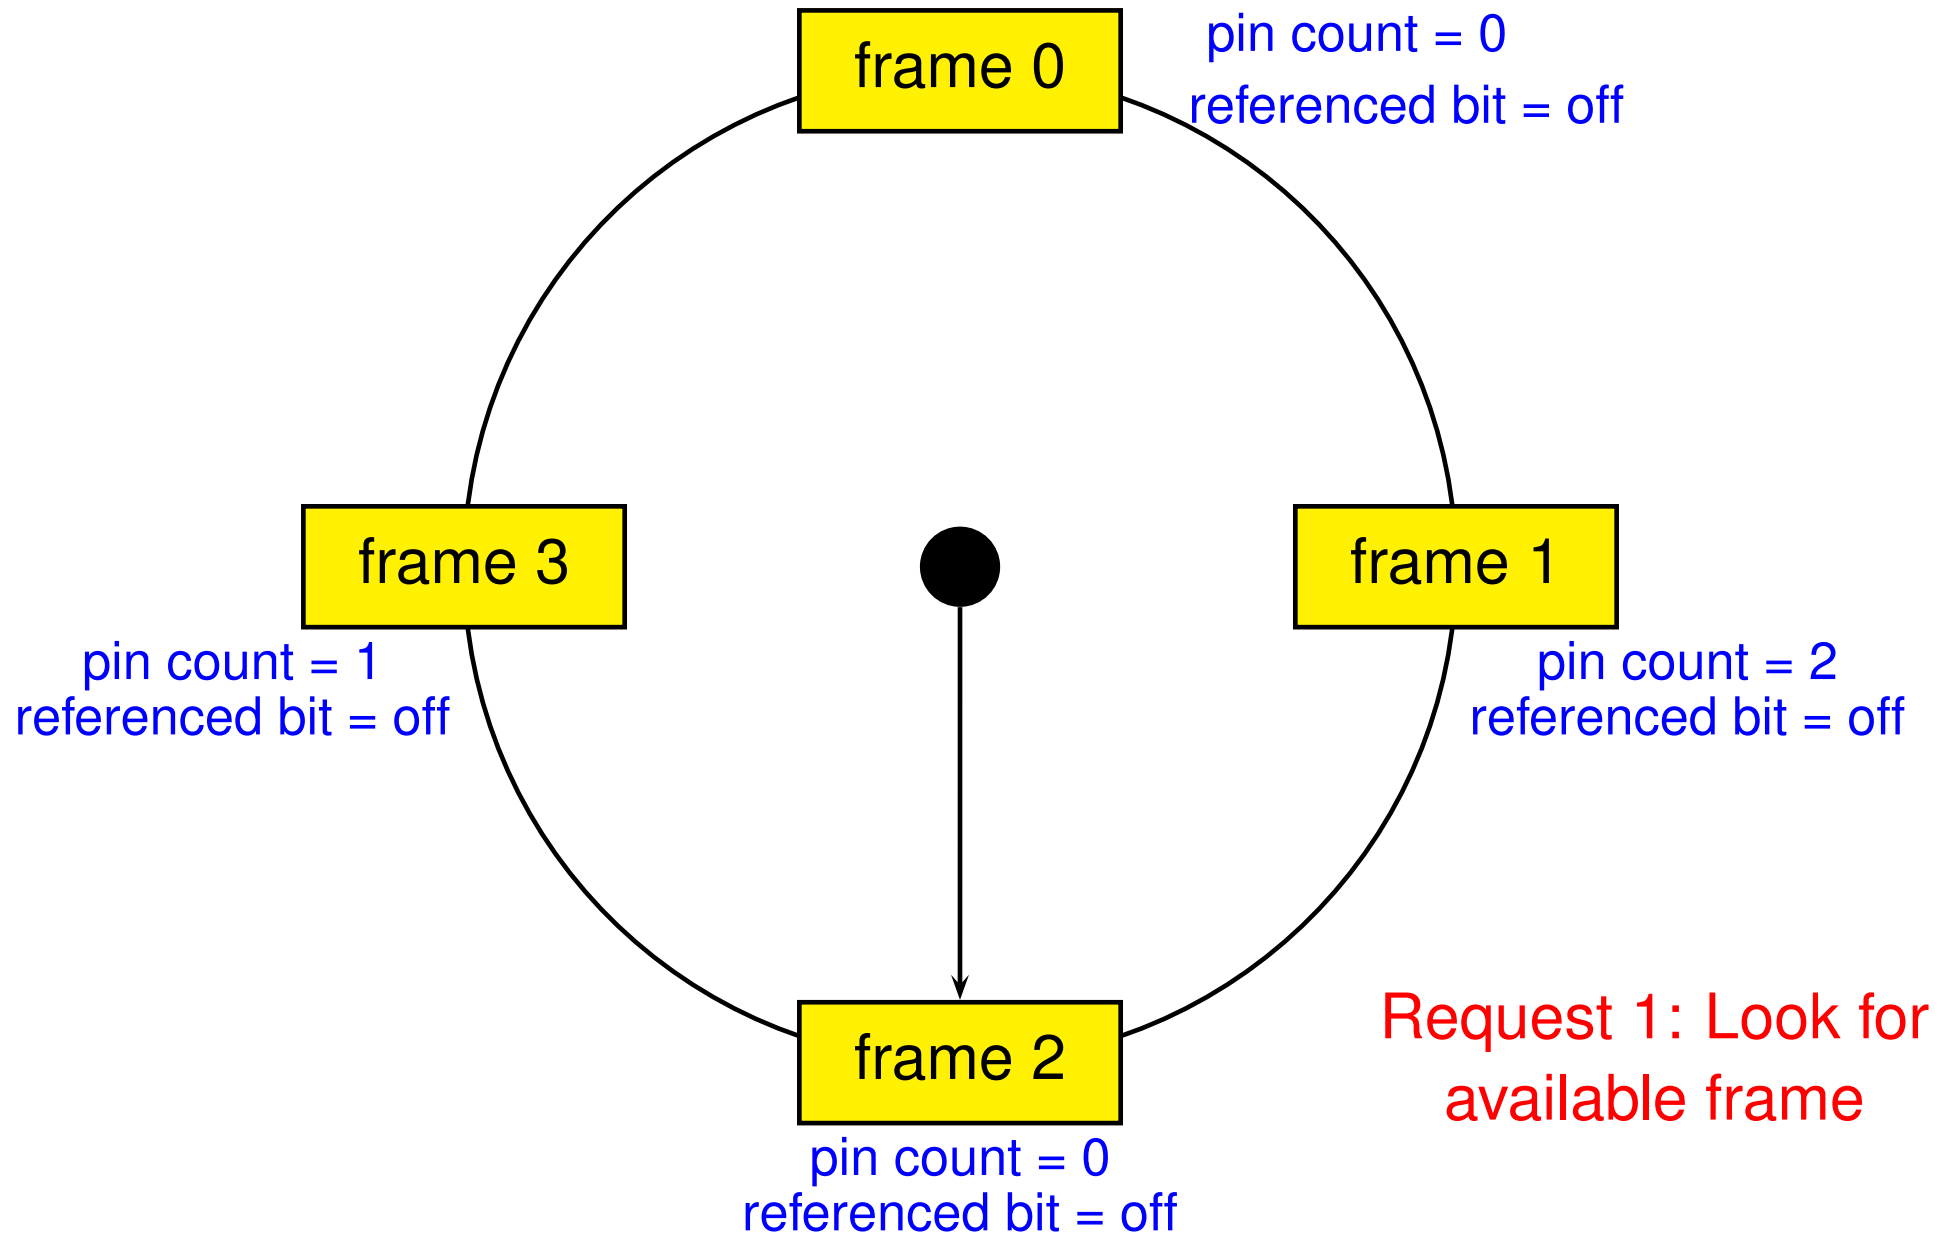
Clock Replacement Policy: Example




Clock Replacement Policy: Example




Clock Replacement Policy: Example




Clock Replacement Policy: Example



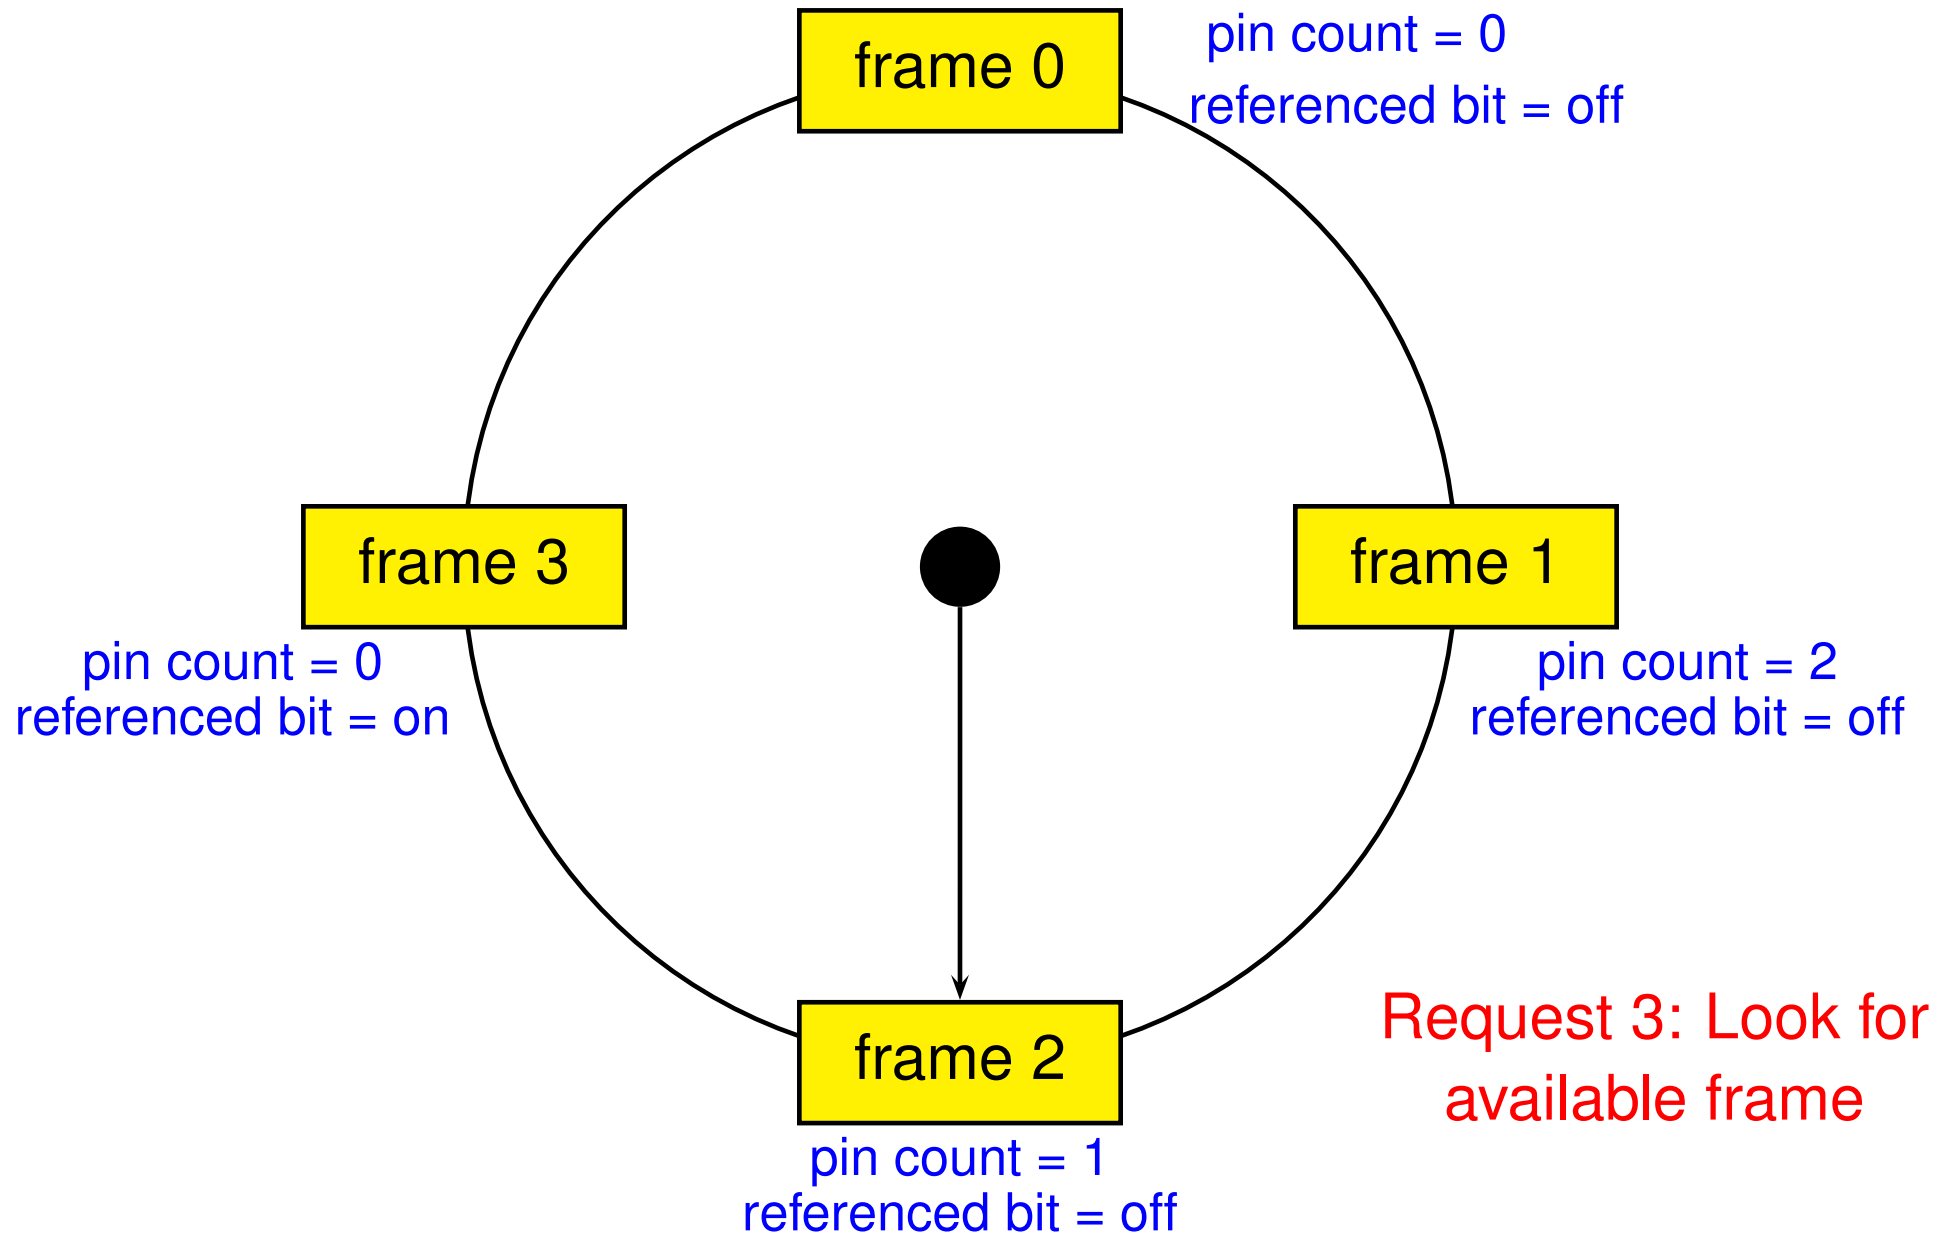
Clock Replacement Policy: Example



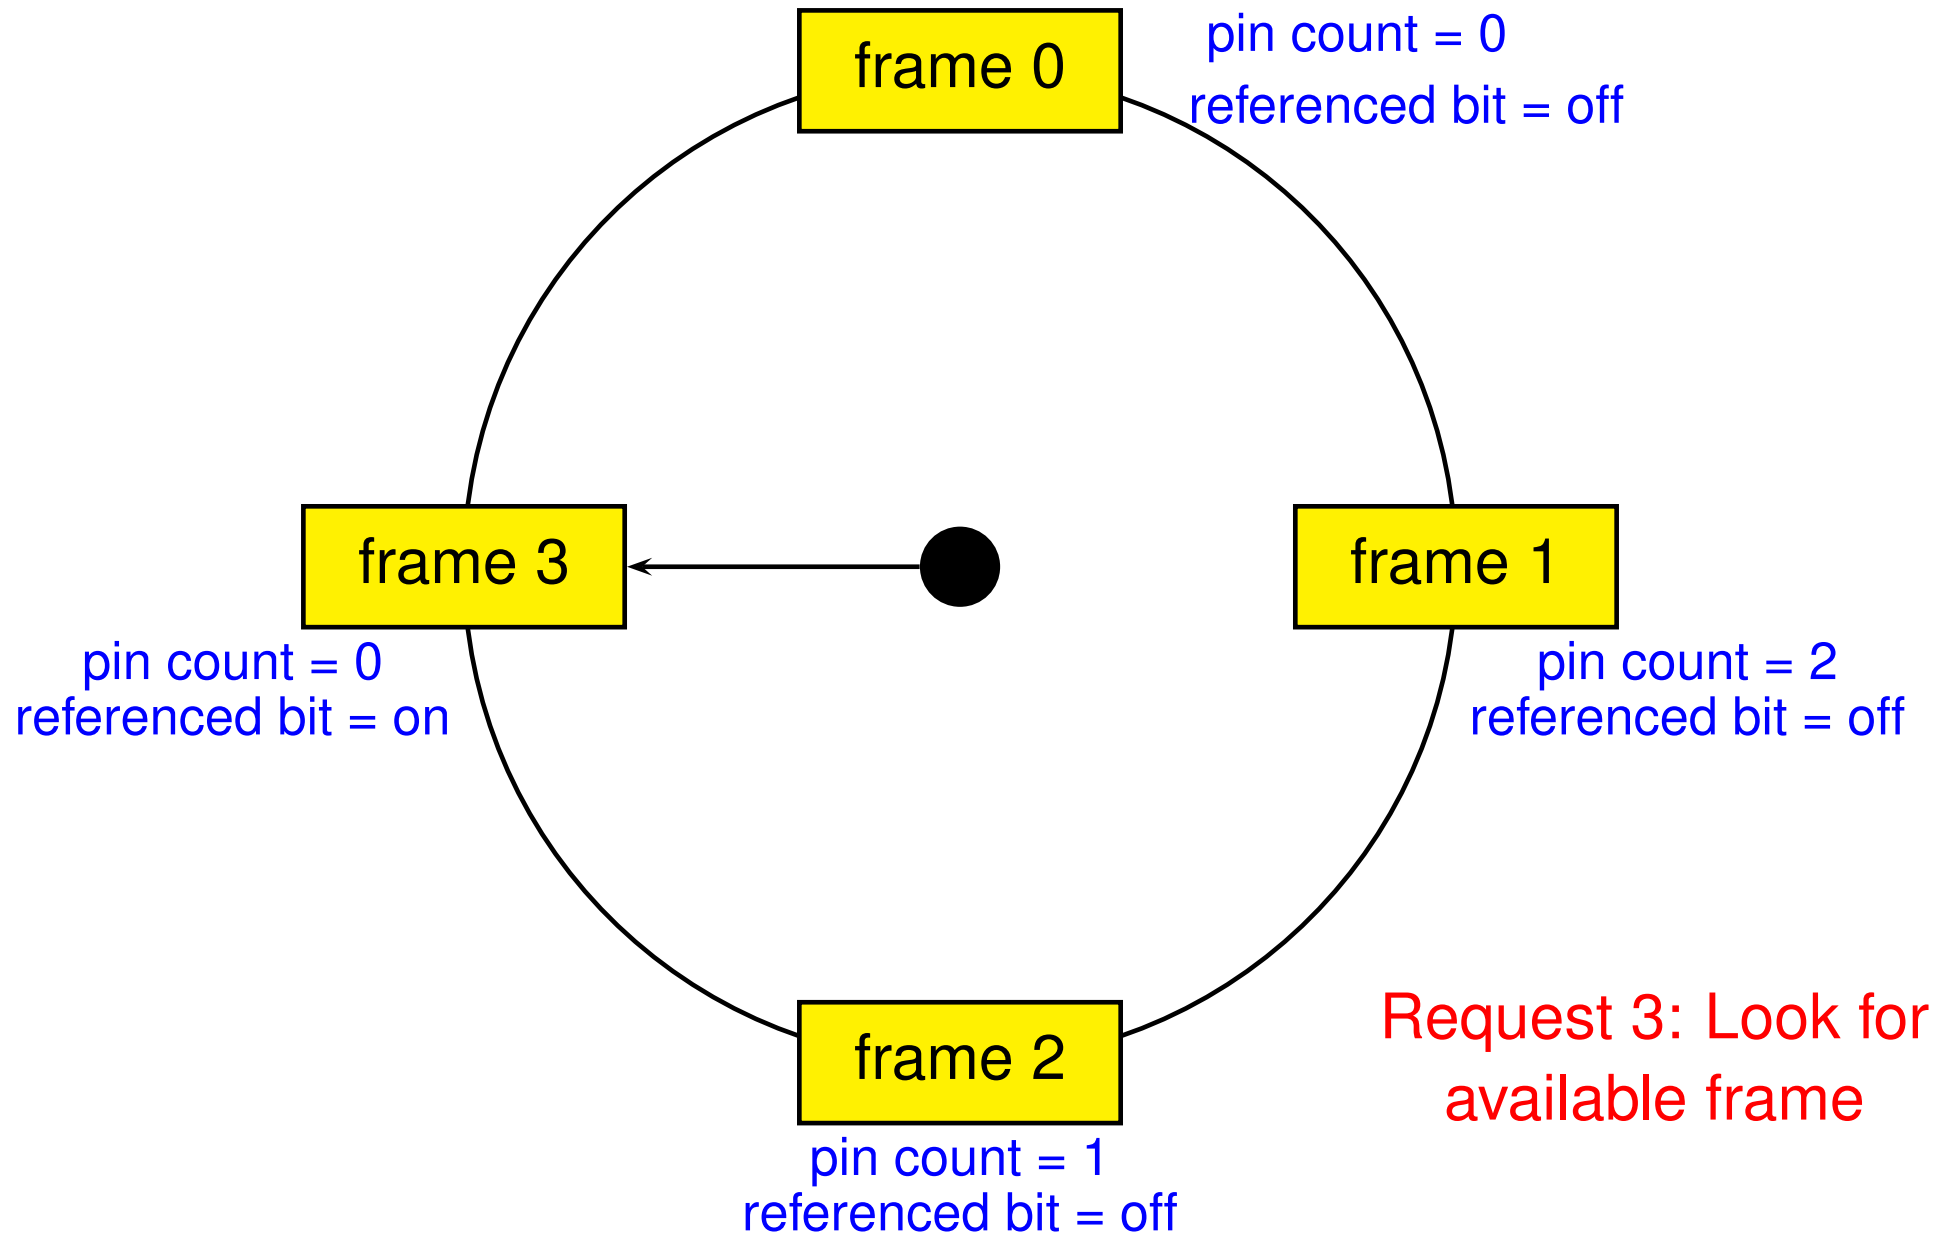
Clock Replacement Policy: Example



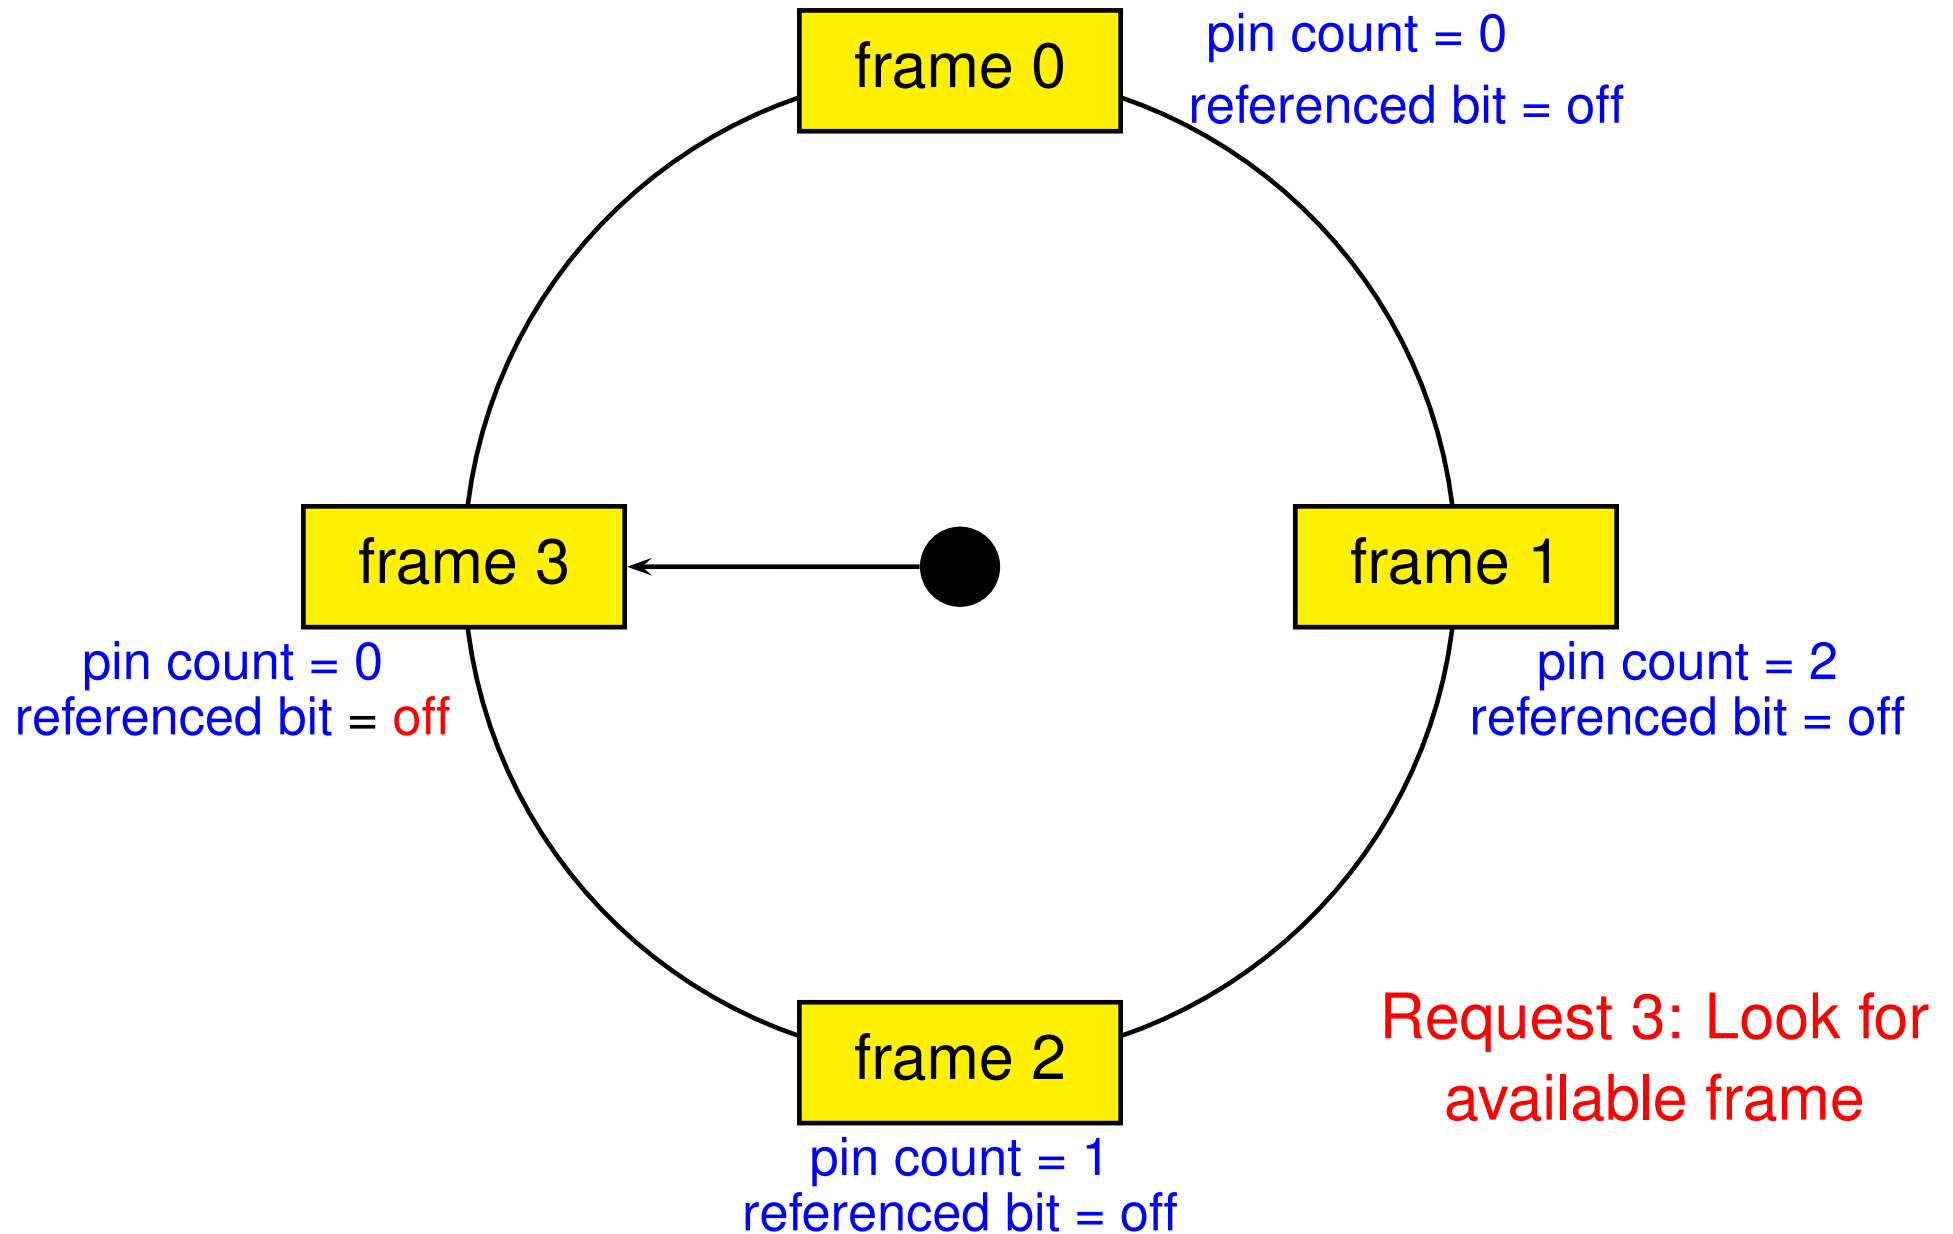
Clock Replacement Policy: Example




Clock Replacement Policy: Example




Clock Replacement Policy: Example



Clock Replacement Policy: Example



Clock Replacement Policy: Example



Files

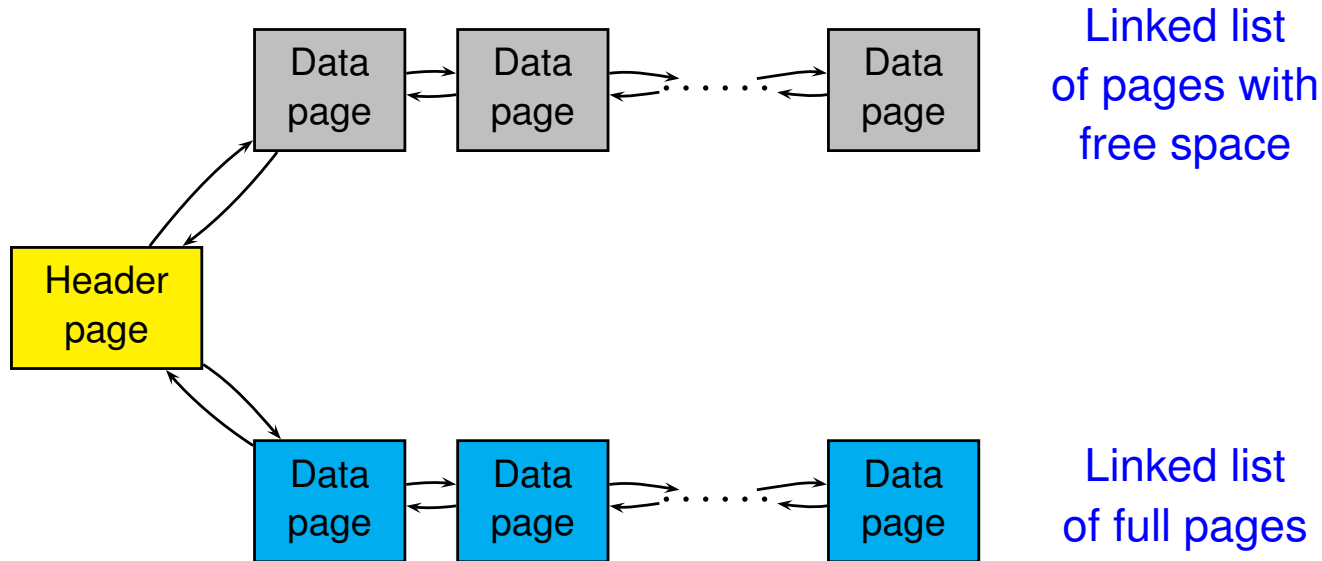
► File abstraction

- Each relation is a file of records
- Each record has a unique record identifier called **RID** (or **TID**)
- Common file operations:
 - ★ create a file
 - ★ delete a file
 - ★ insert a record
 - ★ delete a record with a given RID
 - ★ get a record with a given RID
 - ★ scan all records


► **File organization** = method of arranging data records in a file that is stored on disk

- **Heap file**: unordered file
- **Sorted file**: records are ordered on some search key
- **Hashed file**: records are located in blocks via a hash function

Heap File Implementations





Linked
List
Implementation




Page
Directory
Implementation

Page Formats


- ▶ How are records organized within a page?
- ▶ RID = (page id, slot number)
- ▶ Fixed-Length Records
 - ▶ **Packed organization**: Store records in contiguous slots
 - ▶ **Unpacked organization**: Uses a bit array to maintain free slots



Variable-Length Records: Slotted Page Organization



PostgreSQL's Slotted Page Organization



Source: B. Momjian's slides on PostgreSQL internals

Record Formats

- ▶ How to organize fields within a record?

- ▶ Fixed-Length Records

- ▶ Fields are stored consecutively

F1	F2	F3	F4
----	----	----	----

- ▶ Variable-Length Records

- ▶ Delimit fields with special symbols

F1	\$	F2	\$	F3	\$	F4
----	----	----	----	----	----	----

- ▶ Use an array of field offsets

o_1	o_2	o_3	o_4	F1	F2	F3	F4
-------	-------	-------	-------	----	----	----	----

Each o_i is an offset to beginning of field F_i