

CS3223 Lecture 3


Hash-based Indexing

Hashed-Based Indexing

- ▶ Used for equality queries but not for range queries
- ▶ **Hashing Techniques:**
 - ▶ Static hashing
 - ▶ Dynamic hashing
 - ★ Linear hashing
 - ★ Extendible hashing
 - ★ etc

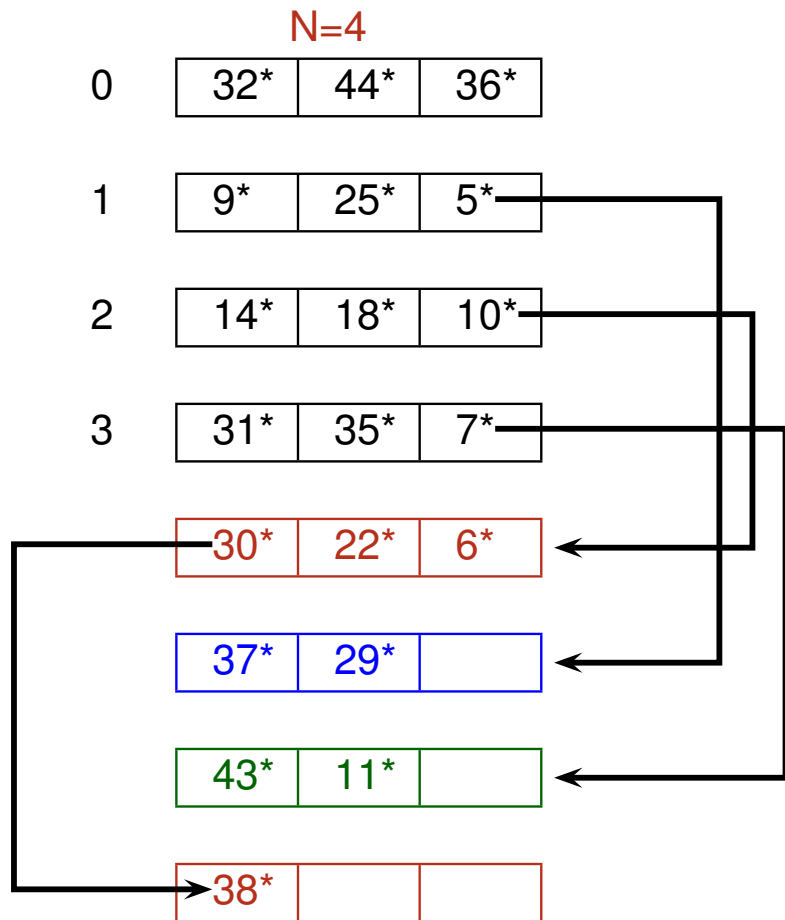
Static Hashing

- ▶ Data is stored in N buckets B_0, B_1, \dots, B_{N-1}
 - ▶ N is fixed at creation time
- ▶ Hashing function $h(.)$ is used to identify the bucket to store a record
 - ▶ A record with search key k is inserted into bucket B_i , where $i = h(k) \bmod N$
 - ▶ $h(k)$ maps the search key value into a bit string
- ▶ Each bucket consists of one **primary data page** & a chain of zero or more **overflow data pages**



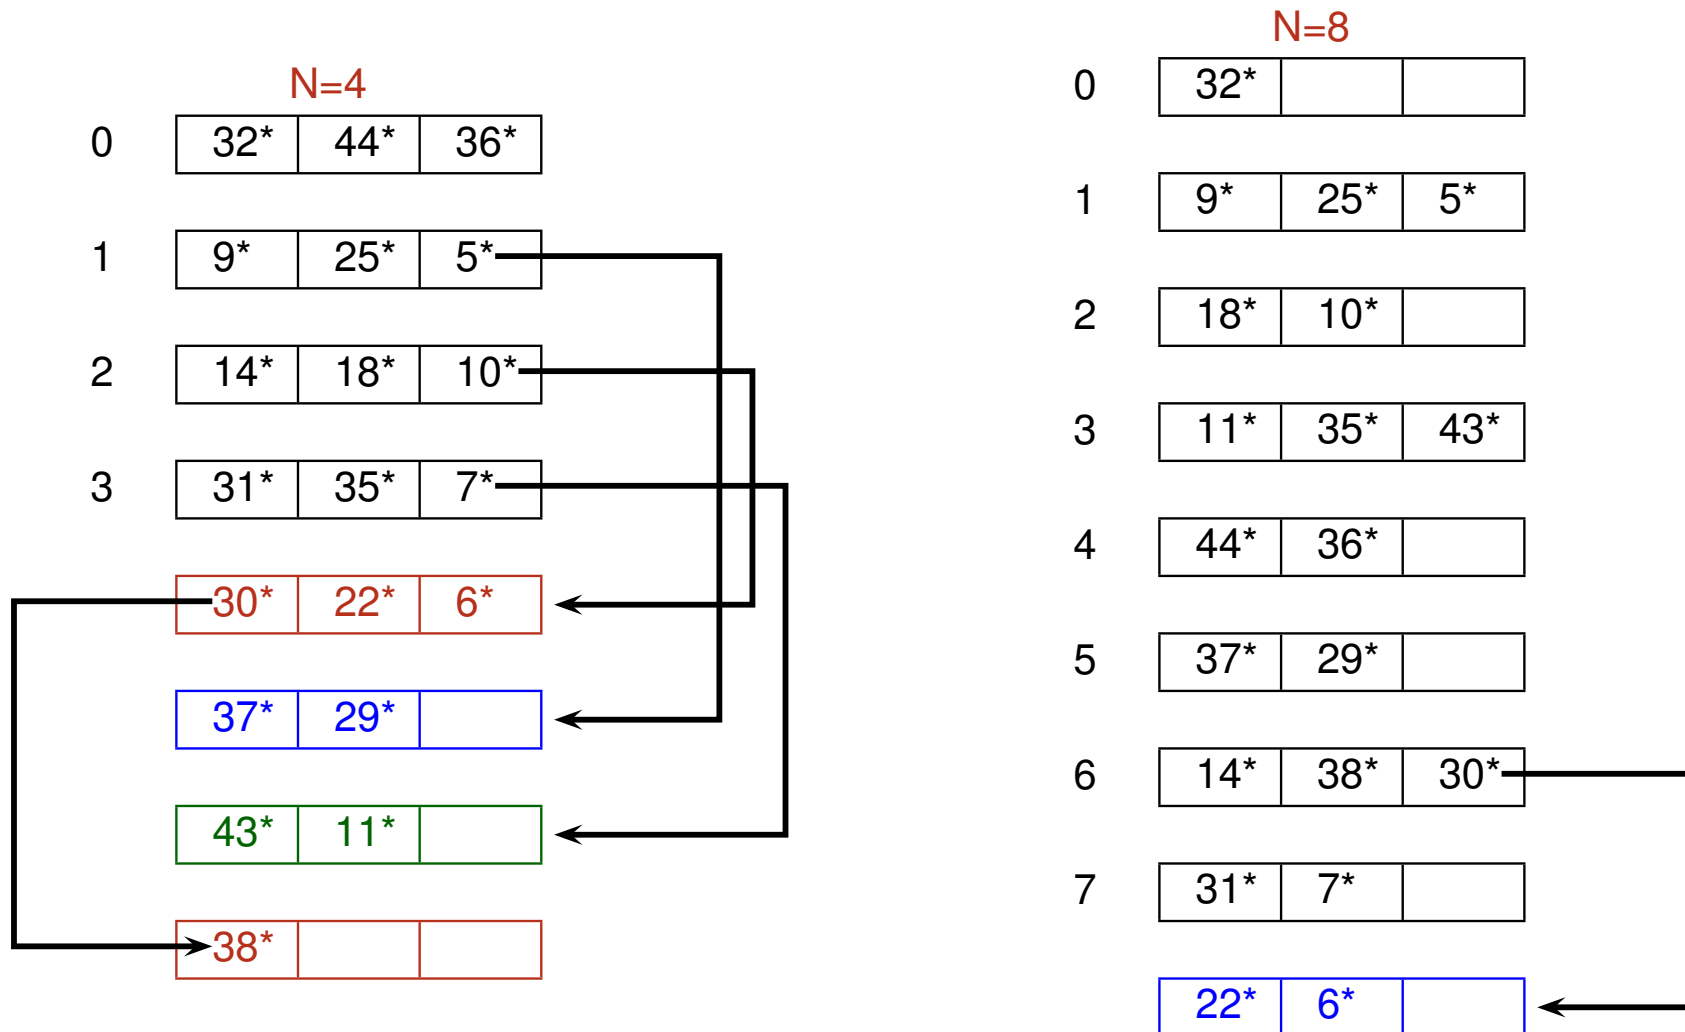
Static Hashing: Example

v^* denote a data entry e with $h(e.key) = v$
 $e.key$ denote the index's search key value of e



Static Hashing: Example

v^* denote a data entry e with $h(e.key) = v$
 $e.key$ denote the index's search key value of e



Linear Hashing

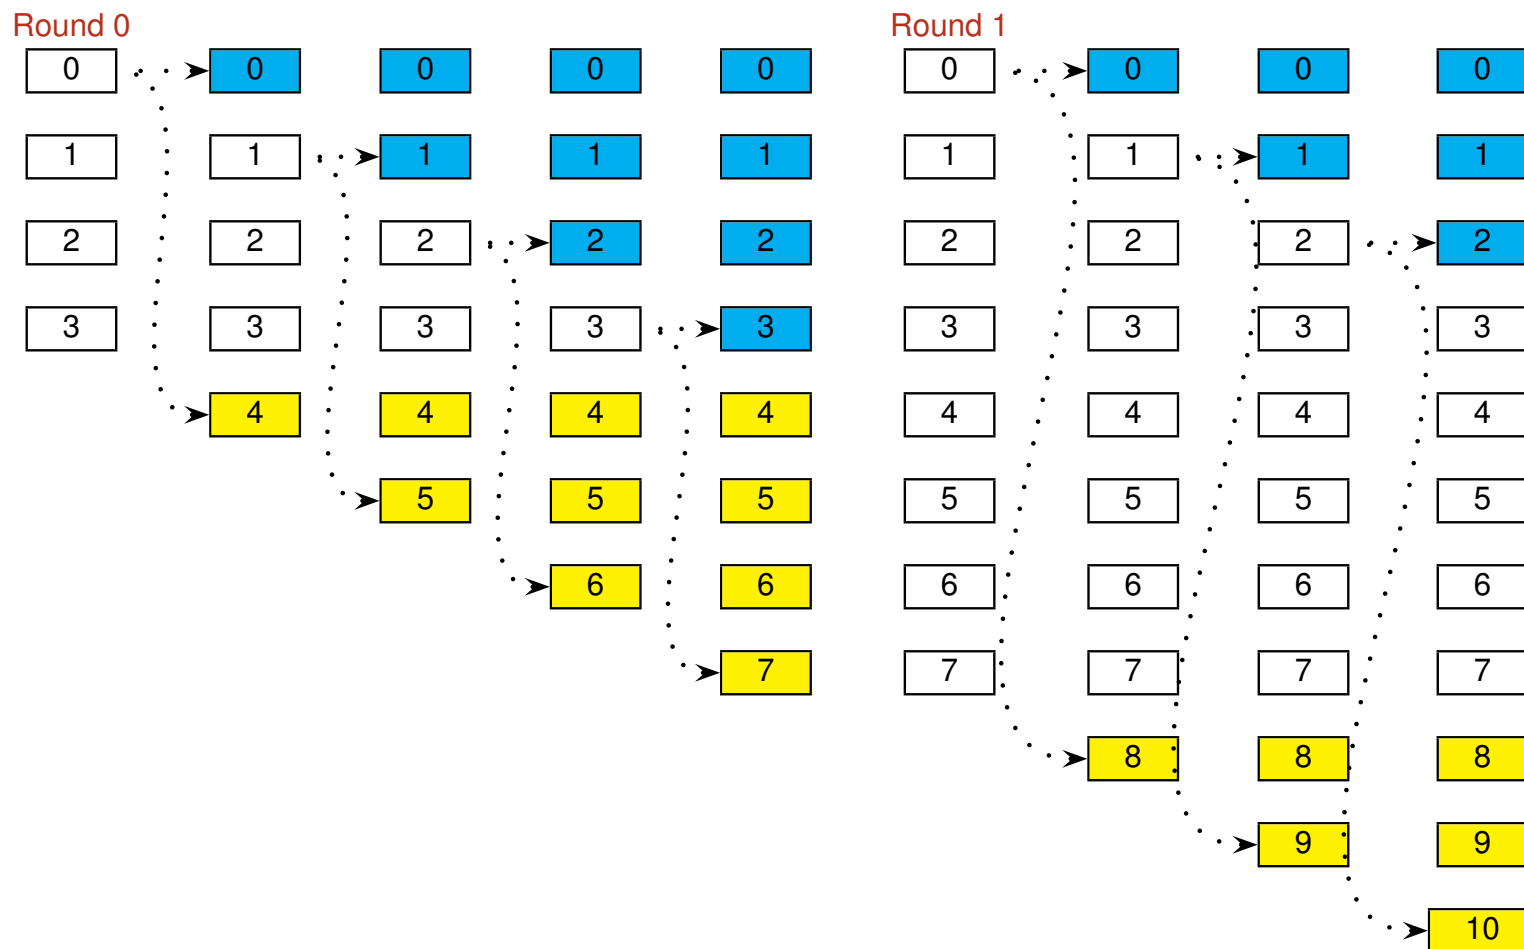
- ▶ Dynamic hashing technique
- ▶ Hash file grows linearly by systematic **splitting** of buckets
- ▶ Overflow pages are needed since an overflowed bucket might not be split immediately
- ▶ Each bucket consists of a primary data page & a chain of zero or more overflow pages
- ▶ An insertion into a bucket B_i overflows if all the pages in B_i (i.e., primary & overflow pages) are full

Linear Hashing (cont.)

- ▶ Assume initial file size of N_0 buckets
 - ▶ Buckets $B_0, B_1, \dots, B_{N_0-1}$
- ▶ File grows linearly by **splitting buckets** in rounds
 - ▶ At each round, buckets are split sequentially: B_i is split before B_{i+1}
- ▶ How to split a bucket B_i ?
 - ▶ Add a new bucket B_j (known as **split image** of B_i)
 - ▶ Redistribute entries in B_i between B_i & B_j
- ▶ File size increases by one bucket after each bucket split
- ▶ At the end of one round of splitting (i.e., every bucket at the start of the round has been split), file size is doubled
- ▶ Let N_i denote the file size at the beginning of round i ($i=0,1,\dots$)
 - ▶ $N_i = 2^i N_0$
- ▶ At the end of round i , N_i new buckets are added: $B_{N_i}, B_{N_i+1}, \dots, B_{2N_i-1}$
 - ▶ In round i , the **split image** of B_j is B_{N_i+j} for $j \in [0, N_i - 1]$

Splitting Buckets

- ▶ Assume $N_0 = 4$
- ▶ In round 0, split image of bucket B_i is B_{4+i} , for $i \in [0, 3]$
- ▶ In round 1, $N_1 = 8$ & split image of bucket B_i is B_{8+i} , for $i \in [0, 7]$



Dynamic Hashing

0	32*	44*	36*	
1	9*	5*		
2	14*	18*	30*	
3	31*	35*		

Before B_0 was split in round 0

0	32*			
1	9*	5*		
2	14*	18*	30*	
3	31*	35*		
4	44*	36*		

After B_0 was split in round 0

- ▶ Recall that v^* denote a data entry e with $h(e.key) = v$ and $e.key$ denote the index's search key value of e
- ▶ Each round uses two hash functions: functions h_i and h_{i+1} for round i
 - ▶ $h_i(k) = h(k) \bmod N_i$
 - ▶ B_x is the bucket for search key k if B_x had not been split, where $x = h_i(k)$
 - ▶ B_y is the bucket for search key k if B_x had been split, where $y = h_{i+1}(k)$
- ▶ Keep track of which bucket to be split next using variable **next**

Redistributing Entries

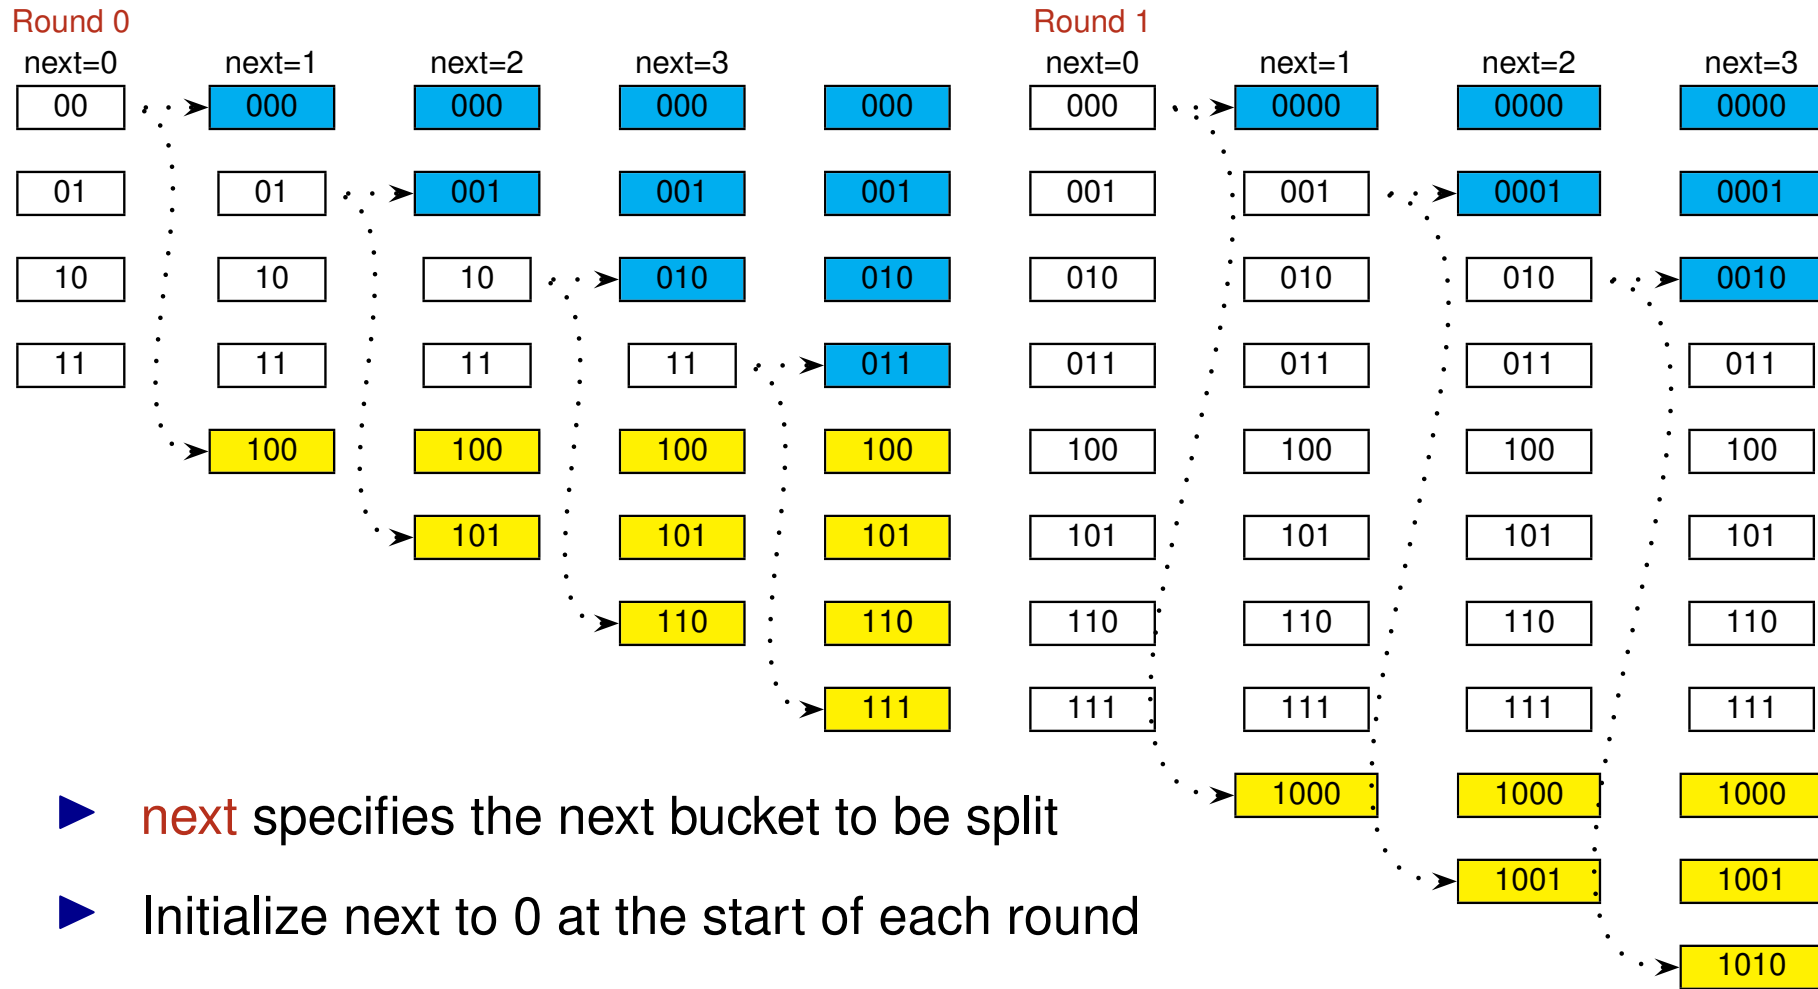
		$h(e.key)$	
0	32* 44* 36*	32 (100000 ₂)	0 32*
		44 (101100 ₂)	
1	9* 5*	36 (100100 ₂)	1 9* 5*
		9 (001001 ₂)	
2	14* 18* 30*	5 (000101 ₂)	2 14* 18* 30*
		14 (001110 ₂)	
3	31* 35*	18 (010010 ₂)	3 31* 35*
		30 (011110 ₂)	
		31 (011111 ₂)	
		35 (100011 ₂)	4 44* 36*

Before B_0 was split in round 0


After B_0 was split in round 0

- ▶ For convenience, $N_0 = 2^m$ for $m \in \{0, 1, \dots\}$
- ▶ $N_i = 2^i N_0 = 2^{m+i}$
- ▶ $h_i(k) = h(k) \bmod N_i = \text{value of last } (m+i) \text{ bits of } h(k)$
- ▶ Consider the splitting of bucket B_j in round i
 - ▶ Before splitting, all entries e in B_j have $h_i(e.key) = j$ (i.e., same last $(m+i)$ bits)
 - ▶ After splitting, e remains in B_j iff the last $(m+i+1)$ bit of $h(e.key)$ is 0

Splitting Buckets




- ▶ **next** specifies the next bucket to be split
- ▶ Initialize next to 0 at the start of each round
- ▶ Buckets (in round i) can be classified into three regions:



When to split a bucket?

- ▶ The time to split the next bucket can be decided with various criteria:
 - ▶ Split whenever some bucket overflows
 - ▶ Split whenever space utilization of file is above some threshold
 - ▶ etc.
- ▶ Overflow pages are needed since an overflowed bucket might not be split immediately
- ▶ We shall assume that a bucket split is triggered whenever some **bucket overflows**
 - ▶ Bucket B_j overflows if an entry is to be inserted into B_j and all the pages in B_j (i.e., primary & overflow pages) are full
- ▶ We use **level** to denote the splitting round number

Inserting data entry with search key k



- **GetBucketNum(k)** returns bucket # where entry with search key k is located

$$GetBucketNum(k) = \begin{cases} h_{level}(k) & \text{if } h_{level}(k) \geq next, \\ h_{level+1}(k) & \text{otherwise.} \end{cases}$$

- **SplitBucket()** splits bucket B_{next}

1. Redistribute the entries in B_{next} into $B_{next+N_{level}}$ using $h_{level+1}()$
2. $next = next + 1$
3. if $(next = N_{level})$ then { level = level + 1; next = 0 }

Linear Hashing: Inserting 43^* (101011)

level = 0, $N_0 = 4$, next=0

00	32*	44*	36*	
----	-----	-----	-----	--

01	9*	25*	5*	
----	----	-----	----	--

10	14*	18*	10*	30*
----	-----	-----	-----	-----

11	31*	35*	7*	11*
----	-----	-----	----	-----

Linear Hashing: Inserting 43^* (101011)

level = 0, $N_0 = 4$, next=0

00	32*	44*	36*	
01	9*	25*	5*	
10	14*	18*	10*	30*
11	31*	35*	7*	11*

level = 0, $N_0 = 4$, next=1

000	32*			
01	9*	25*	5*	
10	14*	18*	10*	30*
11	31*	35*	7*	11*
100	44*	36*		
	43*			

- ▶ Bucket B_{11} overflows
- ▶ Split bucket B_{00} ($32 = 100000$, $36 = 100100$, $44 = 101100$)
- ▶ Increment next to 1
- ▶ Insert 43^* into overflow page

Linear Hashing: Inserting 37^* (100101)

level = 0, $N_0 = 4$, next=1

000

32^*			
--------	--	--	--

01

9^*	25^*	5^*	
-------	--------	-------	--

10

14^*	18^*	10^*	30^*
--------	--------	--------	--------

11

31^*	35^*	7^*	11^*
--------	--------	-------	--------



100

44^*	36^*		
--------	--------	--	--

43^*			
--------	--	--	--



Linear Hashing: Inserting 37^* (100101)



- ▶ No overflow, no splitting
- ▶ Insert 37^* into bucket B_{01}

Linear Hashing: Inserting 29^* (011101)

level = 0, $N_0 = 4$, next=1

000

32^*			
--------	--	--	--

01

9^*	25^*	5^*	37^*
-------	--------	-------	--------

10

14^*	18^*	10^*	30^*
--------	--------	--------	--------

11

31^*	35^*	7^*	11^*
--------	--------	-------	--------

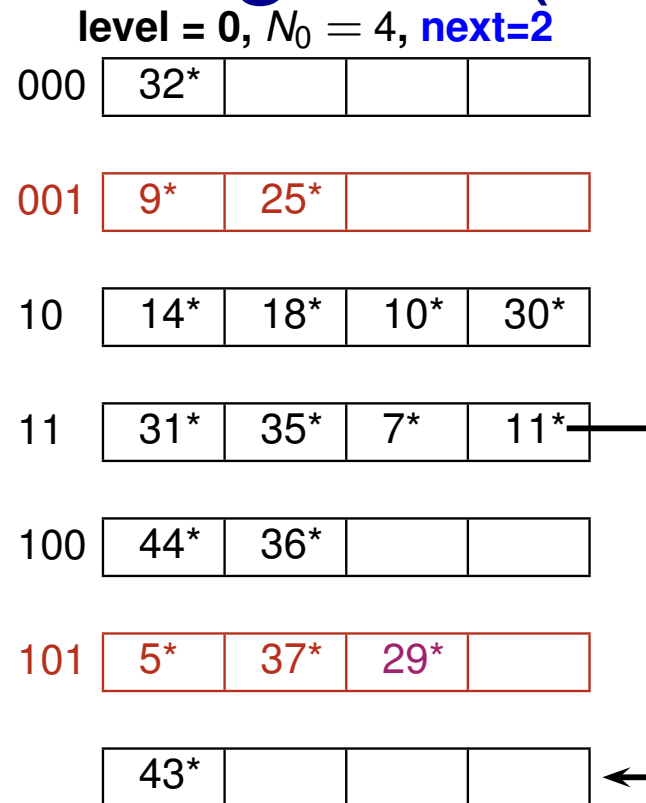

100

44^*	36^*		
--------	--------	--	--

43^*			
--------	--	--	--



Linear Hashing: Inserting 29^* (011101)



- ▶ Bucket B_{01} overflows
- ▶ Split bucket B_{01} ($5 = 101$, $9 = 1001$, $25 = 11001$, $37 = 100101$)
- ▶ Increment next to 2
- ▶ Insert 29^* into bucket B_{101}

Linear Hashing: Inserting 22^* (010110)

level = 0, $N_0 = 4$, next=2

000

32^*			
--------	--	--	--

001

9^*	25^*		
-------	--------	--	--

10

14^*	18^*	10^*	30^*
--------	--------	--------	--------

11

31^*	35^*	7^*	11^*
--------	--------	-------	--------

100

44^*	36^*		
--------	--------	--	--

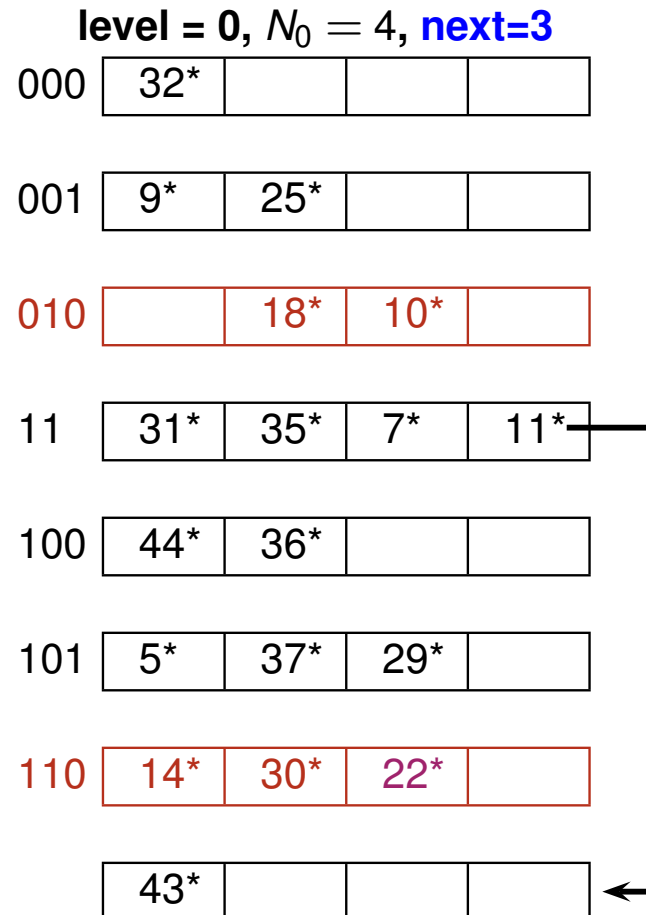

101

5^*	37^*	29^*	
-------	--------	--------	--

43^*			
--------	--	--	--




Linear Hashing: Inserting 22^* (010110)





- ▶ Bucket B_{10} overflows
- ▶ Split bucket B_{10} ($10 = 1010$, $14 = 1110$, $18 = 10010$, $30 = 11110$)
- ▶ Increment next to 3
- ▶ Insert 22^* into bucket B_{110}

Linear Hashing: Inserting 6* (000110)



Linear Hashing: Inserting 6^* (000110)



- ▶ No overflow, no splitting
- ▶ Insert 6^* into bucket B_{110}

Linear Hashing: Inserting 38* (100110)

level = 0, $N_0 = 4$, next=3

000

32*			
-----	--	--	--

001

9*	25*		
----	-----	--	--

010

	18*	10*	
--	-----	-----	--

11

31*	35*	7*	11*
-----	-----	----	-----

100

44*	36*		
-----	-----	--	--

101

5*	37*	29*	
----	-----	-----	--

110


14*	30*	22*	6*
-----	-----	-----	----

43*			
-----	--	--	--




Linear Hashing: Inserting 38* (100110)

level = 0, $N_0 = 4$, next=3



level = 1, $N_0 = 4$, next=0



- ▶ Bucket B_{110} overflows
- ▶ Split bucket B_{11} (7 = 111, 11 = 1011, 31 = 11111, 35 = 100011, 43 = 101011)
- ▶ Increment level to 1; reset next to 0
- ▶ Insert 38* into overflow page

Linear Hashing: Inserting 62^* (111110)

level = 1, $N_0 = 4$, next=0

000

32^*			
--------	--	--	--

001

9^*	25^*		
-------	--------	--	--

010

	18^*	10^*	
--	--------	--------	--

011

43^*	35^*		11^*
--------	--------	--	--------

100

44^*	36^*		
--------	--------	--	--

101

5^*	37^*	29^*	
-------	--------	--------	--

110

14^*	30^*	22^*	6^*
--------	--------	--------	-------

111


31^*	7^*		
--------	-------	--	--

38^*			
--------	--	--	--




Linear Hashing: Inserting 62^* (111110)

level = 1, $N_0 = 4$, next=0



- ▶ No overflow, no splitting
- ▶ Insert 62^* into bucket B_{110}

level = 1, $N_0 = 4$, next=0




Linear Hashing: Deletion

- ▶ Locate bucket and delete entry
- ▶ If the last bucket $B_{N_{level}+next-1}$ becomes empty, it can be removed
- ▶ **Case 1:** If $next > 0$
 - ▶ Decrement $next$ by one

level = 1, $N_0 = 2$, next = 2

000	
001	
10	
11	
100	
101	empty



level = 1, $N_0 = 2$, next = 1


000	
01	
10	
11	
100	

Linear Hashing: Deletion (cont.)

- ▶ **Case 2:** If ($next = 0$) and ($level > 0$)
 - ▶ Update `next` to point to the last bucket in previous level
 $B_{N_{level}-1}$
 - ▶ Decrement `level` by one

level = 1, $N_0 = 2$, next = 0

00	
01	
10	
11	empty



level = 0, $N_0 = 2$, next = 1

00	
1	
10	

Linear Hashing: Performance

- ▶ One disk I/O unless the bucket has overflow pages
 - ▶ On average 1.2 disk I/O for uniform data distribution
 - ▶ Worst case: I/O cost is linear in the number of data entries
- ▶ Poor space utilization with skewed data distribution

Linear Hashing: Summary

- ▶ File grows dynamically by splitting buckets linearly in rounds
 - ▶ Each bucket split adds one new bucket to file
- ▶ The bucket B_j for a search key k is determined using two hash functions (dependent on round)
 - ▶ Assume initial file has 2^m buckets & current round number is level
 - ▶ $j = \text{last } (m + \text{level})$ bits of $h(k)$ if B_j has not yet been split in current round; otherwise, $j = \text{last } (m + \text{level} + 1)$ bits of $h(k)$
- ▶ In this course, we assume a bucket split is triggered whenever some bucket overflows due to an insertion
- ▶ Overflow pages are required for an overflowed bucket if it is not the next bucket to be split

Another Idea for Dynamic Hashing

- ▶ Similar to Linear Hashing
 - ▶ Number of buckets grow dynamically
 - ▶ Uses some number of the least significant bits of $h(k)$ to determine the bucket address for search key k
- ▶ But adds a new bucket (as split image) whenever an existing bucket overflows
 - ▶ No overflow pages! (except when the number of collisions exceed page capacity)
 - ★ Two data entries **collide** if they have the same $h(.)$ value

Another Idea for Dynamic Hashing (cont.)

00	32*	44*	36*	
01	9*	25*	5*	
10	10*	18*	26*	34*
11	31*	35*	7*	11*

Initial state with 4 buckets

00	32*	44*	36*	
01	9*	25*	5*	
010	10*	18*	26*	34*
11	31*	35*	7*	11*
110	6*			

After insertion of 6*


00	32*	44*	36*	
01	9*	25*	5*	
0010	2*	18*	34*	
11	31*	35*	7*	11*
110	6*			
1010	10*	26*		

After insertion of 2*

- ▶ The order in which buckets are split is random
- ▶ How to keep track of the number of last bits used as bucket addresses for different buckets?
- ▶ How to locate a bucket given its bucket address?


Extendible Hashing

- ▶ Uses a directory of pointers to buckets
 - ▶ Directory expands dynamically as buckets overflow
- ▶ Directory has 2^d entries
 - ▶ d is called **global depth** of the hashed file
 - ▶ Each directory entry has a unique d -bit address $b_d b_{d-1} \dots b_2 b_1$
 - ▶ Two directory entries are said to **correspond** if their addresses differ only in the d^{th} bit (i.e., b_d); such entries are called **corresponding entries**
 - ▶ Each bucket maintains a **local depth** (denoted by $\ell \in [0, d]$)
 - ★ All entries in a bucket with local depth ℓ have the same last ℓ bits in $h(.)$




Handling Bucket Overflow (Case 1)

- ▶ When a bucket overflows, it is split
 - ▶ Allocate a new bucket called its **split image**
 - ▶ Redistribute entries (including new entry) between split bucket & its split image
- ▶ Case 1: Split bucket's local depth is equal to global depth
- ▶ Example: Inserting 6^* (110)



Handling Bucket Overflow (Case 1)



- Initially, all buckets have local depth $\ell = d$
- A bucket split causes directory to double only if bucket's $\ell = d$
 - d is incremented by one whenever the directory is doubled
- ℓ is incremented by one for a bucket whenever it splits
- A split bucket & its image have the same local depth



Directory is expanded


Handling Bucket Overflow (Case 1)

- ▶ When the directory is doubled,
 - ▶ Each new directory entry (except for the entry for the split image) points to the same bucket as its corresponding entry
- ▶ Number of directory entries pointing to a bucket = $2^{d-\ell}$
- ▶ Example: Inserting 6^* (110)



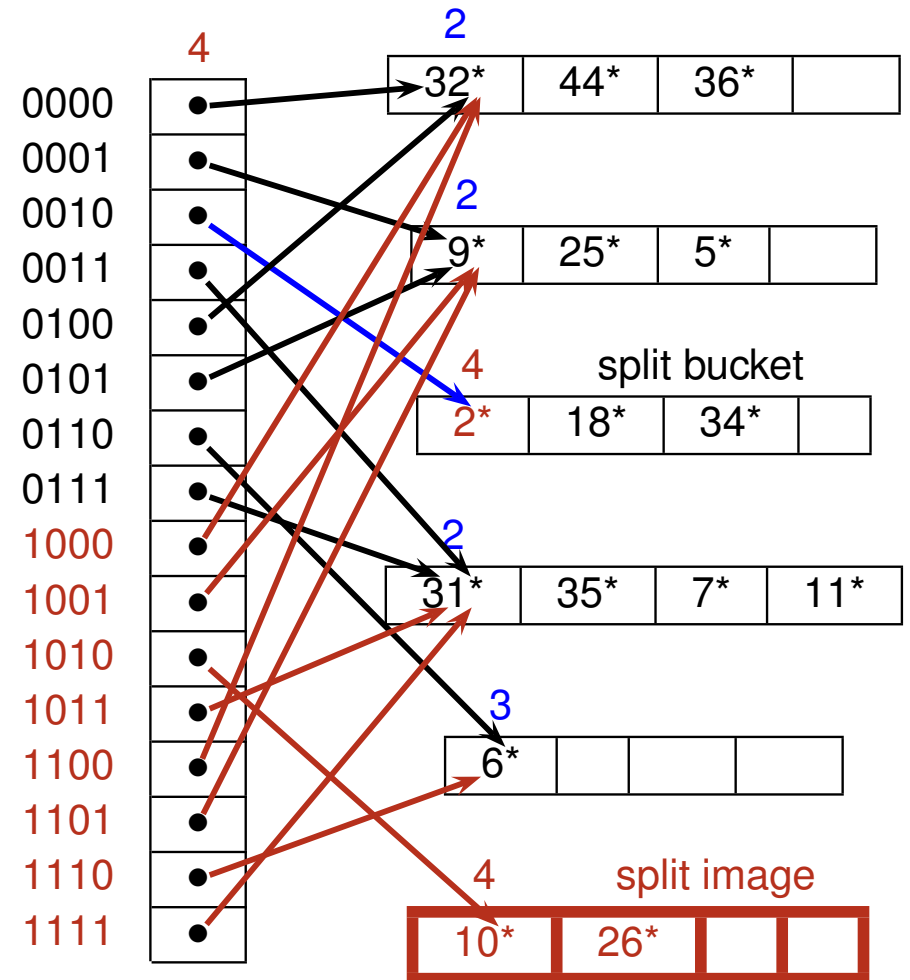

Another Bucket Overflow Example (Case 1)

- Example: Inserting 2* (010)



Another Bucket Overflow Example (Case 1)


► Example: Inserting 2* (010)



Directory is expanded


Handling Bucket Overflow (Case 2)

- ▶ Case 2: Split bucket's local depth < global depth
- ▶ Example: Inserting 22^* (010110)



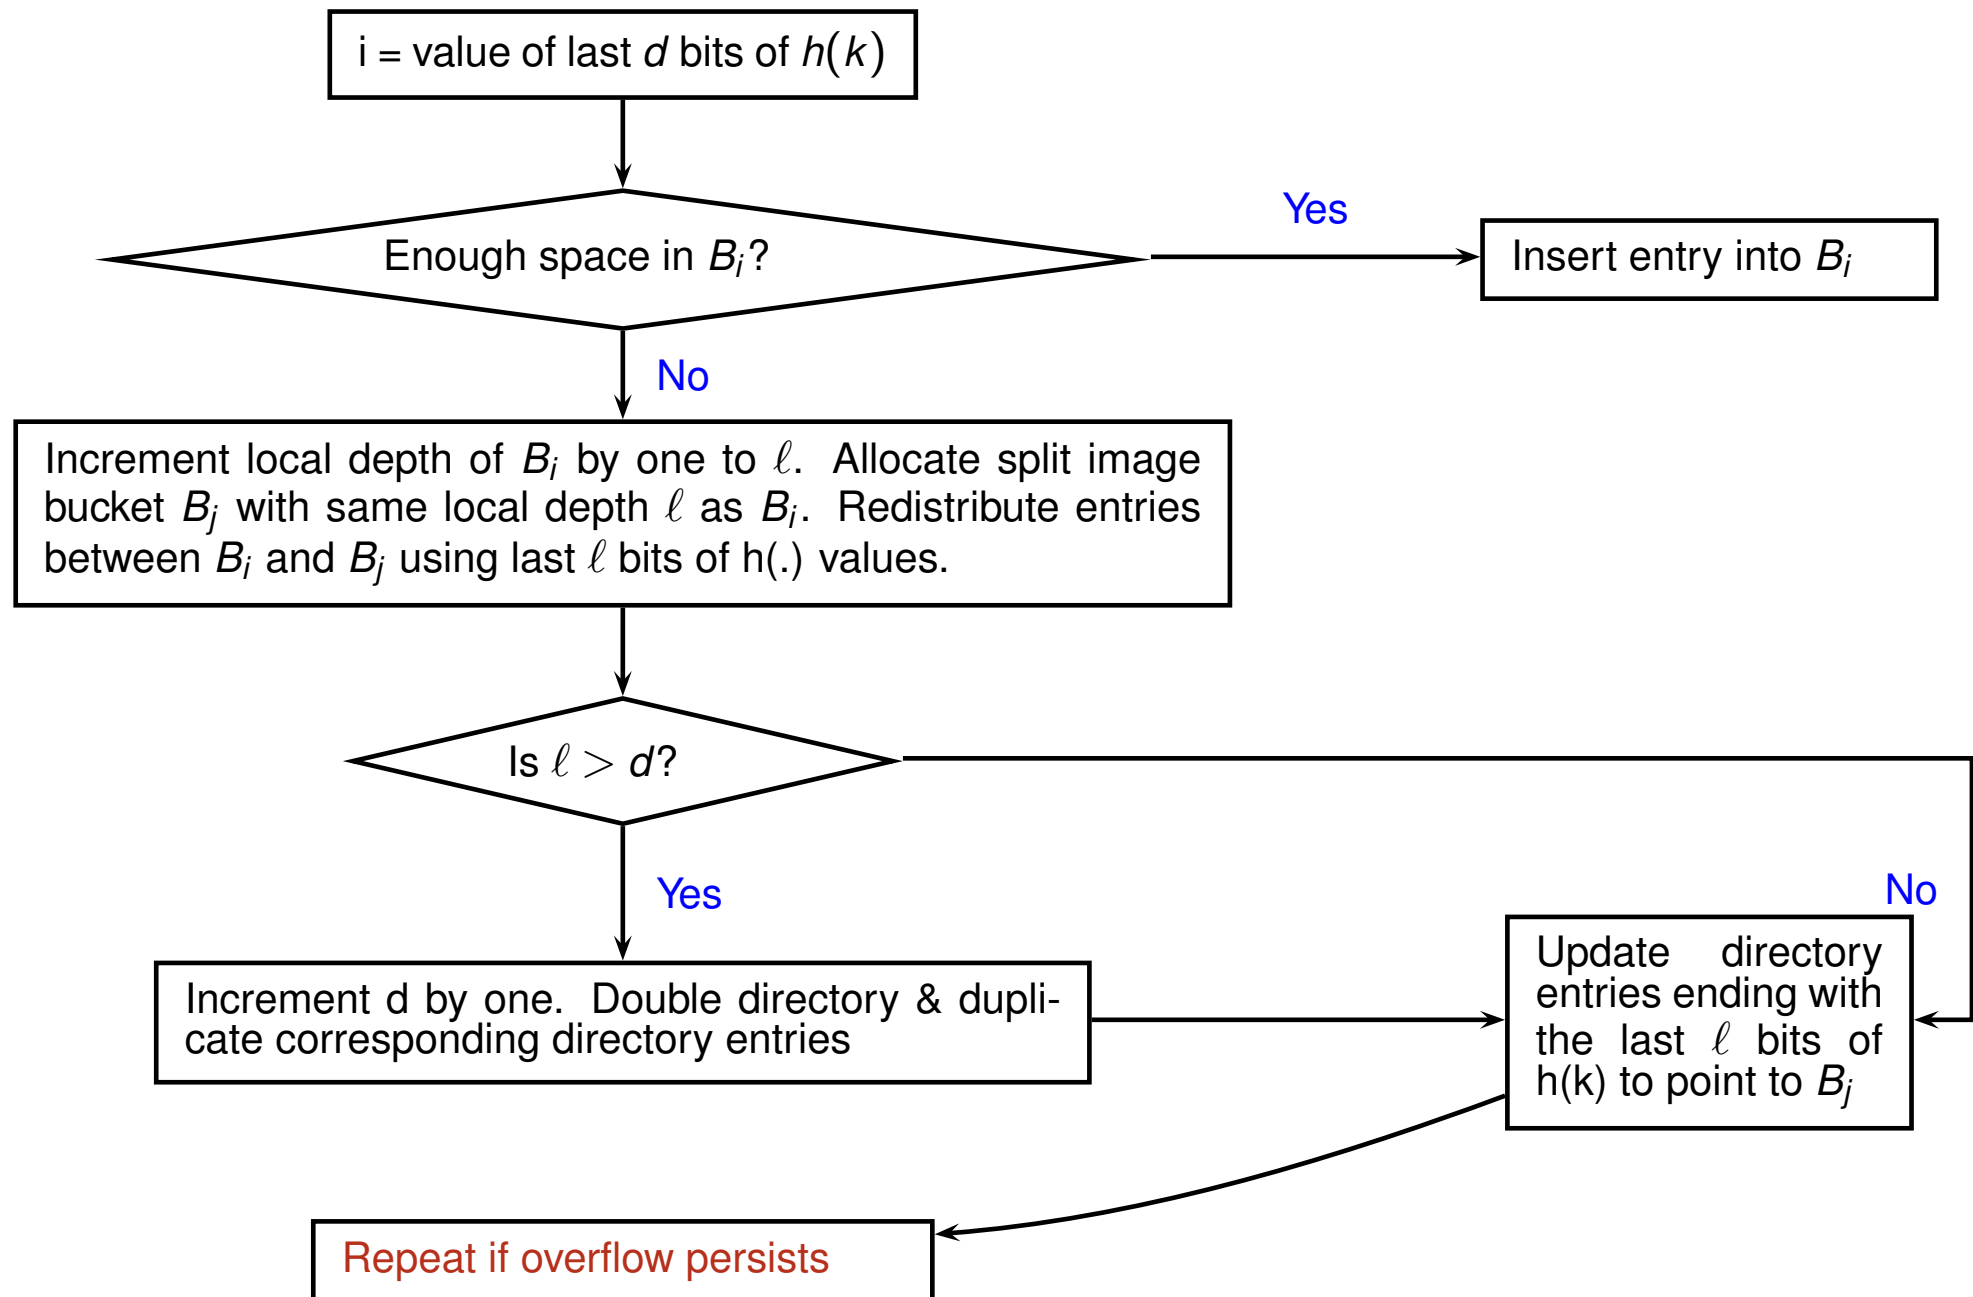
Handling Bucket Overflow (Case 2)

- ▶ Case 2: Split bucket's local depth < global depth
- ▶ Example: Inserting 22^* (010110)



Directory is not expanded




Inserting data entry with search key k



Extendible Hashing: Deletion

- ▶ Locate bucket B_i containing entry & delete entry
- ▶ If B_i becomes empty, B_i can be merged with the bucket B_j where i & j differs only in the ℓ^{th} bit
 - ▶ B_i is deallocated
 - ▶ ℓ is decremented by one
 - ▶ Directory entries that point to B_i are updated to point to B_j
- ▶ More generally, B_i & B_j (where i & j differs only in ℓ^{th} bit) can be merged if their entries can fit within a bucket
- ▶ If each pair of corresponding entries point to the same bucket, directory can be halved
 - ▶ d is decremented by one

Example: Deleting 10* (1010)



Extendible Hashing: Performance

- ▶ At most 2 disk I/Os for equality selection
 - ▶ At most one disk I/O if directory fits in main memory
- ▶ Handling collisions:
 - ▶ Two data entries **collide** if they have the same hashed value
 - ▶ Overflow pages are needed when the number of collisions exceeds page capacity