

Questions to be discussed: 1, 2, 3 & 4

1. Consider the Linear Hashing index shown below

level = 0, $N_0 = 1$, next=0
 0

2*		
----	--	--

Linear Hashing index with one data bucket

where each data bucket can store 3 data entries. Show the Linear Hashing index after inserting the following sequence of data entries:

3*, 5*, 7*, 11*, 17*, 19*, 23*, 0*, 4*, 29*, 33*, 13*, 6*

Assume that a bucket split is triggered whenever some bucket overflows.

2. (Exercise 11.9, R&G) Consider the snapshot of the Linear Hashing index shown below.

level = 0, $N_0 = 4$, next=0
 00

64*	44*		
-----	-----	--	--

 01

9*	25*	5*	
----	-----	----	--

 10

10*			
-----	--	--	--

 11

31*	15*	7*	3*
-----	-----	----	----

Assume that a bucket split is triggered whenever some bucket overflows.

- What is the maximum number of data entries that can be inserted (given the best possible distribution of keys) before you have to split a bucket? Explain very briefly.
 - Show the file after inserting a *single* record whose insertion causes a bucket split.
 - Consider the original Linear Hashing index shown.
 - What is the minimum number of record insertions that will cause a split of all four buckets? Explain very briefly.
 - What is the value of **next** after making these insertions?
 - What can you say about the number of pages in the fourth bucket shown after this series of record insertions?
3. Consider the Extendible Hashing index shown below

0 0
 0

•

 →

2*		
----	--	--

Extendible Hashing index with one data bucket

where each data bucket can store 3 data entries. Show the Extendible Hashing index after inserting the following sequence of data entries:

3*, 5*, 7*, 11*, 17*, 19*, 23*, 0*, 4*, 29*, 33*, 13* 6*

Note that k^* denotes a data entry having a hashed value of k .

4. (Adapted from Exercise 11.1, R&G) Consider the Extendible Hashing index shown below which was created by inserting a sequence of 13 data entries without any deletions. Which bucket could be the last bucket that was split by the insertions?

