

Questions to be discussed: 1, 2 & 3.

1. (Adapted from Exercise 14.3, R&G) Consider processing the following SQL projection query:

SELECT DISTINCT title, dname FROM Executives

You are given the following information:

- **Executives** has attributes **ename**, **title**, **dname**, and **address**; all are string fields of the same length.
- The **ename** attribute is a candidate key.
- The relation contains 10,000 pages.
- There are 10 buffer pages.

Consider the optimized version of the sorting-based projection algorithm: The initial sorting pass reads the input relation and creates sorted runs of tuples containing only attributes **dname** and **title**. Subsequent merging passes eliminate duplicates while merging the initial runs to obtain a single sorted result (as opposed to doing a separate pass to eliminate duplicates from a sorted result containing duplicates).

In this question, the cost metric is the number of page I/Os.

- (a) How many sorted runs are produced in the first pass? What is the average length of these runs? What is the I/O cost of this sorting pass?
- (b) How many additional merge passes are required to compute the final result of the projection query? What is the I/O cost of these additional passes?
- (c) (i) Suppose that a clustered B⁺-tree index on **title** is available. Is this index likely to offer a cheaper alternative to sorting? Would your answer change if the index were unclustered? Would your answer change if the index were a hash index?
- (ii) Suppose that a clustered B⁺-tree index on **dname** is available. Is this index likely to offer a cheaper alternative to sorting? Would your answer change if the index were unclustered? Would your answer change if the index were a hash index?
- (iii) Suppose that a clustered B⁺-tree index on (**dname**, **title**) is available. Is this index likely to offer a cheaper alternative to sorting? Would your answer change if the index were unclustered? Would your answer change if the index were a hash index?
- (d) Suppose that the query is as follows:

SELECT title, dname FROM Executives

That is, you are not required to do duplicate elimination. How would your answers to the previous questions change?

2. (Exercise 14.4 R&G) Consider the join $R \bowtie_{R.a=S.b} S$, given the following information about the relations to be joined. The cost metric is the number of page I/Os unless otherwise noted, and the cost of writing out the result should be uniformly ignored.
- Relation R contains 10,000 tuples and has 10 tuples per page.
 - Relation S contains 2000 tuples and also has 10 tuples per page.
 - Attribute b of relation S is the primary key for S.
 - Both relations are stored as simple heap files.
 - Neither relation has any indexes built on it.
 - 52 buffer pages are available.
- (a) What is the cost of joining R and S using a **page-oriented simple nested loops join**? What is the minimum number of buffer pages required for this cost to remain unchanged?
- (b) What is the cost of joining R and S using a **block nested loops join**? What is the minimum number of buffer pages required for this cost to remain unchanged?
- (c) What would be the lowest possible I/O cost for joining R and S using **any join algorithm**, and how much buffer space would be needed to achieve this cost? Explain briefly.
- (d) How many tuples does the join of R and S produce, at most, and how many pages are required to store the result of the join back on disk?
- (e) Would your answers to any of the previous questions in this exercise change if you were told that R.a is a foreign key that refers to S.b?

3. (Exercise 14.5 R&G) Consider the join $R \bowtie_{R.a=S.b} S$, given the following information about the relations to be joined. The cost metric is the number of page I/Os unless otherwise noted, and the cost of writing out the result should be uniformly ignored.
- Relation R contains 10,000 tuples and has 10 tuples per page.
 - Relation S contains 2000 tuples and also has 10 tuples per page.
 - Attribute b of relation S is the primary key for S.
 - Both relations are stored as simple heap files.
 - Neither relation has any indexes built on it.
 - 52 buffer pages are available.
 - Assume that all indexes are format 2, any B⁺-tree index on R has two levels of internal nodes, and any B⁺-tree index on S has one level of internal nodes.
- (a) With 52 buffer pages, if unclustered B⁺-tree indexes existed on R.a and S.b, would either provide a cheaper alternative for performing the join (using an **index nested loops join**) than a **block nested loops join**? Explain.
1. Would your answer change if only five buffer pages were available?
 2. Would your answer change if S contained only 10 tuples instead of 2000 tuples?
- (b) With 52 buffer pages, if clustered B⁺-tree indexes existed on R.a and S.b, would either provide a cheaper alternative for performing the join (using the **index nested loops algorithm**) than a **block nested loops join**? Explain.
1. Would your answer change if only five buffer pages were available?
 2. Would your answer change if S contained only 10 tuples instead of 2000 tuples?