

CS3223 Lecture 7

Transaction Management

Transactions

- ▶ A transaction is an abstraction representing a logical unit of work
- ▶ **Example:** `Transfer (x, y, amount)`
 - ▶ transaction to transfer `amount` from account `x` to account `y`

```
BEGIN TRANSACTION;  
SELECT balance INTO :Balx FROM Account WHERE accountId = :x;  
  
SELECT balance INTO :Baly FROM Account WHERE accountId = :y;  
  
If (:Balx < amount) then ABORT;  
  
UPDATE Account SET balance = :Baly + :amount WHERE accountId = :y;  
  
UPDATE Account SET balance = :Balx - :amount WHERE accountId = :x;  
  
COMMIT;
```

Transaction Management

- ▶ Ensures four properties of transactions (Xacts) to maintain data in the face of concurrent access and system failures
 1. **Atomicity**: Either all or none of the actions in Xact happen
 2. **Consistency**: If each Xact is consistent, and the DB starts consistent, the DB ends up consistent
 3. **Isolation**: Execution of one Xact is isolated from other Xacts
 4. **Durability**: If a Xact commits, its effects persist
- ▶ The **concurrency control manager** component ensures isolation
- ▶ The **recovery manager** component ensures atomicity and durability

Transactions

- ▶ A transaction (Xact) T_i can be viewed as a sequence of **actions**:
 - ▶ $R_i(O)$ = T_i reads an object O
 - ▶ $W_i(O)$ = T_i writes an object O
 - ▶ $Commit_i$ = T_i terminates successfully
 - ▶ $Abort_i$ = T_i terminates unsuccessfully
- ▶ Each Xact must end with either a commit or an abort action
- ▶ An **active Xact** is a Xact that is still in progress (i.e., has not yet terminated)

Transactions

- ▶ $R(x), R(y), W(y), W(x)$, Commit
- ▶ $R(x), R(y)$, Abort

```
BEGIN TRANSACTION;  
SELECT balance INTO :Balx FROM Account WHERE accountId = :x;  
  
SELECT balance INTO :Baly FROM Account WHERE accountId = :y;  
  
If (:Balx < amount) then ABORT;  
  
UPDATE Account SET balance = :Baly + :amount WHERE accountId = :y;  
  
UPDATE Account SET balance = :Balx - :amount WHERE accountId = :x;  
  
COMMIT;
```

Transaction Schedules

- ▶ **Schedule** = a list of actions from a set of Xacts, where the order of the actions within each Xact is preserved
- ▶ **Example:** Consider two Xacts T_1 and T_2 :
 - ▶ T_1 : $R_1(A)$, $W_1(A)$, $R_1(B)$, $W_1(B)$, $Commit_1$
 - ▶ T_2 : $R_2(A)$, $W_2(A)$, $R_2(B)$, $W_2(B)$, $Commit_2$
- ▶ Some schedules of T_1 and T_2 :
 - S_1 : $R_1(A)$, $W_1(A)$, $R_1(B)$, $W_1(B)$, $Commit_1$, $R_2(A)$, $W_2(A)$, $R_2(B)$, $W_2(B)$, $Commit_2$
 - S_2 : $R_2(A)$, $W_2(A)$, $R_2(B)$, $W_2(B)$, $Commit_2$, $R_1(A)$, $W_1(A)$, $R_1(B)$, $W_1(B)$, $Commit_1$
 - S_3 : $R_1(A)$, $W_1(A)$, $R_2(A)$, $W_2(A)$, $R_1(B)$, $W_1(B)$, $Commit_1$, $R_2(B)$, $W_2(B)$, $Commit_2$
- ▶ A **serial schedule** is a schedule where the actions of Xacts are not interleaved

Read From & Final Write

- ▶ We say that T_j reads O from T_i in a schedule S if the last write action on O before $R_j(O)$ in S is $W_i(O)$
- ▶ We say that T_j reads from T_i if T_j has read some object from T_i
- ▶ We say that T_i performs the final write on O in a schedule S if the last write action on O in S is $W_i(O)$
- ▶ **Example:** Consider the following schedule:

$R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), \text{Commit}_1, R_2(B), W_2(B), \text{Commit}_2$

- ▶ T_1 reads A from the initial database (or T_1 reads A from T_0)
 - ★ Assume that initial database was created by dummy Xact T_0
- ▶ T_2 reads A from T_1
- ▶ T_1 reads B from T_0
- ▶ T_2 reads B from T_1
- ▶ T_2 performs the final write on A
- ▶ T_2 performs the final write on B

Interleaved Transaction Executions

T_1 : Transfer(x,y,100) T_2 : Transfer(y,x,100)

T_1	T_2	Comments
		$x = 400, y = 100$
$R_1(x)$		400
$R_1(y)$		100
$y = y + 100$		
$W_1(y)$		$y = 200$
	$R_2(y)$	200
	$R_2(x)$	400
$x = x - 100$		
$W_1(x)$		$x = 300$
	$x = x + 100$	
	$W_2(x)$	$x = 500$
	$y = y - 100$	
	$W_2(y)$	$y = 100$

Initial database state: $x = 400, y = 100$

Final database state: $x = 500, y = 100$

Correctness of Interleaved Xact Executions

An interleaved Xact execution schedule is **correct** if it is “equivalent” to some serial schedule over the same set of Xacts

Execution Schedule	Serial Schedules	
	S_1	S_2
$R_1(x, 400)$	$R_1(x, 400)$	$R_2(y, 100)$
$R_1(y, 100)$	$R_1(y, 100)$	$R_2(x, 400)$
$W_1(y, 200)$	$W_1(y, 200)$	$W_2(x, 500)$
$R_2(y, 200)$	$W_1(x, 300)$	$W_2(y, 0)$
$R_2(x, 400)$	$R_2(y, 200)$	$R_1(x, 500)$
$W_1(x, 300)$	$R_2(x, 300)$	$R_1(y, 0)$
$W_2(x, 500)$	$W_2(x, 400)$	$W_1(y, 100)$
$W_2(y, 100)$	$W_2(y, 100)$	$W_1(x, 400)$

View Serializable Schedules

- ▶ Two schedules S and S' (over the same set of Xacts) are **view equivalent** (denoted by $S \equiv_v S'$) if they satisfy all the following conditions:

1. If T_i reads A from T_j in S , then T_i must also read A from T_j in S'
2. For each data object A , the Xact (if any) that performs the final write on A in S must also perform the final write on A in S'

- ▶ **Example:**

- ▶ S_3 : $R_1(x), R_1(y), W_1(y), R_2(y), R_2(x), W_1(x), W_2(x), W_2(y)$
- ▶ S_5 : $R_1(x), R_2(y), R_1(y), W_1(y), R_2(x), W_2(x), W_2(y), W_1(x)$
- ▶ S_6 : $R_1(x), R_2(y), R_1(y), R_2(x), W_1(y), W_2(x), W_1(x), W_2(y)$
- ▶ $S_3 \not\equiv_v S_5$ as T_2 reads y from T_1 in S_3 but T_2 reads y from T_0 in S_5
- ▶ Similarly, $S_3 \not\equiv_v S_6$
- ▶ $S_5 \equiv_v S_6$

View Serializable Schedules (cont.)

- ▶ A schedule S is a **view serializable schedule (VSS)** if S is view equivalent to some serial schedule over the same set of Xacts
- ▶ **Example:**
 - ▶ Consider two Xacts T_1 and T_2 :
 - ★ $T_1: R_1(x), R_1(y), W_1(y), W_1(x)$
 - ★ $T_2: R_2(y), R_2(x), W_2(x), W_2(y)$
 - ▶ Serial schedules over $\{T_1, T_2\}$:
 - ★ $S_1: R_1(x), R_1(y), W_1(y), W_1(x), R_2(y), R_2(x), W_2(x), W_2(y)$
 - ★ $S_2: R_2(y), R_2(x), W_2(x), W_2(y), R_1(x), R_1(y), W_1(y), W_1(x)$
 - ▶ S_3 is not view serializable; S_4 is view serializable
 - ★ $S_3: R_1(x), R_1(y), W_1(y), R_2(y), R_2(x), W_1(x), W_2(x), W_2(y)$
 - ★ $S_4: R_1(x), R_1(y), W_1(y), R_2(y), W_1(x), R_2(x), W_2(x), W_2(y)$

Conflicting Actions

- ▶ Two actions on the same object **conflict** if
 1. at least one of them is a write action, and
 2. the actions are from different Xacts
- ▶ **Examples:**
 - ▶ $R_1(x)$ and $R_2(x)$ do not conflict
 - ▶ $R_1(x)$ and $W_1(x)$ do not conflict
 - ▶ $W_1(x)$ and $R_2(y)$ do not conflict
 - ▶ $W_1(x)$ and $R_2(x)$ conflict
 - ▶ $W_1(x)$ and $W_2(x)$ conflict

Anomalies with Interleaved Xact Executions

► Anomalies can arise due to conflicting actions:

1. Dirty read problem (due to WR conflicts)

- ★ T_2 reads an object that has been modified by T_1 and T_1 has not yet committed
- ★ T_2 could see an inconsistent DB state!

2. Unrepeatable read problem (due to RW conflicts)

- ★ T_2 updates an object that T_1 has previously read and T_2 commits while T_1 is still in progress
- ★ T_1 could get a different value if it reads the object again!

3. Lost update problem (due to WW conflicts)

- ★ T_2 overwrites the value of an object that has been modified by T_1 while T_1 is still in progress
- ★ T_1 's update is lost!

Dirty Read Problem: Example

T_1	T_2	Comments
		Initial DB state: $x = 100$
R(x)		100
$x = x + 20$		
W(x)		$x = 120$
	R(x)	120
	$x = x \times 2$	
Abort		
	W(x)	$x = 240$

- For every serial schedule over $\{T_1, T_2\}$, the final value of x is 200

Unrepeatable Read Problem: Example

T_1	T_2	Comments
		Initial DB state: $x = 100$
$R(x)$		100
	$R(x)$	100
	$x = x - 20$	
	$W(x)$	$x = 80$
	Commit	
$R(x)$		80

- For every serial schedule over $\{T_1, T_2\}$, both values read by T_1 are the same

Lost Update Problem: Example

T_1	T_2	Comments
		Initial DB state: $x = 100$
R(x)		100
	R(x)	100
$x = x + 20$		
	$x = x \times 2$	
W(x)		$x = 120$
	W(x)	$x = 200$

- For serial schedule (T_1, T_2) , the final value of x is 240
- For serial schedule (T_2, T_1) , the final value of x is 220

Conflict Serializable Schedules


- ▶ Two schedules S & S' (over the same set of Xacts) are said to be **conflict equivalent** (denoted by $S \equiv_c S'$) if they order every pair of conflicting actions of two committed Xacts in the same way
- ▶ **Example:**
 - ▶ S_3 : $R_1(x), R_1(y), W_1(y), R_2(y), R_2(x), W_1(x), W_2(x), W_2(y)$
 - ▶ S_5 : $R_1(x), R_2(y), R_1(y), W_1(y), R_2(x), W_2(x), W_2(y), W_1(x)$
 - ▶ S_6 : $R_1(x), R_2(y), R_1(y), R_2(x), W_1(y), W_2(x), W_1(x), W_2(y)$
 - ▶ $S_3 \not\equiv_c S_5$ as $W_1(y)$ precedes $R_2(y)$ in S_3 but not in S_5
 - ▶ Similarly, $S_3 \not\equiv_c S_6$
 - ▶ $S_5 \equiv_c S_6$

Conflict Serializable Schedules (cont.)

- ▶ A schedule is a **conflict serializable schedule (CSS)** if it is conflict equivalent to a serial schedule over the same set of Xacts
- ▶ **Example:**
 - ▶ Consider two Xacts T_1 and T_2 :
 - ★ $T_1: R_1(x), R_1(y), W_1(y), W_1(x)$
 - ★ $T_2: R_2(y), R_2(x), W_2(x), W_2(y)$
 - ▶ Serial schedules over $\{T_1, T_2\}$:
 - ★ $S_1: R_1(x), R_1(y), W_1(y), W_1(x), R_2(y), R_2(x), W_2(x), W_2(y)$
 - ★ $S_2: R_2(y), R_2(x), W_2(x), W_2(y), R_1(x), R_1(y), W_1(y), W_1(x)$
 - ▶ S_3 is not conflict serializable; S_4 is conflict serializable
 - ★ $S_3: R_1(x), R_1(y), W_1(y), R_2(y), R_2(x), W_1(x), W_2(x), W_2(y)$
 - ★ $S_4: R_1(x), R_1(y), W_1(y), R_2(y), W_1(x), R_2(x), W_2(x), W_2(y)$

Testing for Conflict Serializability

- ▶ A **conflict serializability graph** for a schedule S (denoted by $CSG(S)$) is a directed graph $G = (V, E)$ such that
 - ▶ V contains a node for each committed Xact in S
 - ▶ E contains (T_i, T_j) if an action in T_i precedes and conflicts with one of T_j 's actions
- ▶ **Example:** Conflict serializability graph for schedule $R_1(A), W_2(A), Commit_2, W_1(A), Commit_1, W_3(A), Commit_3$



- ▶ **Theorem 1:** A schedule is conflict serializable iff its conflict serializability graph is acyclic

Conflict Serializable Schedules

- ▶ **Theorem 2:** A schedule that is conflict serializable is also view serializable
- ▶ For convenience, we use **serializable** to mean conflict serializable

Blind Writes

- ▶ A write on object O by T_i is call a **blind write** if T_i did not read O prior to the write
 - ▶ Consider the schedule: $R_1(x), W_2(y), W_1(x)$
 - ▶ $W_2(y)$ is a blind write
 - ▶ $W_1(x)$ is a non-blind write
- ▶ **Theorem 3:** If S is view serializable and S has no blind writes, then S is also conflict serializable

Cascading Aborts

- ▶ For correctness, if T_i has read from T_j , then T_i must abort if T_j aborts

T_1	T_2	T_1	T_2
$W_1(x)$		$W_1(x)$	
	$R_2(x)$		$R_2(x)$
	$W_2(y)$		$W_2(y)$
<i>Abort₁</i>		<i>Abort₁</i>	
			<i>Abort₂</i>

- ▶ We say that T_1 's abort is cascaded to T_2
- ▶ Recursive aborting process is known as **cascading abort**

Recoverable Schedules

T_1	T_2
$W_1(x)$	$R_2(x)$ $W_2(y)$ <i>Commit₂</i>

Non-Recoverable
Schedule

- ▶ A schedule S is said to be a **recoverable schedule** if for every Xact T that commits in S , T must commit after T' if T reads from T'

Cascadeless Schedules

- ▶ While recoverable schedules guarantee that committed Xacts will not be aborted, cascading aborts of active Xacts are possible
 - ▶ **Example:** if T_i reads from T_j and T_j aborts, T_i must also abort
- ▶ Cascading aborts are undesirable because of the cost of bookkeeping to identify them and the performance penalty incurred
- ▶ To avoid cascading aborts (or to be **cascadeless**), DBMS must permit reads only from committed Xacts
- ▶ A schedule S is a **cascadeless schedule** if whenever T_i reads from T_j in S , $Commit_j$ must precede this read action
- ▶ **Theorem 4:** A cascadeless schedule is also a recoverable schedule

Recovery using Before-Images

- ▶ An efficient approach to undo the actions of aborted Xacts is to restore *before-images* for writes
- ▶ **Example:** Consider the following schedule S :

$W_1(A), W_2(A), Abort_2$

where T_1 updates A to be 100 and T_2 updates A to be 200

- ▶ Assume that the initial value of A is 50
- ▶ Before performing $W_1(A)$, its before-image “ $A=50$ ” is logged
- ▶ Before performing $W_2(A)$, its before-image “ $A=100$ ” is logged
- ▶ To recover from $Abort_2$, $W_2(A)$ is undone by restoring the before-image of A (i.e., the value of A is restored to 100)

Recovery using Before-Images (cont.)

- ▶ However, before-image recovery doesn't always work!
- ▶ **Example:**

$W_1(A)$, $W_2(A)$, $Abort_1$

- ▶ Here, undoing $W_1(A)$ by restoring A to its before-image of 50 is incorrect!

Strict Schedules

- ▶ To enable the use of before-images for recovery, we use *strict schedules*
- ▶ A schedule S is a **strict schedule** if for every $W_i(O)$ in S , O is not read or written by another T_j until T_i either aborts or commits

- ▶ **Example:**

$S:$ $W_1(A)$, $W_2(A)$, $Abort_2$

$S':$ $W_1(A)$, $W_2(A)$, $Abort_1$

Both S and S' are not strict schedules

- ▶ Performance Tradeoff:
 - ▶ recovery (using before-images) is more efficient
 - ▶ concurrent executions become more restrictive

- ▶ **Theorem 5:** A strict schedule is also a cascadeless schedule