

Questions to be discussed: 3, 4 & 5.

1. (Exercise 17.2, R&G) Consider the following classes of schedules: recoverable, cascadeless, and strict. For each of the following schedules, state which of the schedule classes it belongs to.
- (a)  $R_1(X), R_2(X), W_1(X), W_2(X), Commit_1, Commit_2$
  - (b)  $W_1(X), R_2(Y), R_1(Y), R_2(X), Commit_2, Commit_1$
  - (c)  $R_1(X), R_2(Y), W_3(X), R_2(X), R_1(Y), Commit_1, Commit_2, Commit_3$
  - (d)  $R_1(X), R_1(Y), W_1(X), R_2(Y), W_3(Y), W_1(X), R_2(Y), Commit_1, Commit_2, Commit_3$
  - (e)  $R_1(X), W_2(X), W_1(X), Abort_2, Commit_1$
  - (f)  $R_1(X), W_2(X), W_1(X), Commit_2, Commit_1$
  - (g)  $W_1(X), R_2(X), W_1(X), Abort_2, Commit_1$
  - (h)  $W_1(X), R_2(X), W_1(X), Commit_2, Commit_1$
  - (i)  $W_1(X), R_2(X), W_1(X), Commit_2, Abort_1$
  - (j)  $R_2(X), W_3(X), Commit_3, W_1(Y), Commit_1, R_2(Y), W_2(Z), Commit_2$
  - (k)  $R_1(X), W_2(X), Commit_2, W_1(X), Commit_1, R_3(X), Commit_3$
  - (l)  $R_1(X), W_2(X), W_1(X), R_3(X), Commit_1, Commit_2, Commit_3$

**Solution:**

- (a) recoverable, cascadeless, not strict.
- (b) not recoverable, not cascadeless, not strict.
- (c) not recoverable, not cascadeless. not strict.
- (d) not recoverable, not cascadeless. not strict.
- (e) recoverable, cascadeless. not strict.
- (f) recoverable, cascadeless. not strict.
- (g) recoverable, not cascadeless, not strict.
- (h) not recoverable, not cascadeless. not strict.
- (i) not recoverable, not cascadeless. not strict.
- (j) recoverable, cascadeless, strict.
- (k) recoverable, cascadeless, strict.
- (l) recoverable, not cascadeless. not strict.

2. (Exercise 17.3, R&G) Consider the following two concurrency control protocols: 2PL and Strict 2PL. For each of the following schedules, state which of these protocols allows it, that is, allows the actions to occur in exactly the order shown. Assume that each lock acquisition step must occur right before the read/write action (e.g., it's not permissible to request for all the required locks at the beginning of the transaction).
- (a)  $R_1(X), R_2(X), W_1(X), W_2(X), Commit_1, Commit_2$
  - (b)  $W_1(X), R_2(Y), R_1(Y), R_2(X), Commit_2, Commit_1$
  - (c)  $R_1(X), R_2(Y), W_3(X), R_2(X), R_1(Y), Commit_1, Commit_2, Commit_3$
  - (d)  $R_1(X), R_1(Y), W_1(X), R_2(Y), W_3(Y), W_1(X), R_2(Y), Commit_1, Commit_2, Commit_3$
  - (e)  $R_1(X), W_2(X), W_1(X), Abort_2, Commit_1$
  - (f)  $R_1(X), W_2(X), W_1(X), Commit_2, Commit_1$
  - (g)  $W_1(X), R_2(X), W_1(X), Abort_2, Commit_1$
  - (h)  $W_1(X), R_2(X), W_1(X), Commit_2, Commit_1$
  - (i)  $W_1(X), R_2(X), W_1(X), Commit_2, Abort_1$
  - (j)  $R_2(X), W_3(X), Commit_3, W_1(Y), Commit_1, R_2(Y), W_2(Z), Commit_2$
  - (k)  $R_1(X), W_2(X), Commit_2, W_1(X), Commit_1, R_3(X), Commit_3$
  - (l)  $R_1(X), W_2(X), W_1(X), R_3(X), Commit_1, Commit_2, Commit_3$

**Solution:**

- (a) **Not 2PL, not S2PL.**
- (b) **2PL, not S2PL.**
- (c) **not 2PL, not S2PL.**
- (d) **Not 2PL, not S2PL.**
- (e) **Not 2PL, not S2PL.**
- (f) **Not 2PL, not S2PL.**
- (g) **Not 2PL, not S2PL.**
- (h) **Not 2PL, not S2PL.**
- (i) **Not 2PL, not S2PL.**
- (j) **Not 2PL, not S2PL.**
- (k) **Not 2PL, not S2PL.**
- (l) **Not 2PL, not S2PL.**

3. Give an example schedule that is conflict serializable but not a 2PL schedule.

**Solution:** Consider the schedule  $W_1(x), W_2(x), Commit_2, W_3(y), Commit_3, R_1(y), Commit_1$ . It is conflict equivalent to  $(T_3, T_1, T_2)$  but it is not a 2PL schedule as  $T_1$  need to release its exclusive lock on  $x$  after  $W_1(x)$  (to enable  $W_2(x)$ ) and later request a shared lock for  $R_1(y)$ .

4. (Exercise 17.4, R&G) Consider the following two sequences of actions, listed in the order they are submitted to the DBMS.

Sequence S1:  $R_1(X), W_2(X), W_2(Y), W_3(Y), W_1(Y), Commit_1, Commit_2, Commit_3$

Sequence S2:  $R_1(X), W_2(Y), W_2(X), W_3(Y), W_1(Y), Commit_1, Commit_2, Commit_3$

For each sequence and for each of the following concurrency control mechanisms, describe how the concurrency control mechanism handles the sequence. Assume that the start timestamp of transaction  $T_i$  is  $i$ , and a transaction  $T_i$  has a higher priority than another transaction  $T_j$  if  $i < j$ . For lock-based concurrency control mechanisms, add lock and unlock requests to the previous sequence of actions as per the locking protocol. The DBMS processes actions in the order shown. If a transaction is blocked, assume that all its actions are queued until it is resumed; the DBMS continues with the next action (according to the listed sequence) of an unblocked transaction.

- (a) Strict 2PL with timestamps used for Wait-die deadlock prevention policy.
- (b) Strict 2PL with timestamps used for Wound-wait deadlock prevention policy.
- (c) Strict 2PL with deadlock detection. (Show the waits-for graph in case of deadlock.)

**Solution:**

- (a) In Wait-die policy, whenever a lock request can't be granted, the requesting transaction is blocked if it has higher priority; and is aborted, otherwise.

- **Sequence S1:**

- $T_1$  acquires S-lock on X for  $R_1(X)$ .
- When  $T_2$  requests for X-lock on X,  $T_2$  is aborted since  $T_2$  has lower priority than  $T_1$ ,
- $T_3$  acquires X-lock on Y for  $W_3(Y)$ .
- When  $T_1$  requests for X-lock on Y for  $W_1(Y)$ ,  $T_1$  is blocked since  $T_1$  has higher priority than  $T_3$ ,
- $T_3$  commits and releases its X-lock on Y.
- $T_1$  acquires X-lock on Y and resumes execution.
- $T_1$  commits and releases its S-lock on X and X-lock on Y.
- $T_2$  now can be restarted successfully.

- **Sequence S2:**

- $T_1$  acquires S-lock on X for  $R_1(X)$ .
- $T_2$  acquires X-lock on Y for  $W_2(Y)$ .
- When  $T_2$  requests for X-lock on X,  $T_2$  is aborted since  $T_2$  has lower priority than  $T_1$ . Thus,  $T_2$  releases its X-lock on Y.
- $T_3$  acquires X-lock on Y for  $W_3(Y)$ .
- When  $T_1$  requests for X-lock on Y,  $T_1$  is blocked since  $T_1$  has higher priority than  $T_3$ .
- $T_3$  commits and releases its X-lock on Y.
- $T_1$  now acquires X-lock on Y and resumes execution.
- $T_1$  commits and releases its S-lock on X and X-lock on Y.
- $T_2$  can now be restarted successfully.

- (b) In Wound-wait policy, whenever a lock request can't be granted, the requesting transaction aborts the other transaction if the requesting transaction has higher priority; and gets blocked, otherwise.

- **Sequence S1:**

- $T_1$  acquires S-lock on X for  $R_1(X)$ .
- When  $T_2$  requests for X-lock on X,  $T_2$  is blocked since  $T_2$  has lower priority than  $T_1$ .
- $T_3$  acquires X-lock on Y for  $W_3(Y)$ .
- When  $T_1$  requests for X-lock on Y,  $T_3$  is aborted since  $T_1$  has higher priority than  $T_3$ . Thus,  $T_3$  releases its X-lock on Y.
- $T_1$  now acquires X-lock on Y.
- $T_1$  commits and releases its S-lock on X and X-lock on Y.
- $T_2$  now acquires X-lock on X for  $W_2(X)$  and acquires X-lock on Y for  $W_2(Y)$
- $T_2$  commits and releases its X-locks on X & Y.
- $T_3$  now can be restarted successfully.

- **Sequence S2:**

- $T_1$  acquires S-lock on X for  $R_1(X)$ .
- $T_2$  acquires X-lock on Y for  $W_2(Y)$ .
- When  $T_2$  requests for X-lock on X,  $T_2$  is blocked since  $T_2$  has lower priority than  $T_1$ .
- When  $T_3$  requests for X-lock on Y,  $T_3$  is blocked since  $T_3$  has lower priority than  $T_2$ .
- When  $T_1$  requests for X-lock on Y,  $T_2$  is aborted since  $T_1$  has higher priority than  $T_2$ . Thus,  $T_2$  releases its X-lock on Y and  $T_1$  acquires X-lock on Y for  $W_1(Y)$ .
- $T_1$  commits and releases its S-lock on X and X-lock on Y.
- $T_2$  restarts and acquires X-lock on Y for  $W_2(Y)$ .
- $T_2$  acquires X-lock on X for  $W_2(X)$ .
- $T_2$  commits and releases its X-locks on X & Y.
- $T_3$  resumes execution and acquires X-lock on Y for  $W_3(Y)$ .
- $T_3$  commits and releases its lock on Y.

(c) • **Sequence S1:**

- $T_1$  acquires S-lock on X for  $R_1(X)$ .
- When  $T_2$  requests for X-lock on X,  $T_2$  is blocked and  $T_2 \rightarrow T_1$  is added to WFG.
- $T_3$  acquires X-lock on Y for  $W_3(Y)$ .
- When  $T_1$  requests for X-lock on Y,  $T_1$  is blocked and  $T_1 \rightarrow T_3$  is added to WFG.
- $T_3$  commits and releases its X-lock on Y;  $T_1 \rightarrow T_3$  is removed from WFG.
- $T_1$  acquires X-lock on Y and resumes execution.
- $T_1$  commits and releases its S-lock on X and X-lock on Y;  $T_2 \rightarrow T_1$  is removed from WFG.
- $T_2$  now acquires exclusive-locks on X and Y and commits after its execution.

- **Sequence S2:**

- $T_1$  acquires S-lock on X for  $R_1(X)$ .
- $T_2$  acquires X-lock on Y for  $W_2(Y)$ .
- When  $T_2$  requests for X-lock on X,  $T_2$  is blocked and  $T_2 \rightarrow T_1$  is added to WFG.
- When  $T_3$  requests for X-lock on Y,  $T_3$  is blocked and  $T_3 \rightarrow T_2$  is added to WFG.
- When  $T_1$  requests for X-lock on Y,  $T_1$  is blocked and  $T_1 \rightarrow T_2$  is added to WFG.
- A deadlock between  $T_1$  and  $T_2$  is detected.

- Assume that  $T_2$  is aborted to resolve the deadlock. Thus,  $T_2$  releases its X-lock on Y;  $T_2 \rightarrow T_1$ ,  $T_3 \rightarrow T_2$ , and  $T_1 \rightarrow T_2$  are removed from WFG.
- $T_1$  acquires X-lock on Y for  $W_1(Y)$ .
- $T_1$  commits and releases its S-lock on X and X-lock on Y.
- $T_2$  resumes execution and acquires X-lock on Y for  $W_2(Y)$ .
- $T_2$  acquires X-lock on X for  $W_2(X)$ .
- When  $T_3$  resumes execution and requests for X-lock on Y,  $T_3$  is blocked and  $T_3 \rightarrow T_2$  is added to WFG.
- $T_2$  commits and releases its locks on X & Y;  $T_3 \rightarrow T_2$  is removed from WFG.
- $T_3$  resumes execution and acquires X-lock on Y for  $W_3(Y)$ .
- $T_3$  commits and releases its X-lock on Y.

5. For each of the following pairs of transactions  $T_1$  and  $T_2$ , state whether every interleaved execution of  $T_1$  and  $T_2$  is conflict serializable. If your answer is “No”, write down an example schedule that is not conflict serializable.

In this question, assume the following:

1. Each transaction commits and is executed at the READ COMMITTED isolation level: short duration locks are acquired for read operations and long duration locks are acquired for write operations; transactions are allowed to acquire further locks after releasing locks.
  2. Short duration locks are released right after the associated operations are performed.
  3. Lock conversions are not used.
- (a)  $T_1$ :  $R_1(x), W_1(y), Commit_1$   
 $T_2$ :  $W_2(x), Commit_2$
- (b)  $T_1$ :  $R_1(x), W_1(y), Commit_1$   
 $T_2$ :  $R_2(y), W_2(x), Commit_2$
- (c)  $T_1$ :  $W_1(x), Commit_1$   
 $T_2$ :  $R_2(y), W_2(x), Commit_2$
- (d)  $T_1$ :  $R_1(x), W_1(y), Commit_1$   
 $T_2$ :  $R_2(x), W_2(x), W_2(y), Commit_2$

**Solution:**

- (a) Yes. The only conflicting pair of actions is between  $R_1(x)$  and  $W_2(x)$ . If  $T_1$  is granted the lock on  $x$  before  $T_2$ , then the execution is conflict equivalent to the serial schedule  $(T_1, T_2)$ ; otherwise, the execution is the serial schedule  $(T_2, T_1)$ .
- (b) No. Example:  
 $S_1(x), R_1(x), U_1(x), S_2(y), R_2(y), U_2(y), X_2(x), W_2(x), Commit_2, U_2(x), X_1(y)W_1(y), Commit_1, U_1(y)$ .  
 Here, we have the conflicting action pairs  $(R_1(x), W_2(x))$  and  $(W_2(y), W_1(y))$ .
- (c) Yes. The only conflicting pair of actions is between  $W_1(x)$  and  $W_2(x)$ . If  $T_1$  is granted the lock on  $x$  before  $T_2$ , then the execution is conflict equivalent to the serial schedule  $(T_1, T_2)$ ; otherwise, the execution is the serial schedule  $(T_2, T_1)$ .
- (d) No. Example:  
 $S_1(x), R_1(x), U_1(x), S_2(x), R_2(x), U_2(x), X_2(x), W_2(x), X_2(y), W_2(y), Commit_2, U_2(x), U_2(y), X_1(y), W_1(y), Commit_1, U_1(y)$ .  
 Here, we have the conflicting action pairs  $(R_1(x), W_2(x))$  and  $(W_2(y), W_1(y))$ .

6. For each transaction  $T_i$  in a 2PL schedule  $S$ , let  $lastlocktime(T_i)$  denote the time that  $T_i$  acquires its last lock in  $S$ .

For example, consider the following schedule  $S$ :

$$W_2(X), R_1(Y), R_2(Y), R_1(X), Commit_1, Commit_2$$

Schedule  $S$  is a 2PL schedule since it can be executed as follows:

$$X_2(X), W_2(X), S_1(Y), R_1(Y), S_2(Y), R_2(Y), U_2(X), S_1(X), R_1(X), U_1(Y), U_1(X), Commit_1, U_2(Y), Commit_2$$

Note that  $lastlocktime(T_2) < lastlocktime(T_1)$  since  $T_2$ 's last acquired lock (i.e.,  $S_2(Y)$ ) precedes  $T_1$ 's (i.e.,  $S_1(X)$ ).

Given a 2PL schedule  $S$ , let  $S'$  denote the serial schedule over the same set of transactions in  $S$ , where the transactions in  $S'$  are serialized in increasing order of their  $lastlocktime()$  in  $S$ . Prove that  $S$  and  $S'$  are conflict equivalent.

**Solution:** Consider a pair of conflicting actions  $(A_i(O), A_j(O))$  in  $S$  on object  $O$ ,  $i \neq j$ , where  $A_i(O)$  precedes  $A_j(O)$  in  $S$ . Since the actions are conflicting,  $T_i$  must release its lock on  $O$  before  $A_j(O)$  in  $S$ . Furthermore,  $T_i$  does not acquire any other lock in  $S$  after releasing its lock on  $O$  before  $A_j(O)$ . Therefore,  $lastlocktime(T_i) < lastlocktime(T_j)$ , and  $T_i$  precedes  $T_j$  in  $S'$ . Hence,  $A_i(O)$  must also precede  $A_j(O)$  in  $S'$ . Therefore,  $S$  and  $S'$  are conflict equivalent.