

Questions to be discussed: 1, 2 & 4.

1. (a) Write down a serializable Snapshot Isolation (SSI) schedule that is not a conflict serializable schedule (CSS).
- (b) Write down a conflict serializable schedule (CSS) that is not a serializable Snapshot Isolation (SSI) schedule.
- (c) Write down a serializable Snapshot Isolation (SSI) schedule S (involving transactions T_1 , T_2 , and T_3) such that $DSG(S)$ consists of two concurrent rw-dependency edges $T_1 \xrightarrow{rw} T_2$ and $T_2 \xrightarrow{rw} T_3$.

Solution:

- (a) The following schedule is a SSI schedule that is multiversion view equivalent to (T_2, T_1) . However, it is not a valid monoversion schedule and therefore not a CSS.

$$\begin{array}{llll} T_1: & W_1(x_1), & & W_1(y_1), C_1, \\ T_2: & & R_2(x_0), R_2(y_0), & C_2 \end{array}$$

- (b) The following schedule is CSS schedule that is conflict equivalent to (T_1, T_2) . However, under SI, it is impossible for both of the transactions to have committed due to their concurrent updates to y . Specifically, T_2 would have been aborted (either when $W_2(y)$ is attempted under FUW or when C_2 is attempted under FCW).

$$\begin{array}{llll} T_1: & R_1(x), & W_1(y), & C_1, \\ T_2: & & R_2(x), & W_2(y), C_2 \end{array}$$

- (c) The following schedule is a SSI schedule that is multiversion view equivalent to (T_1, T_2, T_3) . Observe that $T_1 \xrightarrow{rw} T_2$ (due to $R_1(x_0)$ and $W_2(x_2)$), and $T_2 \xrightarrow{rw} T_3$ (due to $R_2(y_0)$ and $W_3(y_3)$),

$$\begin{array}{llllll} T_1: & R_1(x), & & & C_1, & \\ T_2: & & W_2(x), & R_2(y), & & C_2, \\ T_3: & & & & W_3(y), & C_3 \end{array}$$

2. For each of the following multiversion schedules, state whether it is multiversion view serializable (MVSS). If the schedule is MVSS, write down a serial monoversion schedule that is multiversion view equivalent to it.

- (a)
- | | | | | | | | | | |
|---------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| T_1 : | | $R_1(x_4)$ | | $W_1(x_1)$ | | | $R_1(y_0)$ | $W_1(z_1)$ | |
| T_2 : | | | $R_2(y_0)$ | | | | $R_2(x_0)$ | | $R_2(z_3)$ |
| T_3 : | | | | | $W_3(x_3)$ | $W_3(z_3)$ | | | |
| T_4 : | $W_4(y_4)$ | $W_4(x_4)$ | | | | | | | |
- (b)
- | | | | | | | | | | |
|---------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| T_1 : | | $R_1(x_4)$ | | $W_1(x_1)$ | | | $R_1(y_4)$ | $W_1(z_1)$ | |
| T_2 : | | | $R_2(y_0)$ | | | | $R_2(x_1)$ | | $R_2(z_3)$ |
| T_3 : | | | | | $W_3(x_3)$ | $W_3(z_3)$ | | | |
| T_4 : | $W_4(y_4)$ | $W_4(x_4)$ | | | | | | | |
- (c)
- | | | | | | | | | | |
|---------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| T_1 : | | $R_1(x_4)$ | | $W_1(x_1)$ | | | $R_1(y_4)$ | $W_1(z_1)$ | |
| T_2 : | | | $R_2(y_0)$ | | | | $R_2(x_3)$ | | $R_2(z_3)$ |
| T_3 : | | | | | $W_3(x_3)$ | $W_3(z_3)$ | | | |
| T_4 : | $W_4(y_4)$ | $W_4(x_4)$ | | | | | | | |
- (d)
- | | | | | | | | | | |
|---------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| T_1 : | | $R_1(x_3)$ | | $W_1(x_1)$ | | | $R_1(y_0)$ | $W_1(z_1)$ | |
| T_2 : | | | $R_2(y_4)$ | | | | $R_2(x_4)$ | | $R_2(z_1)$ |
| T_3 : | | | | | $W_3(x_3)$ | $W_3(z_3)$ | | | |
| T_4 : | $W_4(y_4)$ | $W_4(x_4)$ | | | | | | | |

Solution:

To determine whether a schedule S is MVSS, we can construct a directed graph (denoted by $MVSG(S)$) to capture the read-from relations among the transactions. Each node in $MVSG(S)$ represents a transaction with the following edges: (1) (T_j, T_i) if T_i reads from T_j , and (2) (T_j, T_i) if T_j read some object O from the initial database & T_i update object O . If $MVSG(S)$ is cyclic, then S is not MVSS. If $MVSG(S)$ is acyclic, then S is MVSS iff there exists a serial monoversion schedule produced from a topological ordering of $MVSG(S)$ that is multiversion view equivalent to S .

- (a) Not MVSS. Since $R_1(x_4)$, we have $T_4 \rightarrow T_1$ in $MVSG(S)$. Since $R_1(y_0)$ & $W_4(y_4)$, we have $T_1 \rightarrow T_4$ in $MVSG(S)$. Since $MVSG(S)$ is cyclic, S is not MVSS.
- (b) Not MVSS. Since $R_1(x_4)$, we have $T_4 \rightarrow T_1$ in $MVSG(S)$. Since $R_2(y_0)$ & $W_4(y_4)$, we have $T_2 \rightarrow T_4$ in $MVSG(S)$. Since $R_2(x_1)$, we have $T_1 \rightarrow T_2$ in $MVSG(S)$. Since $MVSG(S)$ is cyclic, S is not MVSS.
- (c) MVSS. Since $R_1(x_4)$, we have $T_4 \rightarrow T_1$ in $MVSG(S)$. Since $R_2(y_0)$ & $W_4(y_4)$, we have $T_2 \rightarrow T_4$ in $MVSG(S)$. Since $R_2(x_3)$, we have $T_3 \rightarrow T_2$ in $MVSG(S)$. The schedule is multiversion view equivalent to serial monoversion schedule (T_3, T_2, T_4, T_1) .
- (d) MVSS. Since $R_1(x_3)$, we have $T_3 \rightarrow T_1$ in $MVSG(S)$. Since $R_1(y_0)$ & $W_4(y_4)$, we have $T_1 \rightarrow T_4$ in $MVSG(S)$. Since $R_2(y_4)$, we have $T_4 \rightarrow T_2$ in $MVSG(S)$. Since $R_2(z_1)$, we have $T_1 \rightarrow T_2$ in $MVSG(S)$. The schedule is multiversion view equivalent to serial monoversion schedule (T_3, T_1, T_4, T_2) .

3. (Exercise 17.3, R&G) For each of the following schedules, state whether it is a Snapshot Isolation schedule under the First-Committer-Wins rule.
- (a) $R_1(X), R_2(X), W_1(X), W_2(X), Commit_1, Commit_2$
 - (b) $W_1(X), R_2(Y), R_1(Y), R_2(X), Commit_1, Commit_2$
 - (c) $R_1(X), R_2(Y), W_3(X), R_2(X), R_1(Y), Commit_1, Commit_2, Commit_3$
 - (d) $R_1(X), R_1(Y), W_1(X), R_2(Y), W_3(Y), W_1(X), R_2(Y), Commit_1, Commit_2, Commit_3$
 - (e) $R_1(X), W_2(X), W_1(X), Commit_2, Commit_1$
 - (f) $W_1(X), R_2(X), W_1(X), Commit_2, Commit_1$
 - (g) $R_2(X), W_3(X), Commit_3, W_1(Y), Commit_1, R_2(Y), W_2(Z), Commit_2$
 - (h) $R_1(X), W_2(X), Commit_2, W_1(X), Commit_1, R_3(X), Commit_3$
 - (i) $R_1(X), W_2(X), W_1(X), R_3(X), Commit_1, Commit_2, Commit_3$

Solution:

- (a) Not SI-FCW: T_2 must abort due to $W_2(X)$.
- (b) SI-FCW.
- (c) SI-FCW.
- (d) SI-FCW.
- (e) Not SI-FCW: T_1 must abort due to $W_1(X)$.
- (f) SI-FCW.
- (g) SI-FCW.
- (h) Not SI-FCW: T_1 must abort due to $W_1(X)$.
- (i) Not SI-FCW: T_2 must abort due to $W_2(X)$.

4. Consider the following schedule involving transactions T_1 , T_2 , T_3 , and T_4 which are executed in a database system using the Snapshot Isolation protocol with the First Committer Wins Rule.

Assume the following values for the state of the database before the start of the schedule:

$$a_0 = 10, \quad b_0 = 20, \quad c_0 = 30.$$

Timestamp	T_3	T_1	T_2	T_4	Comments
1	$R_3(c)$				
2			$R_2(b)$		
3			$W_2(b)$		T_2 updates b to 45.
4			$Commit_2$		
5	$R_3(b)$				
6	$W_3(b)$				T_3 updates b to 60.
7		$R_1(a)$			
8		$R_1(b)$			
9		$W_1(a)$			T_1 updates a to 100.
10				$R_4(b)$	
11				$R_4(a)$	
12				$W_4(b)$	T_4 updates b to 16.
13	$R_3(a)$				
14	$W_3(c)$				T_3 updates c to 200.
15		$W_1(c)$			T_1 updates c to 40.
16		$Commit_1$			
17				$R_4(b)$	
18				$R_4(c)$	

- Write down the value read by $R_3(b)$ at timestamp 5.
- Write down the value read by $R_1(b)$ at timestamp 8.
- Write down the value read by $R_4(b)$ at timestamp 10.
- Write down the value read by $R_4(a)$ at timestamp 11.
- Write down the value read by $R_3(a)$ at timestamp 13.
- Write down the value read by $R_4(b)$ at timestamp 17.
- Write down the value read by $R_4(c)$ at timestamp 18.

Solution: The database snapshot for a transaction T consists of (1) the versions written by transactions that have committed before the start of T , and (2) any versions written by T itself. Thus, T can't read any version written by a concurrent transaction. For each $R_i(O)$, the most recent version of O in T_i 's database snapshot is returned.

- Write down the value read by $R_3(b)$ at timestamp 5. (20)
- Write down the value read by $R_1(b)$ at timestamp 8. (45)
- Write down the value read by $R_4(b)$ at timestamp 10. (45)
- Write down the value read by $R_4(a)$ at timestamp 11. (10)
- Write down the value read by $R_3(a)$ at timestamp 13. (10)
- Write down the value read by $R_4(b)$ at timestamp 17. (16)
- Write down the value read by $R_4(c)$ at timestamp 18. (30)