

CS3223 Lecture 10

Crash Recovery

Recovery

- ▶ **Recovery manager** guarantees atomicity and durability properties of Xacts
 - ▶ **Undo**: remove effects of **aborted Xact** to preserve **atomicity**
 - ▶ **Redo**: re-installing effects of **committed Xact** for **durability**
- ▶ Types of failure
 - ▶ **Transaction failures**: transaction aborts
 - ★ Application rollbacks transaction
 - ★ System rollbacks transaction (e.g. deadlock, violation of integrity constraint)
 - ▶ **System crash**: loss of volatile memory contents
 - ★ Power failure
 - ★ Bug in DBMS/OS
 - ★ Hardware malfunction
 - ▶ **Media failures**: data is lost/corrupted on non-volatile storage
 - ★ Disk head crash / Failure during data transfer

Recovery Manager

Processes three operations:

- ▶ **Commit(T)** - install T's updated pages into database
- ▶ **Abort(T)** - restore all data that T updated to their prior values
- ▶ **Restart** - recover database to a consistent state from system failure
 - ▶ abort all active Xacts at the time of system failure
 - ▶ installs updates of all committed Xacts that were not installed in the database before the failure

Desirable properties:

- ▶ Add little overhead to the normal processing of Xacts
- ▶ Recover quickly from a failure

Interaction of Recovery & Buffer Managers

Two issues with dirty pages in buffer pool:

1. **Can a dirty page updated by Xact T be written to disk before T commits?**
 - ▶ Yes \implies steal policy
 - ▶ No \implies no-steal policy
2. **Must all dirty pages that are updated by Xact T be written to disk when T commits?**
 - ▶ Yes \implies force policy
 - ▶ No \implies no-force policy

Interaction of Recovery & Buffer Managers (cont.)

- ▶ **Steal policy**: allows dirty pages updated by Xact T to be replaced from buffer pool before T commits
- ▶ **Force policy**: requires all dirty pages updated by Xact T to be written to disk when T commits
- ▶ Four possible design options:

	Force	No-force
Steal	undo & no redo	undo & redo
No-steal	no undo & no redo	no undo & redo

- ▶ No-steal policy \implies no undo
- ▶ Force policy \implies no redo

Example: Crash Recovery

Accounts

name	balance
Alice	200
Bob	800
Carol	300
Dave	500
Eve	600
Fred	200

Transactions:

T_1 : Transfer \$100 from Alice to Carol

T_2 : Deposit \$200 to Bob

T_3 : Withdraw \$500 from Eve

T_1 : $R_1(\text{Alice}, 200)$, $R_1(\text{Carol}, 300)$, $W_1(\text{Alice}, 100)$, $W_1(\text{Carol}, 400)$

T_2 : $R_2(\text{Bob}, 800)$, $W_2(\text{Bob}, 1000)$

T_3 : $R_3(\text{Eve}, 600)$, $W_3(\text{Eve}, 100)$


Example: Crash Recovery (cont.)

Schedule


Step	T ₁	T ₂	T ₃
1	R_1 (Alice, 200)		
2	R_1 (Carol, 300)		
3	W_1 (Alice, 100)		
4		R_2 (Bob, 800)	
5		W_2 (Bob, 1000)	
6		$Commit_2$	
7	W_1 (Carol, 400)		
8			R_3 (Eve, 600)
9			W_3 (Eve, 100)
10			$Commit_3$
11		SYSTEM CRASH!	

Example: Crash Recovery (cont.)

Database



Buffer Pool

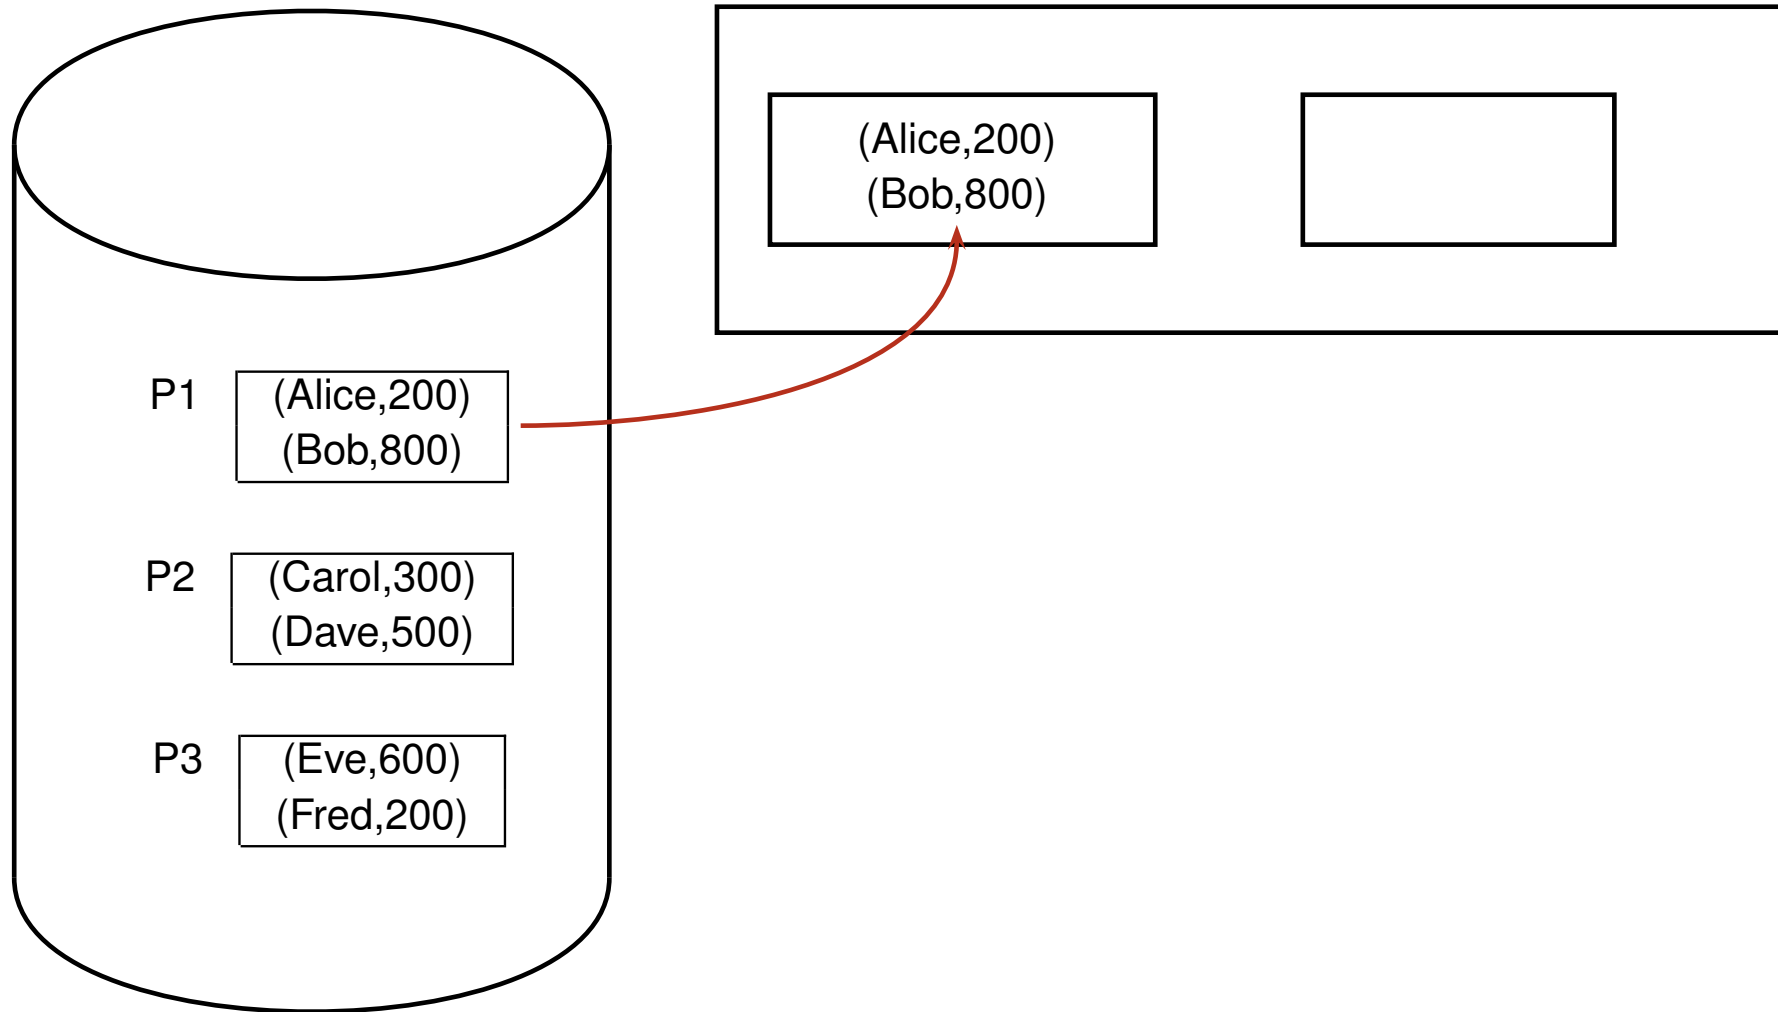


Initial State

Example: Crash Recovery (cont.)

Database

Buffer Pool

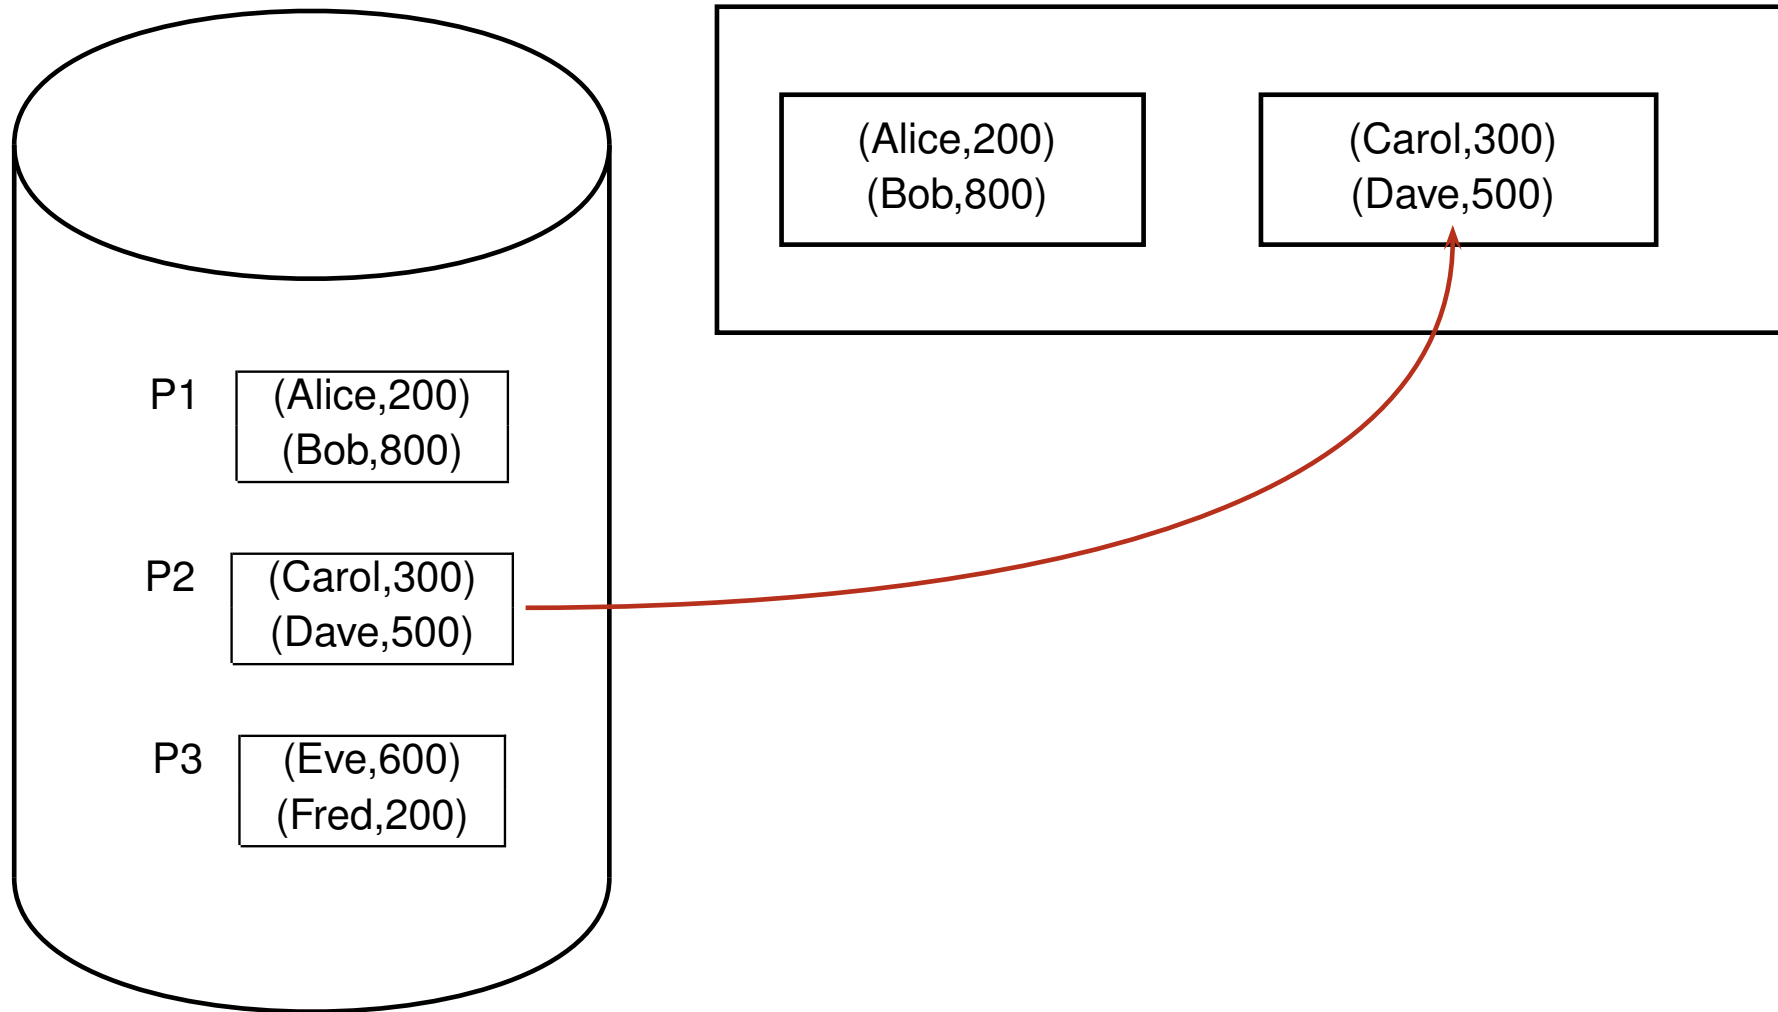


R_1 (Alice,200): P1 is read into buffer pool

Example: Crash Recovery (cont.)

Database

Buffer Pool





R_1 (Carol, 300): P2 is read into buffer pool

Example: Crash Recovery (cont.)

Database

Buffer Pool





W_1 (Alice, 100)

Example: Crash Recovery (cont.)

Database

Buffer Pool


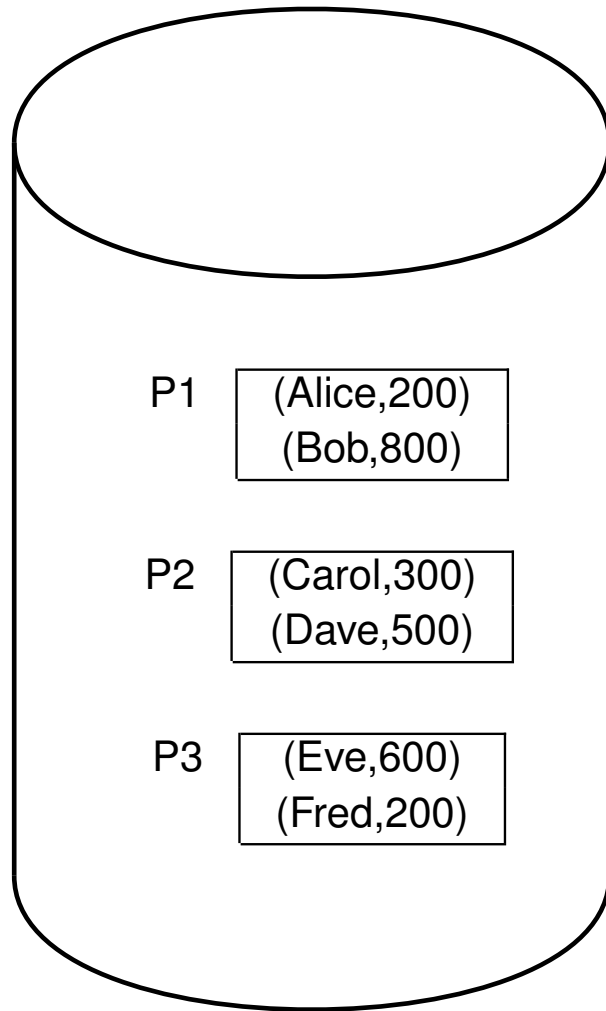


$R_2(\text{Bob}, 800)$

Example: Crash Recovery (cont.)

Database

Buffer Pool





$W_2(\text{Bob}, 1000), \text{Commit}_2$

Example: Crash Recovery (cont.)

Database

Buffer Pool





W_1 (Carol,400)

Example: Crash Recovery (cont.)

Database

Buffer Pool

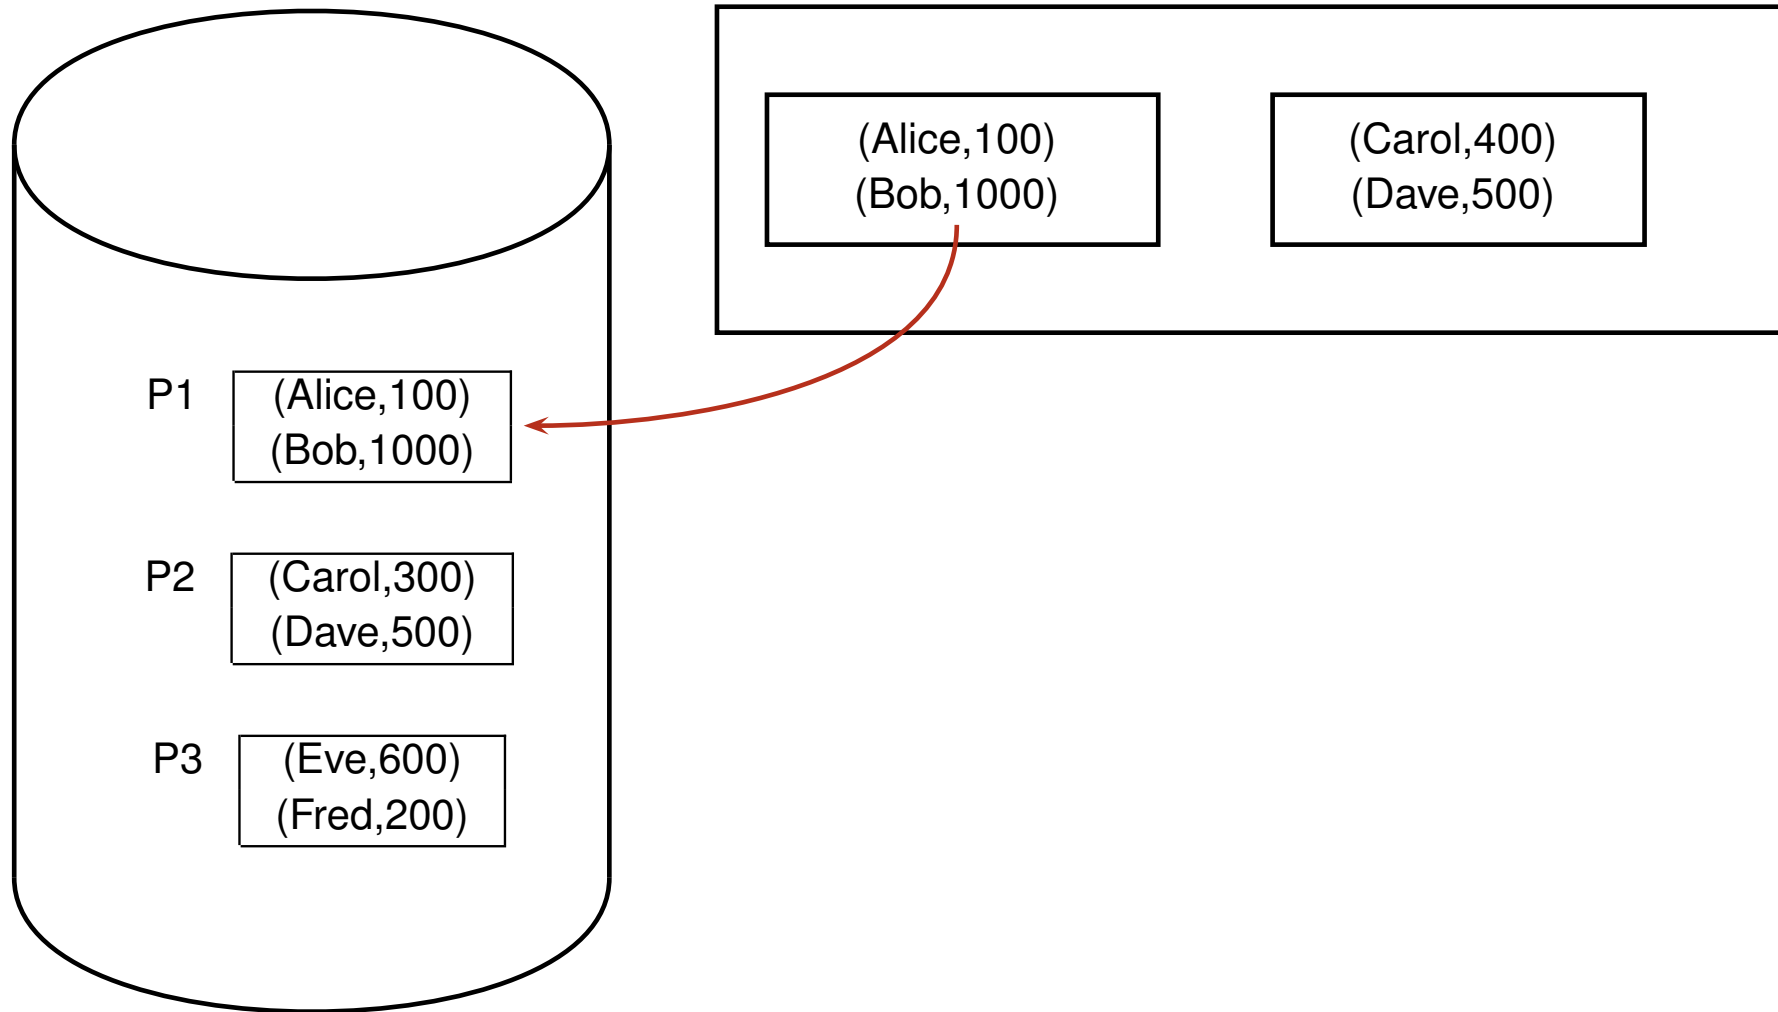


$R_3(\text{Eve},600)$: buffer pool manager picks P1 for replacement

Example: Crash Recovery (cont.)

Database

Buffer Pool

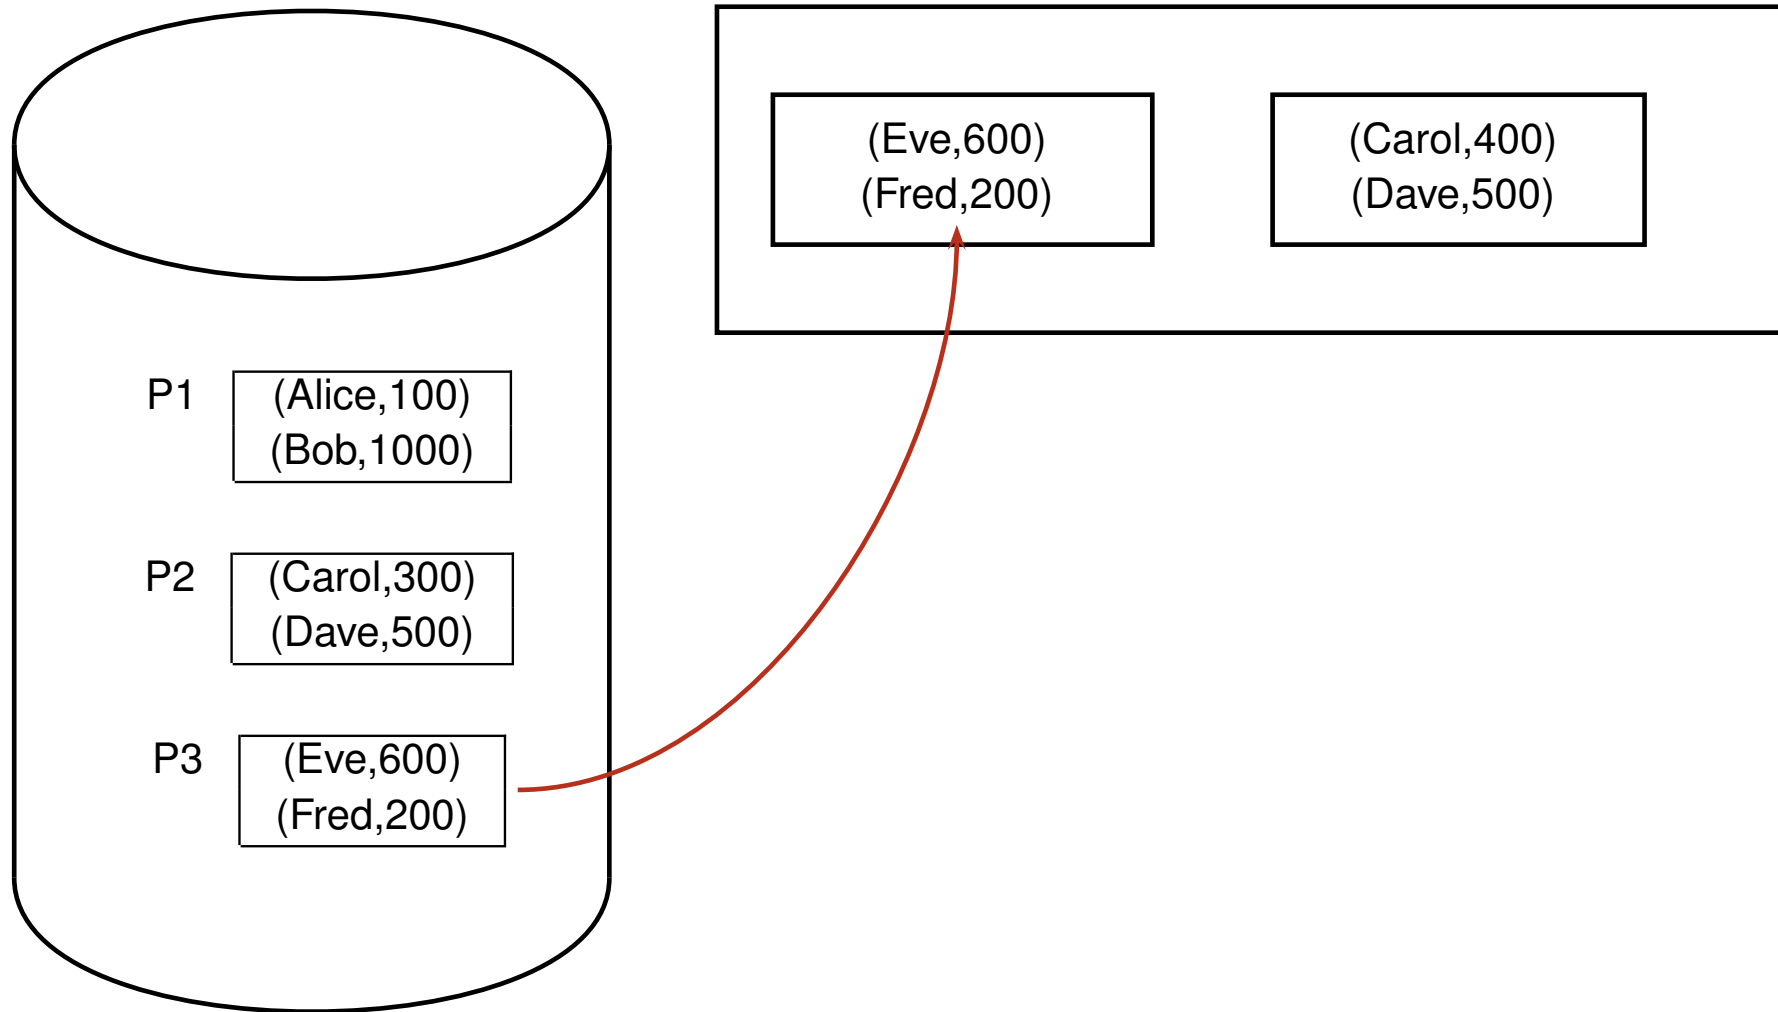


$R_3(\text{Eve},600)$: P1 is written to disk

Example: Crash Recovery (cont.)

Database

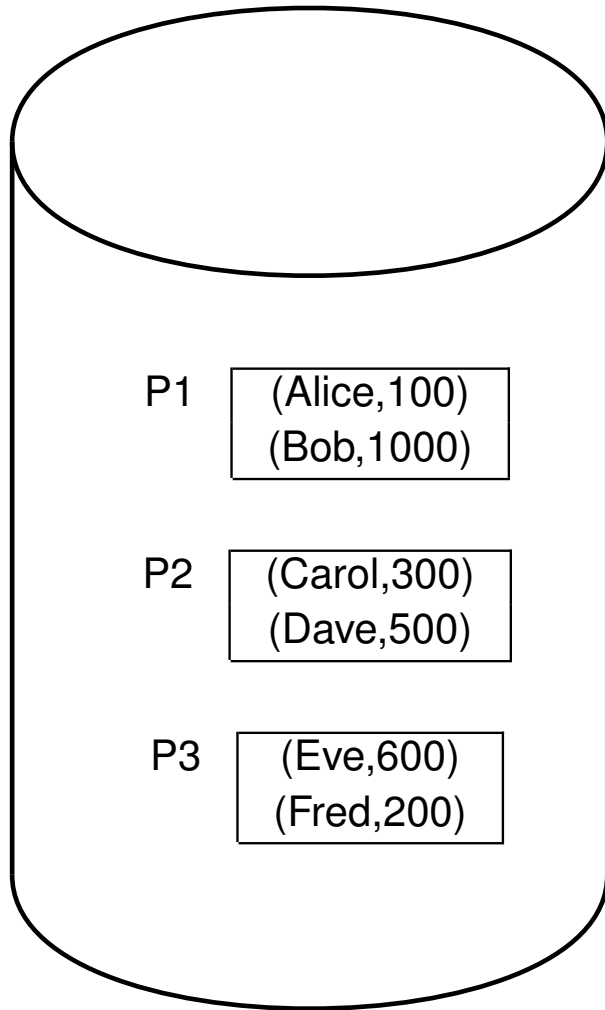
Buffer Pool




$R_3(\text{Eve}, 600)$: P3 is read into buffer pool

Example: Crash Recovery (cont.)

Database

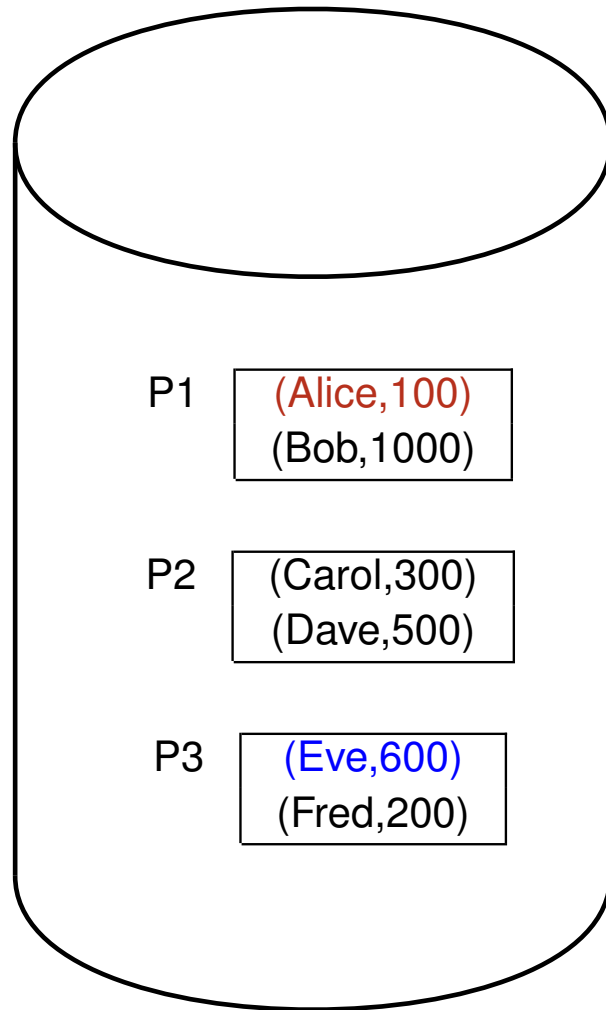


Buffer Pool



$W_3(\text{Eve}, 100), \text{Commit}_3$

Example: Crash Recovery (cont.)



Step	T ₁	T ₂	T ₃
1	$R_1(\text{Alice}, 200)$		
2	$R_1(\text{Carol}, 300)$		
3	$W_1(\text{Alice}, 100)$		
4		$R_2(\text{Bob}, 800)$	
5		$W_2(\text{Bob}, 1000)$	
6		Commit_2	
7	$W_1(\text{Carol}, 400)$		
8			$R_3(\text{Eve}, 600)$
9			$W_3(\text{Eve}, 100)$
10			Commit_3
11		SYSTEM CRASH!	

- ▶ Need to abort T_1 : undo $W_1(\text{Alice}, 100) \longrightarrow (\text{Alice}, 200)$
- ▶ Need to redo $W_3(\text{Eve}, 600) \longrightarrow (\text{Eve}, 100)$

Log-based Database Recovery

- ▶ **Log** (aka trail/journal): history of actions executed by DBMS
 - ▶ Contains a log record for each write, commit, & abort
- ▶ Log is stored as a sequential file of records in stable storage
- ▶ Log is stored in **stable storage** (storage that can survive crashes & media failures)
 - ▶ Implemented by maintaining multiple copies of information (possibly at different locations) on non-volatile storage devices
- ▶ Each log record has a unique identifier called **Log Sequence Number (LSN)**
 - ▶ LSN of an earlier log record is smaller than that of a later log record

ARIES Recovery Algorithm

- ▶ Algorithm for **R**ecovery and **I**solating **E**xploiting **S**emantics
- ▶ Designed to work with a **steal, no-force** approach
- ▶ Assumes strict 2PL for concurrency control
- ▶ ARIES is invented by IBM and used in several products (e.g., DB2, Microsoft SQL Server)

Recovery-related Structures

- ▶ Log file
- ▶ Transaction table (TT)
 - ▶ One entry for each **active Xact**
 - ▶ Each entry contains:
 - ★ XactID: Transaction identifier
 - ★ **lastLSN**: LSN of the most recent log record for this Xact
 - ★ Xact status (C or U)
 - C = Xact has committed
 - U = Xact has not committed
- ▶ Dirty page table (DPT)
 - ▶ One entry for each **dirty page** in buffer pool
 - ▶ Each entry contains:
 - ★ pageID: page ID of dirty page
 - ★ **recLSN**: LSN of the earliest log record for an update that caused the page to be dirty
- ▶ All these structures are updated during normal processing

Log records

Information in log records

- ▶ type of log record (e.g., update, commit, abort)
- ▶ identifier of Xact
- ▶ **prevLSN**: LSN of the previous log record for the same Xact
- ▶ other type-specific information

Update log record

- ▶ identifier of page being updated
- ▶ before-image of update
- ▶ after-image of update

Implementing Abort

- ▶ Undo all updates by Xact to database pages
- ▶ **Write-ahead logging (WAL) protocol**
Do not flush an uncommitted update to the database until the log record containing its before-image has been flushed to the log
- ▶ How to enforce WAL protocol?
 - ▶ Each database page contains the LSN of the most recent log record (a.k.a. **pageLSN**) that describes an update to this page
 - ▶ Before flushing a **database page P** to disk, ensure that all the log records up to the log record corresponding to **P's pageLSN** have been flushed to disk

Implementing Abort (cont.)

- ▶ How to undo all the updates of Xact?
 - ▶ For each log record of Xact in reverse order, restore log record's before-image
- ▶ How to efficiently retrieve Xact's log records in reverse order?
 - ▶ **Xact Table (TT)** - maintains one entry for each active Xact
 - ▶ Each TT entry stores the LSN of most recent log record for Xact (a.k.a. **lastLSN**)
 - ▶ Use **lastLSN** to retrieve the most recent log record for Xact; the other log records for Xact are retrieved via the **prevLSN** of each retrieved log record
- ▶ **Logging Changes During Undo**: Changes made to database while undoing a Xact are also logged to ensure that an action is not repeated in the event of repeated undos

Implementing Commit

- ▶ Need to ensure that all the updates of Xact must be in stable storage (database or log) before Xact is committed
- ▶ **Force-at-commit protocol**
Do not commit a Xact until the after-images of all its updated records are in stable storage (database or log)
- ▶ How to enforce force-at-commit protocol?
 - ▶ Write a commit log record for Xact
 - ▶ Flush all the log records for Xact to disk
- ▶ Xact is considered to have committed if its commit log record has been written to stable storage

Implementing Restart

- ▶ Recovery from system crashes consists of three phases:
 1. **Analysis phase**: identifies dirty buffer pool pages & active Xacts at time of crash
 2. **Redo phase**: redo actions to restore database state to what it was at the time of crash
 3. **Undo phase**: undo actions of Xacts that did not commit
- ▶ **Repeating History During Redo**: During restart following a crash, first restore system to the state before crash, and then undo the actions of Xacts that are active at the time of crash

Normal Transaction Processing

Updating Xact Table (transID, lastLSN, status)

- ▶ When the first log record is created for Xact T , create a new entry for T with status = U
- ▶ When a new log record r is created for Xact T , update `lastLSN` for T 's entry to be r 's LSN
- ▶ If Xact T commits, update `status` for T 's entry to be C
- ▶ When an end log record is generated for Xact T , remove T 's entry

Updating Dirty Page Table (pageID, recLSN)

- ▶ When a page P in buffer pool is updated & DPT has no entry for P , create a new table entry for P with `recLSN` = LSN of log record corresponding to update
- ▶ When a dirtied page P in buffer pool is flushed to disk, remove entry for P

Normal Xact Processing: Example

LOG

10	prevLSN	XactID	type	pageID	length	offset	before image	after image
	-	T_1	update	P500	3	21	ABC	DEF

DIRTY PAGE TABLE

pageID	recLSN
P500	10

XACT TABLE

XactID	lastLSN	status
T_1	10	U

Normal Xact Processing: Example

LOG

	prevLSN	XactID	type	pageID	length	offset	before image	after image
10	-	T_1	update	P500	3	21	ABC	DEF
20	-	T_2	update	P600	3	41	HIJ	KLM

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20

XACT TABLE

XactID	lastLSN	status
T_1	10	U
T_2	20	U

Normal Xact Processing: Example

LOG

	prevLSN	XactID	type	pageID	length	offset	before image	after image
10	-	T_1	update	P500	3	21	ABC	DEF
20	-	T_2	update	P600	3	41	HIJ	KLM
30	20	T_2	update	P500	3	20	GDE	QRS

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20

XACT TABLE

XactID	lastLSN	status
T_1	10	U
T_2	30	U

Normal Xact Processing: Example

LOG

	prevLSN	XactID	type	pageID	length	offset	before image	after image
10	-	T_1	update	P500	3	21	ABC	DEF
20	-	T_2	update	P600	3	41	HIJ	KLM
30	20	T_2	update	P500	3	20	GDE	QRS
40	10	T_1	update	P505	3	21	TUV	WXY

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

XACT TABLE

XactID	lastLSN	status
T_1	40	U
T_2	30	U

Types of Log Records

- ▶ All log records have the following information:
 - (1) type of log record (e.g., update, commit, abort),
 - (2) identifier of Xact, and
 - (3) prevLSN (LSN of the previous log record for the same Xact)
- ▶ **Update log record**
 - ▶ After updating a page P, create an update log record r
 - ▶ Update pageLSN of P = LSN of r
 - ▶ Additional fields in update log record:
 - ★ **pageID**: identifier of page being updated
 - ★ **offset**: byte offset within page indicating the beginning of updated portion
 - ★ **length**: number of bytes for updated portion of data page
 - ★ **before-image**: value of the changed bytes before update
 - ★ **after-image**: value of the changed bytes after update

Types of Log Records (cont.)

► Compensation log record (CLR)

- When the update described by an update log record (ULR) is undone, create a compensation log record
- Additional fields in CLR:
 - ★ page ID
 - ★ `undoNextLSN` = LSN of next log record to be undone (i.e., `prevLSN` in ULR)
 - ★ action taken to undo update

► Commit log record

- When a Xact is to be committed, create a commit log record `r`
- All log records (up to and including `r`) are force-written to stable storage
- Xact is considered committed once `r` has been written to stable storage

Types of Log Records (cont.)

▶ Abort log record

- ▶ When a Xact is to be aborted, create an abort log record
- ▶ Undo is initiated for this Xact

▶ End log record

- ▶ Once the additional follow-up processing initiated by a aborted/committed Xact has completed, create an end log record

▶ Checkpoint log record

- ▶ To be discussed later

▶ Update log records & CLRs are classified as redoable log records

Analysis Phase

- ▶ Performs three tasks:
 1. Determines the point in the log to start the Redo phase
 2. Determines the superset of buffer pages that were dirty at the time of crash
 3. Identifies active Xacts at the time of crash

Analysis Phase (cont.)

- ▶ Initialize **dirty page table** (DPT) & **Xact table** (TT) to be empty
- ▶ Scan the log in forward direction to process each log record r (for Xact T):
 - ▶ If r is an **end log record**
 - ★ Remove T from TT
 - Else
 - ★ Add an entry in TT for T if T is not in TT
 - ★ Update **lastLSN** of entry to be r 's LSN
 - ★ Update **status** of entry to C if r is a commit log record
 - ▶ If (r is a redoable log record for page P) & (P is not in DPT), then
 - ★ Create an entry for P in DPT with `pageID of entry = P's pageID` and `recLSN of entry = r's LSN`
- ▶ At the end of **Analysis phase**:
 - ▶ **Xact table** = list of all active Xacts (with status = U) at time of crash
 - ▶ **dirty page table** = superset of dirty pages at time of crash

Analysis Phase: Example

LOG BUFFER

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U
T_2	30	U

Analysis Phase: Example

LOG BUFFER

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U
T_2	50	C

Analysis Phase: Example

LOG BUFFER

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

Analysis Phase: Example

LOG BUFFER

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2
70	update: (40, T_1 , P700, 3, 51, PQR, IJK)

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40
P700	70

TRANSACTION TABLE

XactID	lastLSN	status
T_1	70	U

Analysis Phase: Example

LOG BUFFER

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2
70	update: (40, T_1 , P700, 3, 51, PQR, IJK)

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40
P700	70

TRANSACTION TABLE

XactID	lastLSN	status
T_1	70	U

- Suppose that the system crashes & all the log records (except for the last record) have been flushed to disk before the crash

Analysis Phase: Example (cont.)

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE


pageID	recLSN

TRANSACTION TABLE

XactID	lastLSN	status

Analysis Phase: Example (cont.)

LOG



10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10

TRANSACTION TABLE

XactID	lastLSN	status
T_1	10	U

Analysis Phase: Example (cont.)

LOG

→	10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
	20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
	30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
	40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
	50	commit: T_2
	60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20

TRANSACTION TABLE

XactID	lastLSN	status
T_1	10	U
T_2	20	U

Analysis Phase: Example (cont.)

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20

TRANSACTION TABLE

XactID	lastLSN	status
T_1	10	U
T_2	30	U

Analysis Phase: Example (cont.)

LOG

→ 10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U
T_2	30	U

Analysis Phase: Example (cont.)

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U
T_2	50	C

Analysis Phase: Example (cont.)

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

Analysis Phase: Example (cont.)

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

- ▶ Active Xacts: T_1
- ▶ Dirty Pages: P500, P600, P505

Redo Phase

- ▶ **RedoLSN** = smallest recLSN among all dirty pages in DPT
- ▶ Let r be the log record with LSN = RedoLSN
- ▶ Scan the log in forward direction starting from r
- ▶ If (r is an update log record) or (r is a CLR) then
 - ▶ fetch page P that is associated with r
 - ▶ If (P 's pageLSN $<$ r 's LSN) then
 - ★ Reapply logged action in r to P
 - ★ update P 's pageLSN = r 's LSN
- ▶ At the end of Redo Phase,
 - ▶ Create end log records for Xacts with status = C in Xact Table & remove their entries from Xact Table
 - ▶ System is restored to state at time of crash

Redo Phase: Example

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

► $\text{RedoLSN} = \min\{10, 20, 40\} = 10$

Redo Phase: Example

LOG

→	10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
	20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
	30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
	40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
	50	commit: T_2
	60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

- ▶ RedoLSN = $\min\{10, 20, 40\} = 10$
- ▶ log record r with LSN = 10: **Update: (-, T_1 , P500, 3, 21, ABC, DEF)**
 - Fetch page P500. Suppose P500's pageLSN = 10
 - r's LSN = P500's pageLSN: No need to redo update on P500

Redo Phase: Example

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

- ▶ RedoLSN = $\min\{10, 20, 40\} = 10$
- ▶ log record r with LSN = 20: **Update: (-, T_2 , P600, 3, 41, HIJ, KLM)**
 - ▶ Fetch page P600. Suppose P600's pageLSN = 20
 - ▶ r's LSN = P600's pageLSN: No need to redo update on P600

Redo Phase: Example

LOG

→ 10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

- ▶ RedoLSN = $\min\{10, 20, 40\} = 10$
- ▶ log record r with LSN = 30: **Update: (20, T_2 , P500, 3, 30, GDE, QRS)**
 - ▶ Fetch page P500. P500's pageLSN = 10
 - ▶ r's LSN > P500's pageLSN: Redo update on P500
 - ▶ Update P500's pageLSN = 30

Redo Phase: Example

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

- ▶ RedoLSN = $\min\{10, 20, 40\} = 10$
- ▶ log record r with LSN = 40: **Update: (10, T_1 , P505, 3, 31, TUV, WXY)**
 - ▶ Fetch page P505. Suppose P505's pageLSN = 40
 - ▶ r's LSN = P505's pageLSN: No need to redo update on P505

Redo Phase: Example

	LOG
10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE	
pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE		
XactID	lastLSN	status
T_1	40	U

► RedoLSN = $\min\{10, 20, 40\} = 10$

Redo Phase: Example

LOG

10	update: (-, T_1 , P500, 3, 21, ABC, DEF)
20	update: (-, T_2 , P600, 3, 41, HIJ, KLM)
30	update: (20, T_2 , P500, 3, 20, GDE, QRS)
40	update: (10, T_1 , P505, 3, 21, TUV, WXY)
50	commit: T_2
60	end: T_2

DIRTY PAGE TABLE

pageID	recLSN
P500	10
P600	20
P505	40

TRANSACTION TABLE

XactID	lastLSN	status
T_1	40	U

- ▶ $\text{RedoLSN} = \min\{10, 20, 40\} = 10$
- ▶ System is restored to state at time of crash

Undo Phase

- ▶ **Goal:** abort active Xacts at time of crash (**loser Xacts**)
 - ▶ Abort loser Xacts by undoing their actions in reverse order
- ▶ Initialize L = set of lastLSNs (with status = U) from TT
- ▶ Repeat until L becomes empty
 - ▶ delete the largest lastLSN from L
 - ▶ let r be the log record corresponding to this lastLSN
 - ▶ if r is an **update log record** for Xact *T* on page *P* then
 - ★ create a CLR r_2 for *T*: r_2 's undoNextLSN = r's prevLSN
 - ★ update *T*'s entry in TT: lastLSN = r_2 's LSN
 - ★ undo the logged action on page *P*
 - ★ update *P*'s pageLSN = r_2 's LSN
 - ★ Update-L-and-TT (r's prevLSN)
 - ▶ else if r is a **CLR** for Xact *T* then
 - ★ Update-L-and-TT (r's undoNextLSN)
 - ▶ else if r is an **abort log record** for Xact *T* then Update-L-and-TT (r's prevLSN)
- ▶ **Update-L-and-TT (lsn)**
 - ▶ if lsn is not null then add lsn to L
 - ▶ else create an end log record for *T* & remove *T*'s entry from TT

Undo Phase: Example

- Recall that at the end of Analysis Phase, T_1 is the only active Xact to be undone (T_1 's lastLSN in TT is 40)

LSN	LOG
10	update: T_1 writes P500
20	update: T_2 writes P600
30	update: T_2 writes P500
40	update: T_1 writes P505
50	commit: T_2
60	end: T_2
	CRASH

Undo Phase: Example

- Recall that at the end of Analysis Phase, T_1 is the only active Xact to be undone (T_1 's lastLSN in TT is 40)

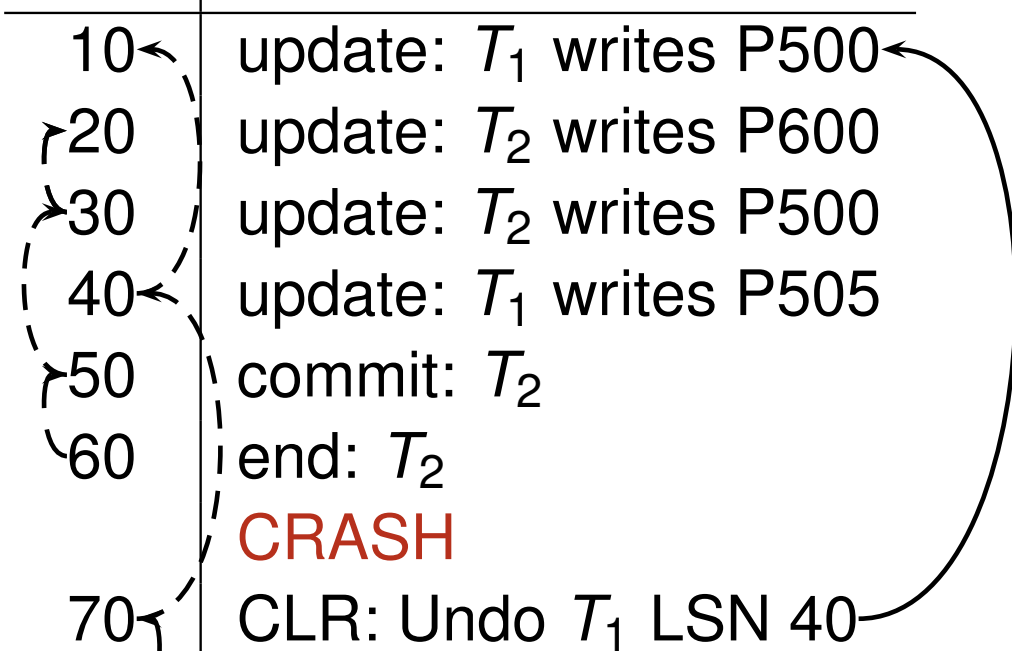
LSN	LOG
10	update: T_1 writes P500
20	update: T_2 writes P600
30	update: T_2 writes P500
40	update: T_1 writes P505
50	commit: T_2
60	end: T_2
	CRASH
70	CLR: Undo T_1 LSN 40

undoNextLSN

Undo Phase: Example

- Recall that at the end of Analysis Phase, T_1 is the only active Xact to be undone (T_1 's lastLSN in TT is 40)

LSN	LOG
10	update: T_1 writes P500
20	update: T_2 writes P600
30	update: T_2 writes P500
40	update: T_1 writes P505
50	commit: T_2
60	end: T_2
	CRASH
70	CLR: Undo T_1 LSN 40
80	CLR: Undo T_1 LSN 10
90	end: T_1



undoNextLSN

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	60	U
T_3	50	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{50, 60\}$

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH
70	CLR: Undo T_2 LSN 60

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	60	U
T_3	50	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{50, 60\}$

- LSN=60: $L = \{20, 50\}$

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH
70	CLR: Undo T_2 LSN 60
80	CLR: Undo T_3 LSN 50
85	end: T_3

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	60	U
T_3	50	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{50, 60\}$

► LSN=60: $L = \{20, 50\}$

► LSN=50: $L = \{20\}$

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH
70	CLR: Undo T_2 LSN 60
80	CLR: Undo T_3 LSN 50
85	end: T_3
	CRASH

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	60	U
T_3	50	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{50, 60\}$

► LSN=60: $L = \{20, 50\}$

► LSN=50: $L = \{20\}$

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH
70	CLR: Undo T_2 LSN 60
80	CLR: Undo T_3 LSN 50
85	end: T_3
	CRASH

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	60	U
T_3	50	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{50, 60\}$

► LSN=60: $L = \{20, 50\}$

► LSN=50: $L = \{20\}$

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	70	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{70\}$

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH
70	CLR: Undo T_2 LSN 60
80	CLR: Undo T_3 LSN 50
85	end: T_3
	CRASH

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	60	U
T_3	50	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{50, 60\}$

► LSN=60: $L = \{20, 50\}$

► LSN=50: $L = \{20\}$

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	70	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{70\}$

► LSN=70: $L = \{20\}$

Undo Phase: Another Example

LSN	LOG
10	update: T_1 writes P5
20	update: T_2 writes P3
30	abort: T_1
40	CLR: Undo T_1 LSN 10
45	end: T_1
50	update: T_3 writes P1
60	update: T_2 writes P5
	CRASH
70	CLR: Undo T_2 LSN 60
80	CLR: Undo T_3 LSN 50
85	end: T_3
	CRASH
90	CLR: Undo T_2 LSN 20
95	end: T_2

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	60	U
T_3	50	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{50, 60\}$

► LSN=60: $L = \{20, 50\}$

► LSN=50: $L = \{20\}$

► Analysis Phase:

pageID	recLSN
P1	50
P3	20
P5	10

XactID	lastLSN	status
T_2	70	U

► Redo Phase: RedoLSN = 10

► Undo Phase: $L = \{70\}$

► LSN=70: $L = \{20\}$

► LSN=20: $L = \{\}$

Checkpointing

- ▶ Perform **checkpoint operations** periodically to speed up restart recovery
- ▶ Checkpointing synchronizes state of log with database state

Simple Checkpointing

1. Stop accepting any new update, commit, & abort operations
 2. Wait till all active update, commit, & abort operations have finished
 3. Flush all dirty pages in buffer
 4. Write a **checkpoint log record** containing **Xact table**
 5. Resume accepting new update, commit, & abort operations
- ▶ During restart recovery, Analysis Phase begins with the latest **checkpoint log record** (CPLR)
 - ▶ Initialize **Xact table** with CPLR's Xact table
 - ▶ Initialize **dirty page table** to be empty

Fuzzy Checkpointing in ARIES

1. Let **DPT'** be the **dirty page table** & **TT'** be the **Xact table**
2. Write a **begin_checkpoint log record**
3. Write a **end_checkpoint log record** containing **DPT'** & **TT'**
4. Write a special **master record** containing the LSN of the **begin_checkpoint log record** to a known place on stable storage

- ▶ During restart recovery, Analysis Phase starts with the **begin_checkpoint log record** (BCPLR) identified by the **master record**
 - ▶ Let ECPLR denote the **end_checkpoint log record** (ECPLR) corresponding to BCPLR
 - ▶ Assume that there are no log records between BCPLR & ECPLR
 - ▶ Initialize **Xact table** with ~~BCPLR's~~ ECPLR's Xact table
 - ▶ Initialize **dirty page table** with ~~BCPLR's~~ ECPLR's dirty page table

Revisiting Redo Phase

- ▶ **RedoLSN** = smallest recLSN among all dirty pages in DPT
- ▶ Let r be the log record with LSN = RedoLSN
- ▶ Scan the log in forward direction starting from r
- ▶ If (r is an update log record) or (r is a CLR) then
 - ▶ fetch page P that is associated with r
 - ▶ If (P 's pageLSN $<$ r 's LSN) then
 - ★ Reapply logged action in r to P
 - ★ update P 's pageLSN = r 's LSN
- ▶ At the end of Redo Phase,
 - ▶ Create end log records for Xacts with status = C in Xact Table & remove their entries from Xact Table
 - ▶ System is restored to state at time of crash

Revisiting Redo Phase: Optimization

- ▶ Goal: Exploit information in DPT to avoid retrieving P
- ▶ **Optimization condition:**
(P is not in DPT) or (P 's recLSN in DPT $>$ r 's LSN)
- ▶ If optimization condition holds, then
 - ▶ the update of r has already been applied to P
 - ▶ r can be ignored & no need to fetch P !
- ▶ If (r is redoable) and (optimization condition does not hold) then
 - ▶ fetch page P that is associated with r
 - ▶ If (P 's pageLSN $<$ r 's LSN) then
 - ★ Reapply logged action in r to P
 - ★ update P 's pageLSN = r 's LSN
 - Else
 - ★ // recLSN \leq r 's LSN \leq P 's pageLSN
 - Update P 's entry in DPT: recLSN = P 's pageLSN + 1