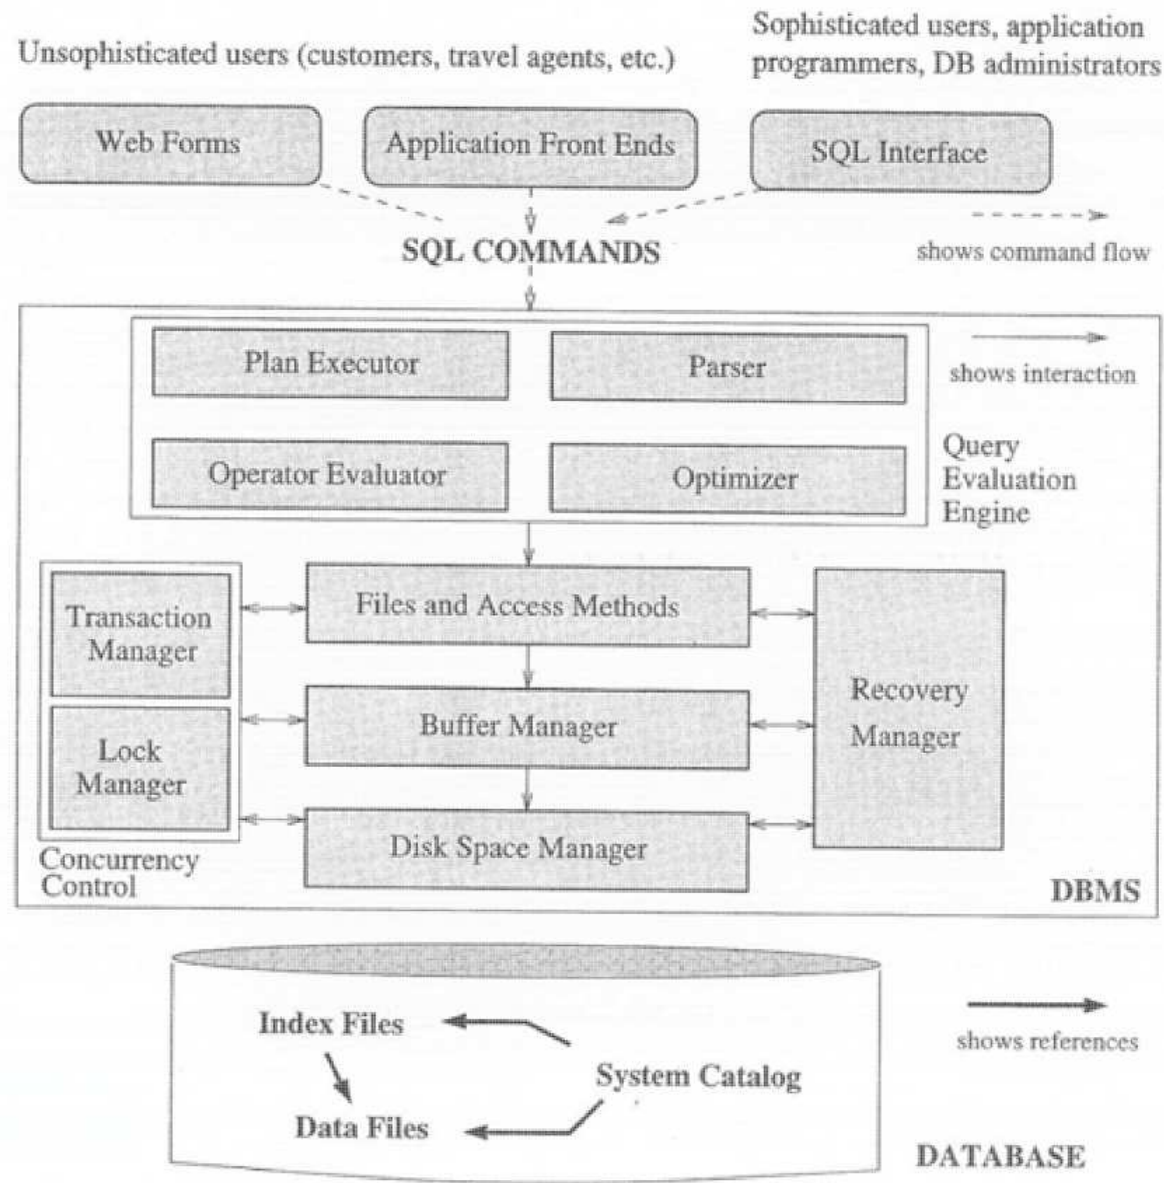


CS3223 Lecture 11

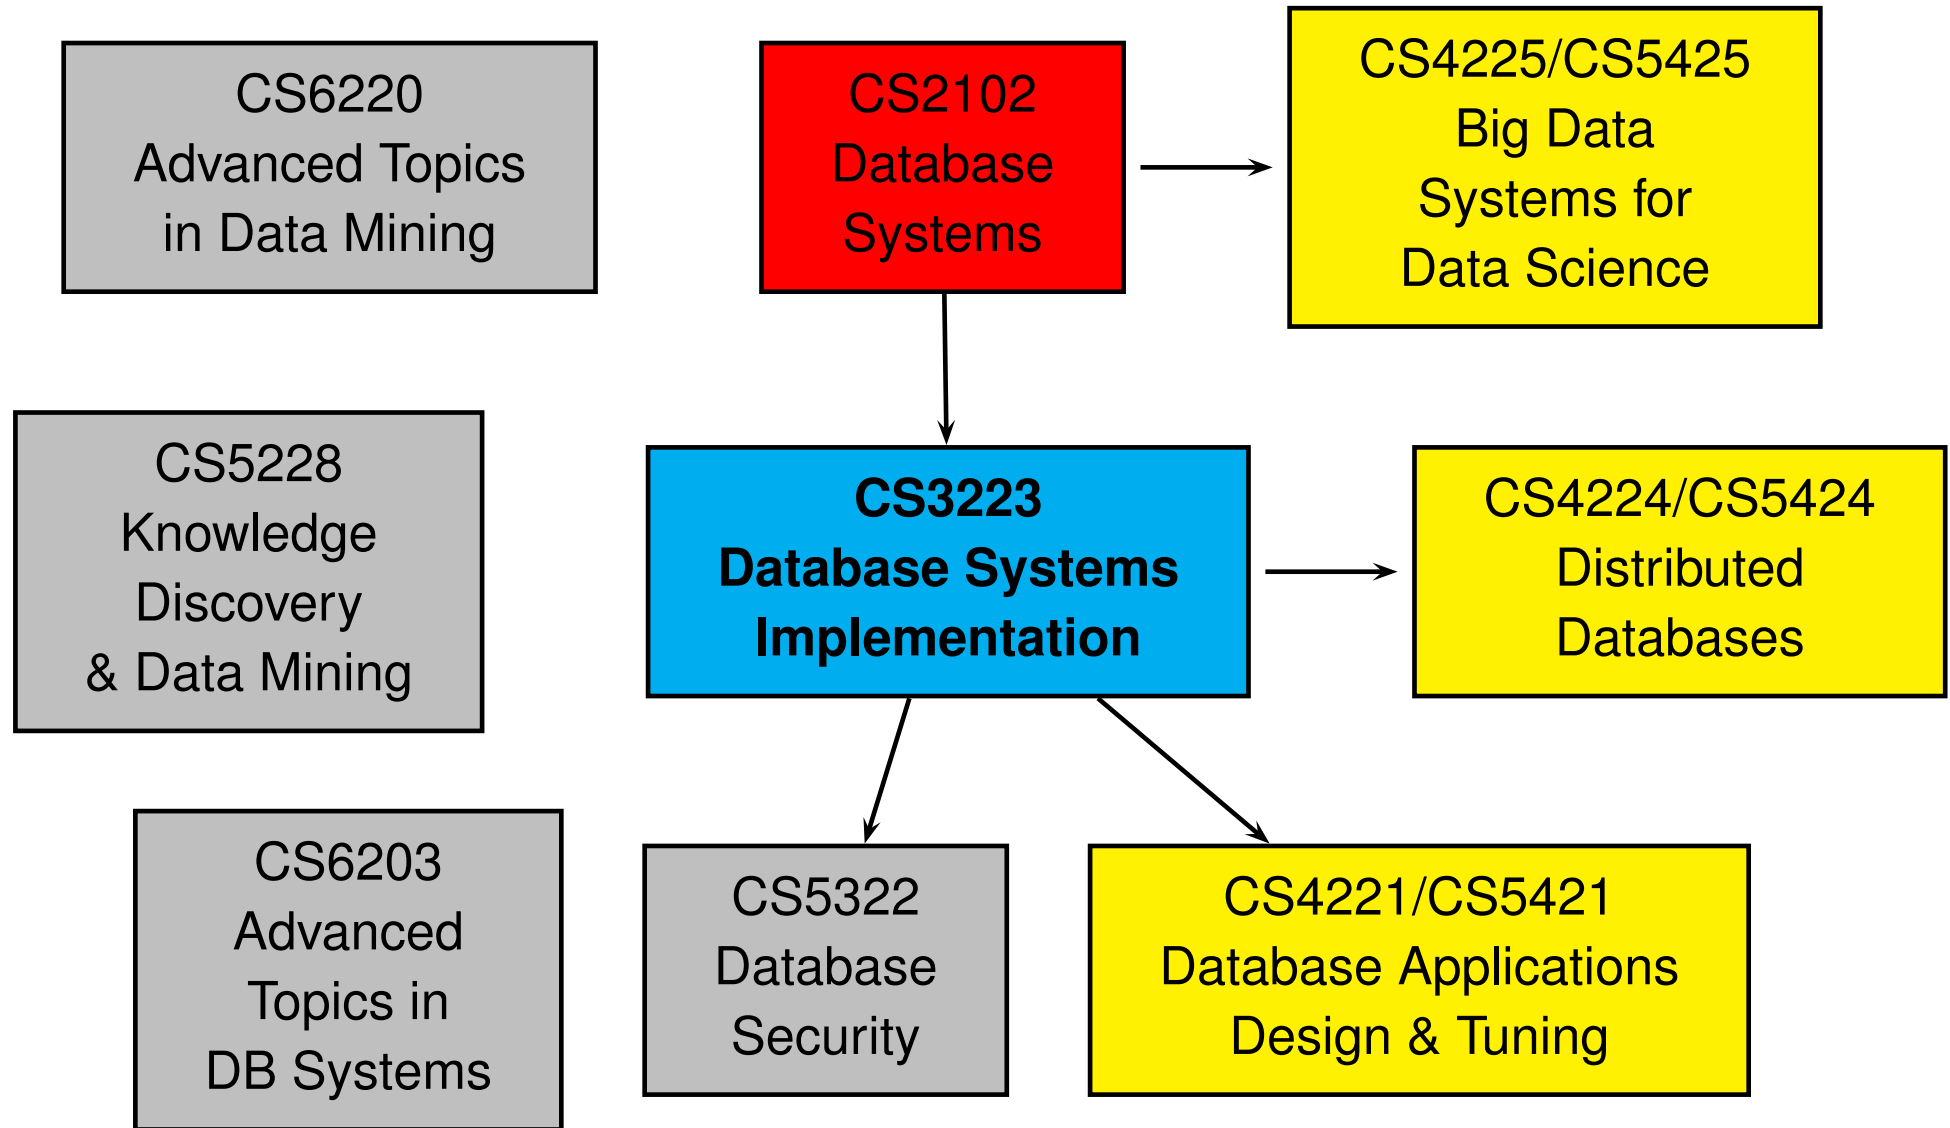
Review

Architecture of DBMS



(Ramakrishnan & Gehrke, Figure 1.3)

Database Courses @ SoC



Part 1: Review

Storage & Access Methods

► Buffer Manager

- ▶ Pin / unpin interface
- ▶ Replacement policies: Clock

► Storage Manager

- ▶ Heap file: linked-list vs page directory
- ▶ Page formats
 - ★ fixed-length records: packed vs unpacked organization
 - ★ variable-length records: slotted-page organization
- ▶ Record formats: delimiters vs offsets for variable-length records

► Access Methods

- ▶ Access patterns: sequential vs random I/O
- ▶ B⁺-tree
- ▶ Linear hashing, Extendible hashing

Query Evaluation

- ▶ **Query plan evaluation**: materialized / pipelined
- ▶ **Sorting** - external merge sort
- ▶ **Selection** - table scan, index scan / intersection
 - ▶ primary conjuncts, covered conjuncts
- ▶ **Projection**
 - ▶ Sort-based, optimized sort-based,
 - ▶ hash-based
- ▶ **Join**
 - ▶ block nested loop, index nested loop,
 - ▶ sort-merge join, optimized sort-merge join,
 - ▶ Grace hash join
- ▶ **Set operations & Aggregation**
 - ▶ Sort-based, hash-based
- ▶ Possibility of exploiting index for operator evaluation

Query Optimization

► Search Space

- linear plans
- bushy plans

► Plan Enumeration

- Query rewriting based on equivalence rules
- Dynamic programming algorithms
 - ★ Basic dynamic programming approach
 - ★ Enhanced dynamic programming approach

► Cost Model

- Size estimation assumptions: uniformity, independence, inclusion
- Histograms: equiwidth, equidepth, equidepth with MCV

Transaction Management

- ▶ Atomicity, Consistency, Isolation, Durability
- ▶ Concurrency Control Manager: Isolation
 - ▶ Anomalies: dirty read, unrepeatable read, lost update, phantom read
- ▶ Recovery Manager: Atomicity, Durability
 - ▶ Deals with commit, abort, and restart operations
 - ▶ Recovery with before-images
 - ▶ Schedules: Strict \subsetneq Cascadeless \subsetneq Recoverable

Concurrency Control

► Theory

- View / Multiversion View Serializability
- Conflict Serializability
 - ★ Conflicting actions, CSG
- $CSS \subsetneq VSS \subsetneq MVSS$

► Protocols

- **Lock-based**: Two-Phase Locking (2PL), Strict 2PL (S2PL)
 - ★ lock modes, multigranularity locking
 - ★ lock upgrades
 - ★ deadlocks - detection, prevention
- **Multiversion**: Snapshot Isolation (SI), Serializable SI
 - ★ First-Committer-Wins Rule, First-Updater-Wins Rule

► Isolation Levels

- Read uncommitted, Read committed, Repeatable read, Serializable
- Lock-based protocols: Short/Long duration locks

Log-based Crash Recovery

► Interactions with buffer manager

- ▶ Steal policy: requires undo recovery operations
- ▶ No-force policy: requires redo operations

► Protocols

- ▶ Write-ahead logging
- ▶ Force-at-commit

► Checkpointing

- ▶ Simple Checkpointing
- ▶ Fuzzy Checkpointing

► ARIES

- ▶ Three phases: Analysis, Redo, Undo
- ▶ Data structures: transaction table & dirty page table
- ▶ Concepts: CLRs, PageLSN, prevLSN chain, RedoLSN, undoNextLSN chain

Part 2: Q&A

Pipelined Query Evaluation

- ▶ Materialize intermediate result only when it's beneficial
- ▶ Query Q1: `select * from R join S on R.a = S.a where R.c < 50`
 - ▶ Assume $|R| < |S|$
 - ▶ Consider evaluating Q1 using **Block Nested Loop Join** (i.e., R is outer relation)
 - ▶ No benefit to materialize $\sigma_{R.c < 50}(R)$
- ▶ Query Q2: `select * from R join S on R.a = S.a where R.c < 50 and S.d = 100`
 - ▶ Assume $|R| < |S|$
 - ▶ Consider evaluating Q2 using **Block Nested Loop Join** (i.e., R is outer relation)
 - ▶ No benefit to materialize $\sigma_{R.c < 50}(R)$
 - ▶ May be beneficial to materialize $\sigma_{S.d = 100}(S)$

Object Versions

- ▶ x_i denote that x_i is a version of x created by T_i
- ▶ Consider two versions of x : x_i & x_j , where $i < j$
 - ▶ Incorrect to conclude that x_i is an earlier version than x_j just because $i < j$!
- ▶ x_i is an earlier version than x_j iff T_i commits before T_j

SI: Transaction Dependencies

- ▶ **Transaction dependencies** are not the same as **conflicting actions**
- ▶ ww & wr dependencies must be non-concurrent
- ▶ rw dependencies could be non-concurrent or current
- ▶ $T_i \xrightarrow{rw} T_j$: order of $R_i(x)$ and $W_j(x)$ does not matter as T_i & T_j are concurrent

- ▶ Example 1:

T_1 : $R_1(a), R_1(x), R_1(b), Commit_1$
 T_2 : $W_2(x), Commit_2$

- ▶ Example 2:

T_1 : $R_1(a), R_1(x), R_1(b), Commit_1$
 T_2 : $W_2(x), Commit_2$